

# PHASM501 – Techniques of High-Performance Computing



# Formalities

- 30 hours lectures
- 10 hours tutorials

Times:

Monday 2 to 4pm  
Thursday 4 to 6pm

Assessment:

2 Courseworks (20% each)

One course project

Assignments: Hand in Week 6 and 10 (Friday 12 lunch time)

Project: Three phases with submission in weeks 5, 8 and 13 (also Friday 12 lunch time)

Tutorials scheduled as needed in these hours.



# Syllabus

## Part 1: HPC Technologies (6 hours)

- Design of HPC Systems
- Shared vs distributed memory parallelization
- Overview of HPC Programming Languages
- Python for High-Performance Computing
- An overview of PETSc and its Python interface

## Part 2: Large-scale computational grids (6 hours)

- Types of grids
- Setting up a grid, data structures, connectivity, etc.
- Grid Partitioning, balancing
- Reference elements
- Evaluating integrals over grids
- Example: Setting up a Laplace problem
- Software for working with grids

# Syllabus...

## Part 3: Sparse matrices (6 hours)

- From grids to sparse matrices
- Data storage schemes for sparse matrices
- Sparse Matrix Vector Products
- Sparse LU Decomposition
- Parallel sparse matrix operations

## Part 4: Iterative Linear Algebra (6 hours)

- Parallel Iterative Solvers
- Overview of Preconditioning techniques
- Algebraic multigrid preconditioning
- Iterative solvers and preconditioners in PETSc

## Part 5: Full worked out parallel implementation of the following problem types (6 hours)

- Solving a transport equation with finite volume methods
- Finite element solution of a boundary value problem
- Finite difference methods for acoustic wave propagation

# A big disclaimer

- This is only the second time this module is being run.
- We have a lot of freedom to shape the module not only for this term but also for the future.
- If you **do not like** something or have problems tell me straight away.
- Also tell me straight away if you **do like** something.

# Performance of modern hardware



Samsung Galaxy S6 Edge  
Linpack Benchmark: 0.8 GFlop/s

1GFlop/s =  $10^9$   
floating point  
operations per second

4Ghz Intel Core i7-6700K  
Linpack Benchmark: 250 GFlop/s



Intel Xeon E5-2699v3 (2.3Ghz, 18  
Cores)  
Linpack Benchmark:  $10^3$   
Gflop/s

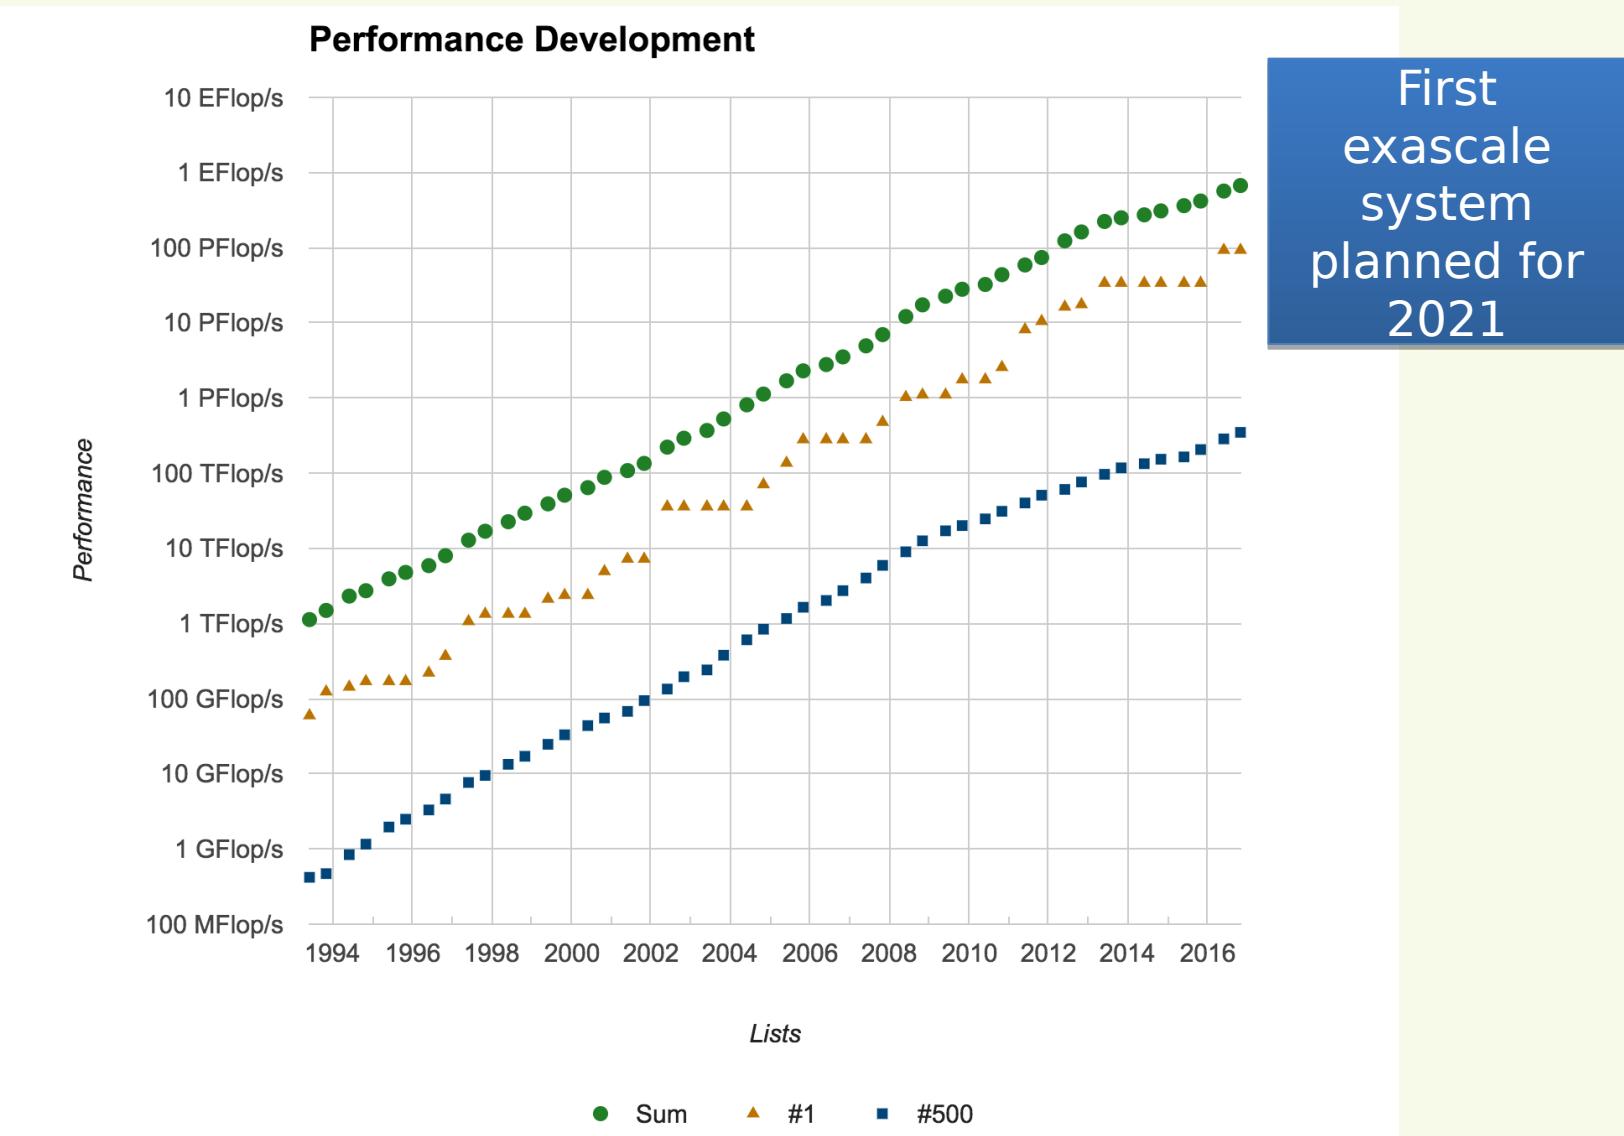
Sunway TaihuLight (Fastest computer in  
the world)  
Linpack Benchmark:  $93 \times 10^6$  GFlop/s



# The Top500 List

Rank	Site	System	Cores	Rmax (TFlop/s)	Rpeak (TFlop/s)	Power (kW)
1	National Supercomputing Center in Wuxi China	<b>Sunway TaihuLight</b> - Sunway MPP, Sunway SW26010 260C 1.45GHz, Sunway NRCPC	10,649,600	93,014.6	125,435.9	15,371
2	National Super Computer Center in Guangzhou China	<b>Tianhe-2 (MilkyWay-2)</b> - TH-IVB-FEP Cluster, Intel Xeon E5-2692 12C 2.200GHz, TH Express-2, Intel Xeon Phi 31S1P NUDT	3,120,000	33,862.7	54,902.4	17,808
3	DOE/SC/Oak Ridge National Laboratory United States	<b>Titan</b> - Cray XK7 , Opteron 6274 16C 2.200GHz, Cray Gemini interconnect, NVIDIA K20x Cray Inc.	560,640	17,590.0	27,112.5	8,209
4	DOE/NNSA/LLNL United States	<b>Sequoia</b> - BlueGene/Q, Power BQC 16C 1.60 GHz, Custom IBM	1,572,864	17,173.2	20,132.7	7,890
5	DOE/SC/LBNL/NERSC United States	<b>Cori</b> - Cray XC40, Intel Xeon Phi 7250 68C 1.4GHz, Aries interconnect Cray Inc.	622,336	14,014.7	27,880.7	3,939
6	Joint Center for Advanced High Performance Computing Japan	<b>Oakforest-PACS</b> - PRIMERGY CX1640 M1, Intel Xeon Phi 7250 68C 1.4GHz, Intel Omni-Path Fujitsu	556,104	13,554.6	24,913.5	2,719
7	RIKEN Advanced Institute for Computational Science (AICS) Japan	K computer, SPARC64 VIIIfx 2.0GHz, Tofu interconnect Fujitsu	705,024	10,510.0	11,280.4	12,660

# Performance over time



# Amdahl's Law

$$T = T_{parallel} + T_{serial}$$

$T$  : Overall time

$$= pT + (1 - p)T$$

$T_{parallel}$  : Execution time of parallelizable code

$T_{serial}$  : Execution time of serial code

$p$  : Fraction of parallelizable code

$$T(s) = \frac{p}{s}T + (1 - p)T$$

$s$  : Speedup of parallel execution

Maximum possible speedup:

$$\lim_{s \rightarrow \infty} \frac{T}{T(s)} = \frac{1}{1 - p}$$

Amdahl's law is too pessimistic as it assumes speedup for a constant program size. In practice larger parallel systems are used for larger problems.

# Gustafson's Law

$$W = (1 - p)W_{\text{serial}} + pW_{\text{parallel}}$$

$T$  : Overall time

$W_{\text{parallel}}$  : Parallelizable workload

$W_{\text{serial}}$  : Serial workload

$p$  : Fraction of parallelizable workload

$s$  : Increase in parallel workload at constant execution time

$$W(s) = (1 - p)W + spW$$

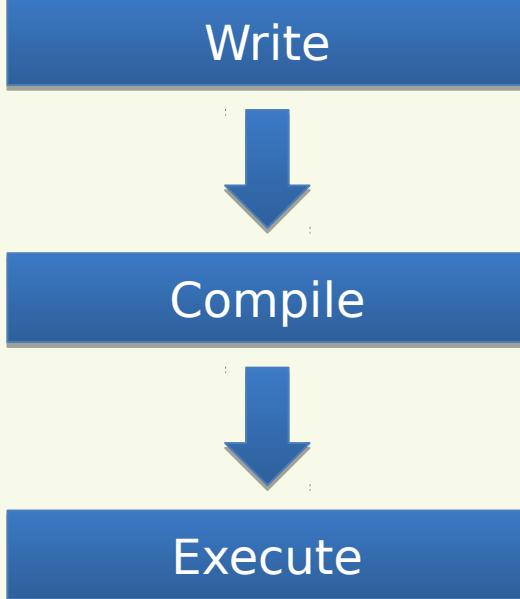
$$\frac{W(s)}{W} = 1 - p + sp$$

A parallel program may read parameters for a computation from disk (serial execution) and then distribute the workload onto a number of fully parallel threads.

# Programming Languages

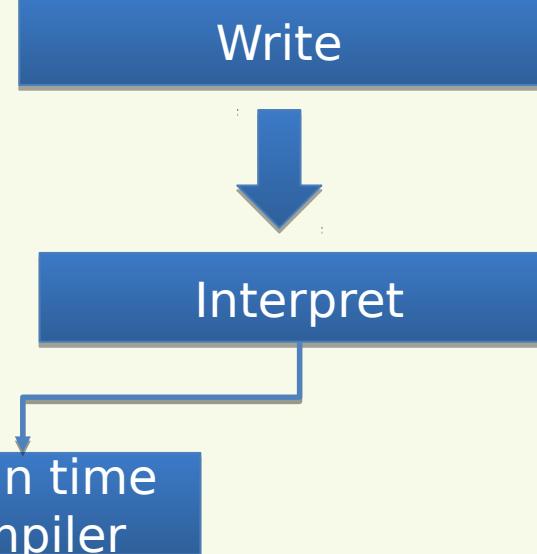
Compiled  
Languages

Fortran, C/C++,  
Java, C#



Interpreted  
Languages

Python, Julia,  
Matlab, R



# Fortran

- Oldest language still used for HPC
- Many traditional HPC libraries written in Fortran or have Fortran interfaces
- Big difference between Fortran 77 and newer Fortran variants
- Fortran is legacy. Do not use for new projects.
- Compilers: gfortran, ifort, pgfortran,

```
program hello
    implicit none
    write (*,*) "Hello,
world."
end program hello
```

...

# C

- Mostly used for low-level libraries and embedded sector (with some notable exceptions such as PETSc or OpenCL)
- C programs compile fast and are easy to get to compile on virtually any C compiler
- Compiler optimizations harder than Fortran (though modern compilers do a great job)
- Compilers: gcc, clang, icc, pgcc, etc.

```
#include <stdio.h>
```

```
int main(void) {  
    printf("Hello  
World\n");  return 0;  
}
```

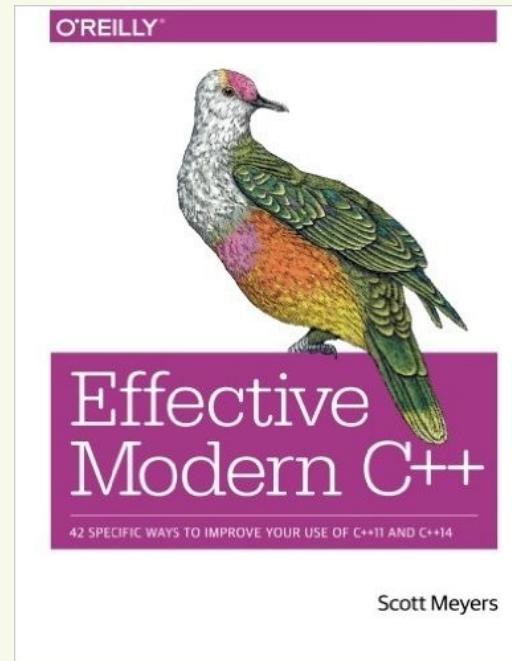
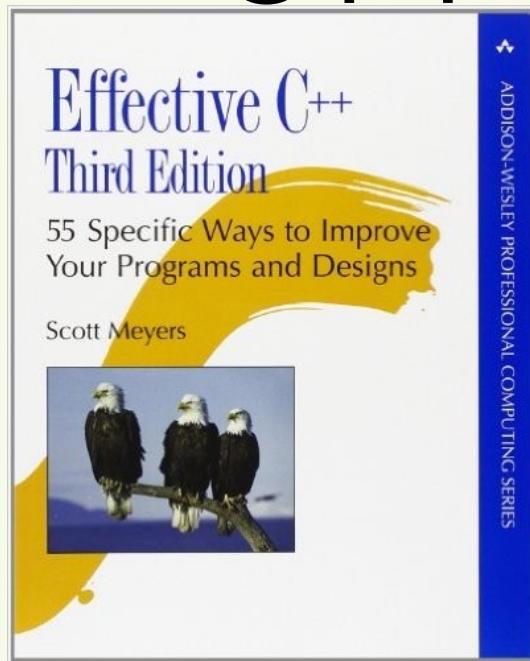
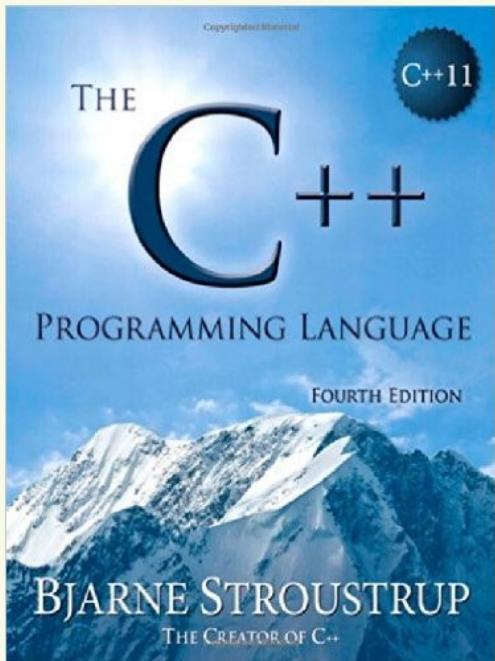
*“C makes it easy to shoot yourself in the foot; C++ makes it harder, but when you do it blows your whole leg off.”*

# C++

- Almost fully backward compatible to C.
- Implements modern object and templated design principles.
- Most new HPC projects use C++.
- Major redesign in C++11 standard
- Compilers: g++, clang, icpc, pgc++, ...

```
#include <iostream>
int main(void) {
    std::cout << "Hello World"
    return 0;
```

# Books for safe and modern C++



# Java and C#

- Rarely used for HPC applications
- Java and C# compile into an intermediate virtual machine language and are typically just-in-time compiled during execution
- Virtual machine environment encapsulates from the hardware but allows much fewer optimizations
- Widely used for business applications

# New compiled languages

- Google Go is a systems programming language. Designed by Google it has fast become popular
- Rust is designed as a much less error prone replacement for C++. Used by Mozilla for their new browser engine
- Swift is the new language by Apple for iOS development
- None of these languages have any traction in the HPC sector (maybe with the exception of Go for some cloud computing tools, but not for fast algorithms)

# Python

- High-Level dynamic programming language
- Supports procedural, object and functional programming
- Very easy to interface with C/C++ and Fortran
- Python itself is slow. But it's extension modules make it extremely powerful for HPC

```
print("Hello world")
```

# Julia

- Modern dynamic high-performance programming language
- Still very young, but fast getting a lot of traction, especially in big data
- Julia is designed as fast, just-in-time compiled language.
- Claims to achieve near C-level performance

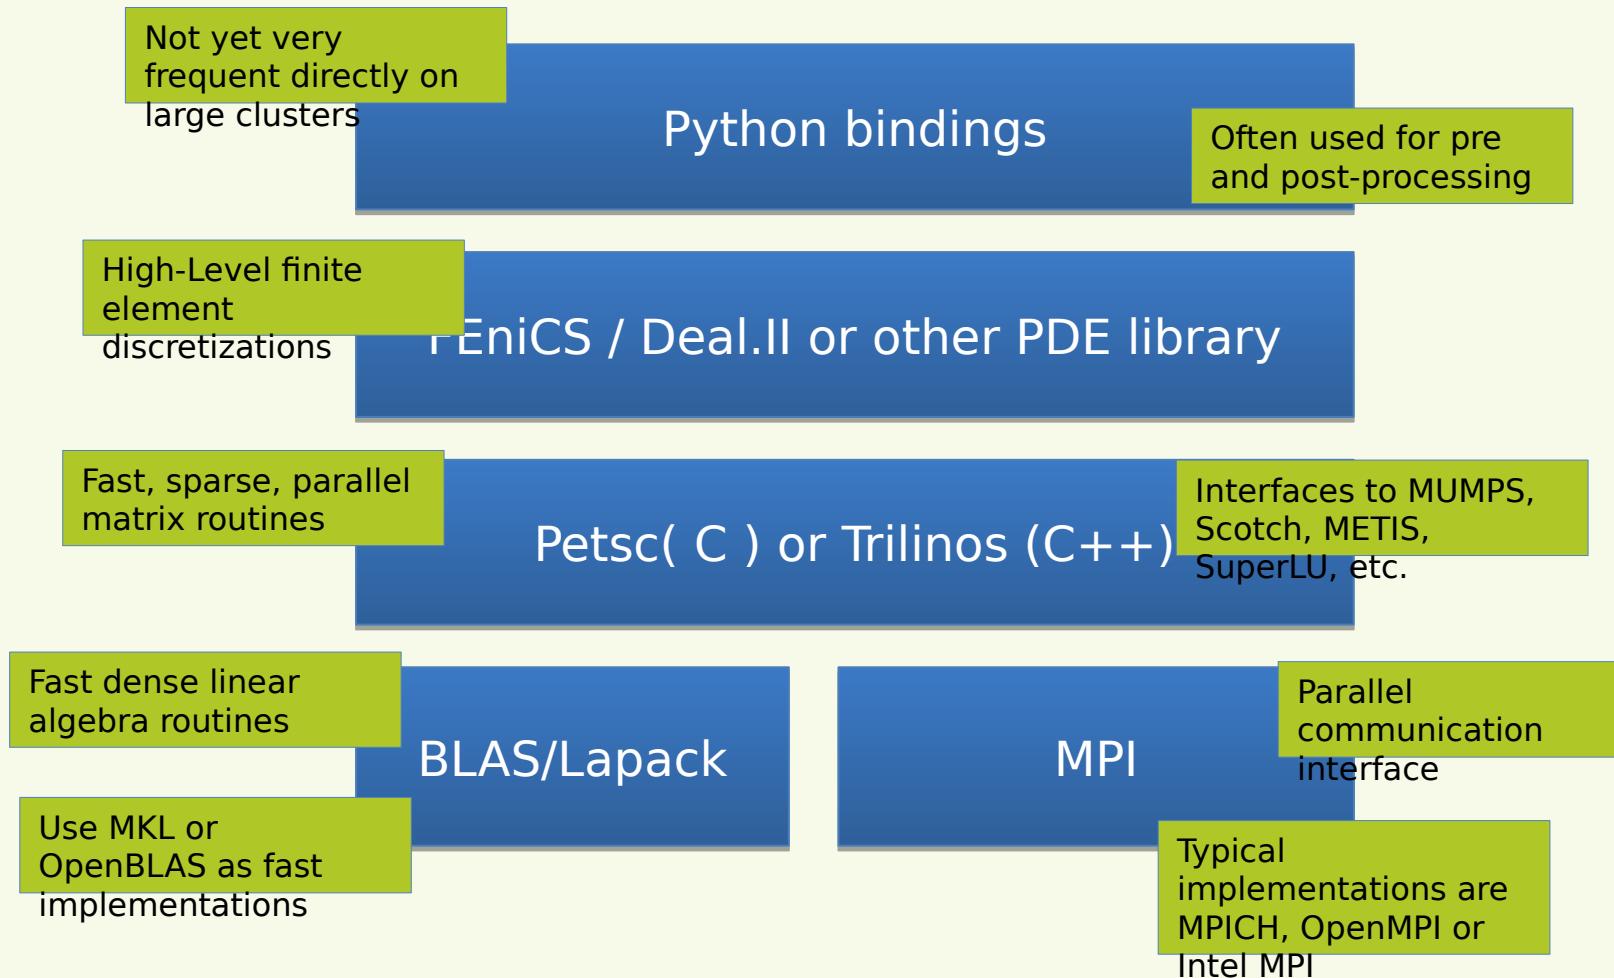
```
println("Hello  
world")
```

# Matlab

- One of the first dynamic numerical computing language
- Originally designed as convenient interface to Fortran libraries
- Matlab extremely wide spread and suitable for a range of HPC applications
- Huge amount of toolboxes for various industry sectors
- Not open-source

```
disp("Hello world")
```

# A typical HPC Stack

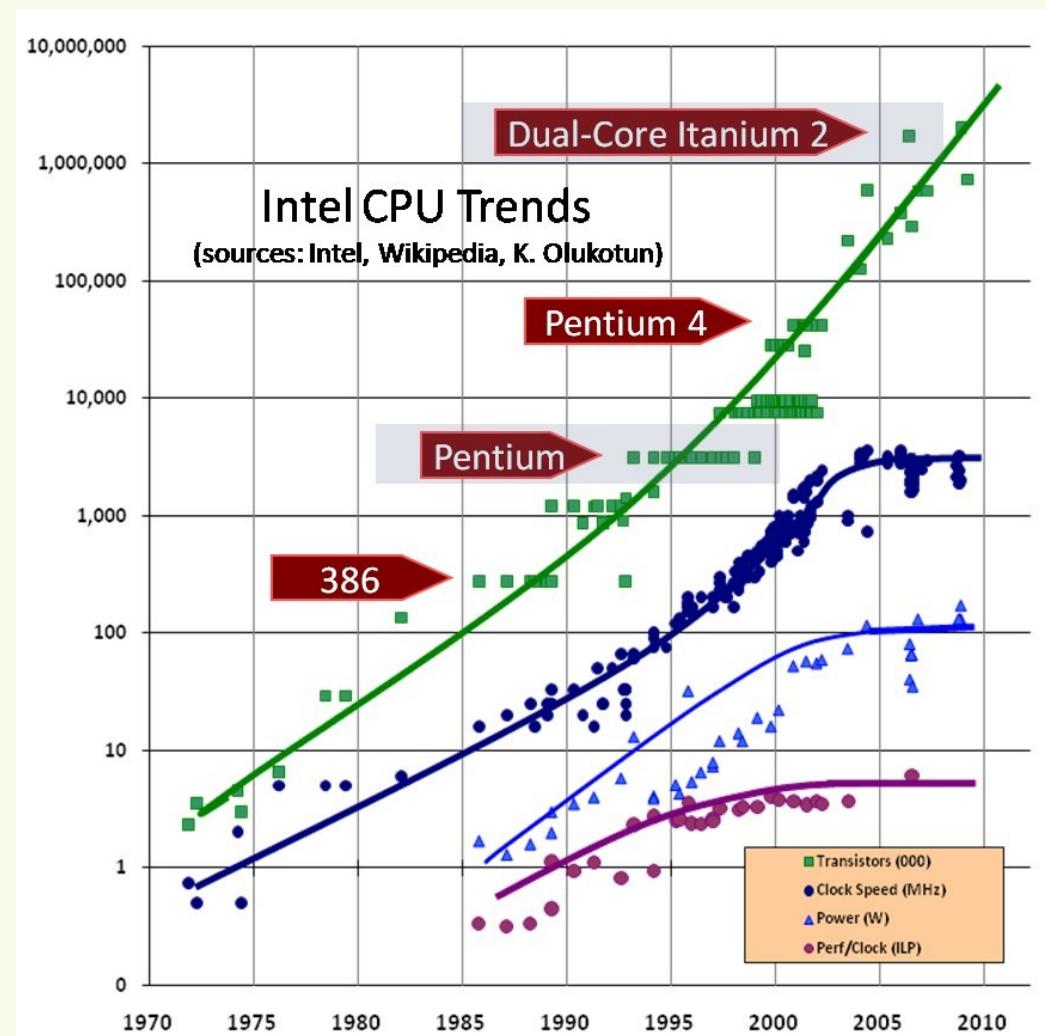


# How do I get it?

- Easiest way: Download Ubuntu Linux. All relevant software available in repositories
- Hard way: Compile everything yourself together with performance libraries (e.g. Intel MKL) of your choice
- Unfortunately, we cannot give you a Petascale machine for playing around...
- For Lecture: Download Anaconda or Intel Python (more in the first tutorial)

# Why parallel processing?

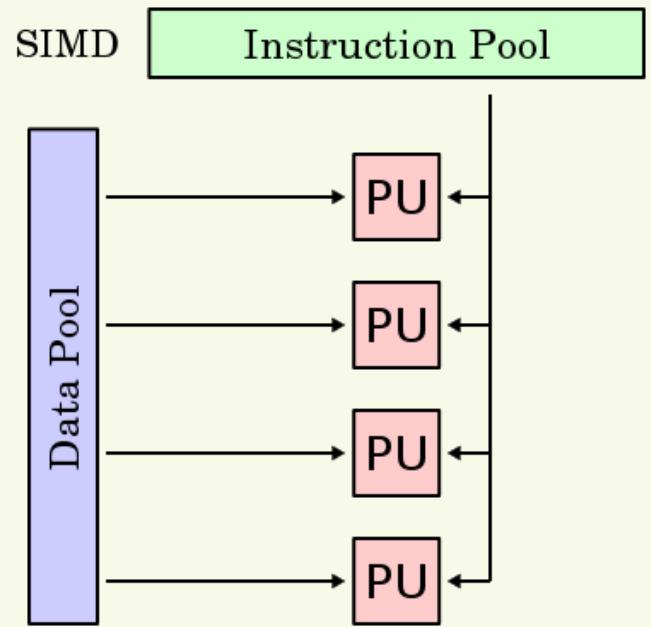
- Number of instructions per second on CPUs cannot scale arbitrarily (energy consumption, heat dissipation, material design limits, etc.)
- Most CPUs run at around 2 to 3GHz. This has not changed for years
- Further performance increase only through massive parallelism



# Parallel Concepts...

SIMD: Single Instruction – Multiple Data

One set of instructions is executed in parallel on a vector of data units (e.g. numbers or matrix rows).



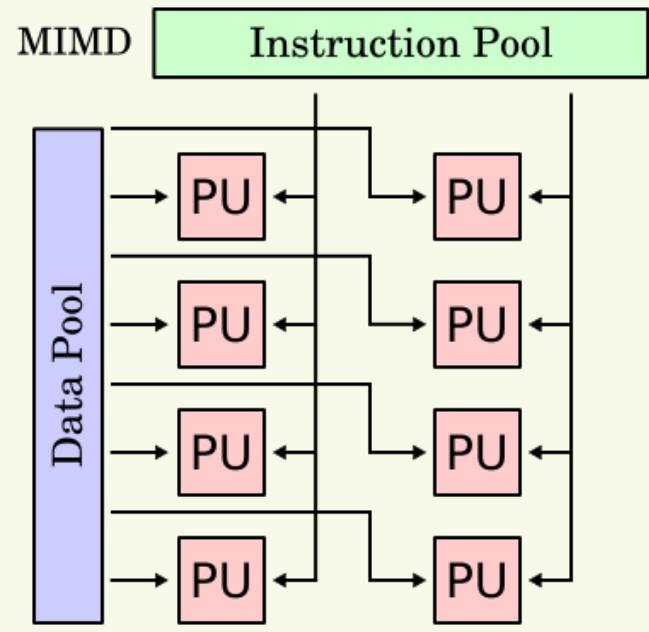
Wikipedia: Flynn's Taxonomy

- Popular concept until the 1990s in HPC Architectures (Vector processors)
- Came out of fashion due to cluster computing
- Recent revival with streaming instruction sets in modern CPUs (Intel SSE, etc.)

# Parallel Concepts...

MIMD: Multiple Instructions - Multiple Data

Each CPU unit can asynchronously execute different types of commands on different data sets.



Wikipedia: Flynn's Taxonomy

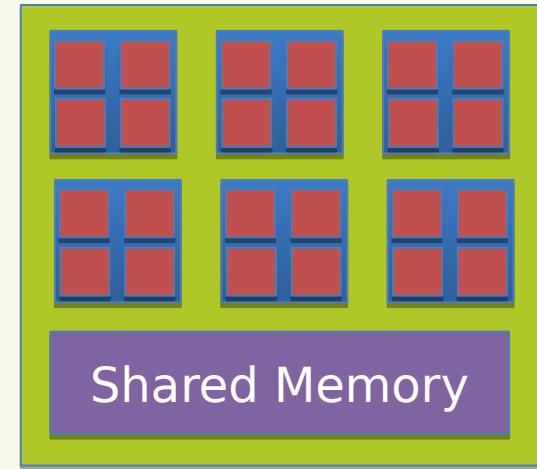
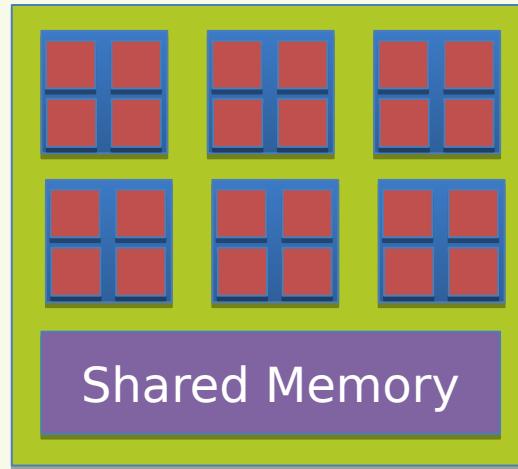
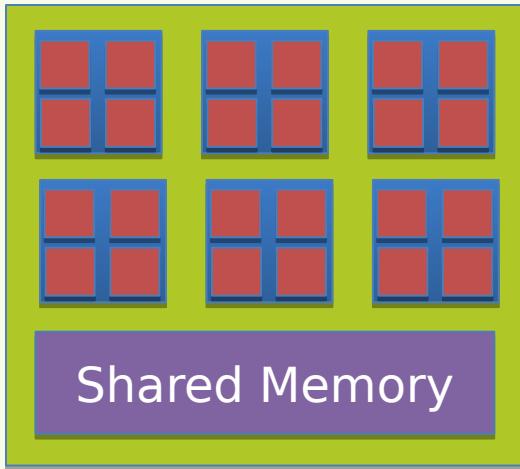
- All modern HPC systems are based on the MIMD design
- All modern multi-core CPUs use MIMD

Cluster Node

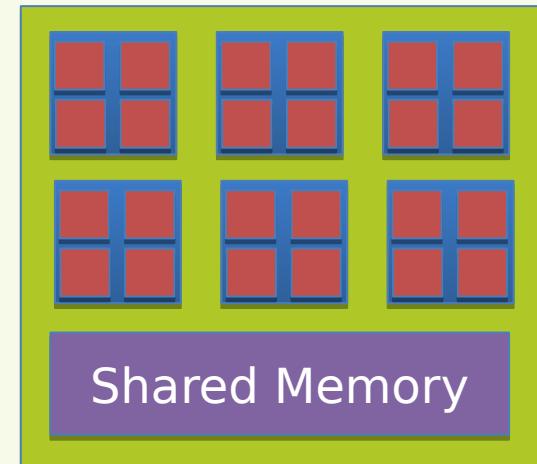
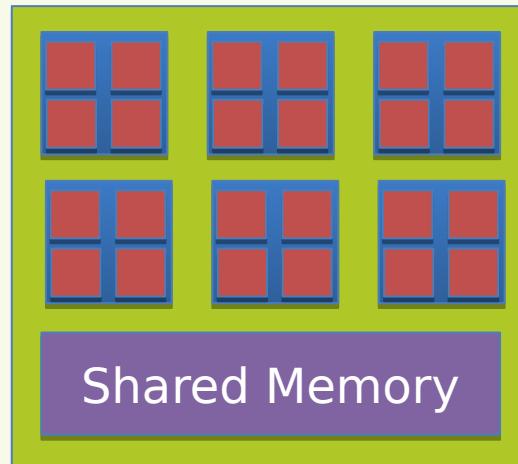
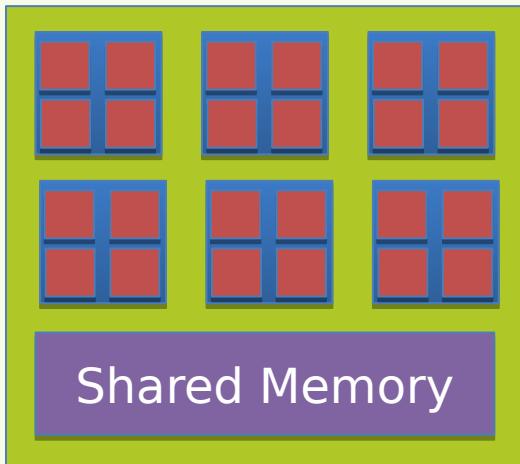
CPU

CPU Core with vector  
unit

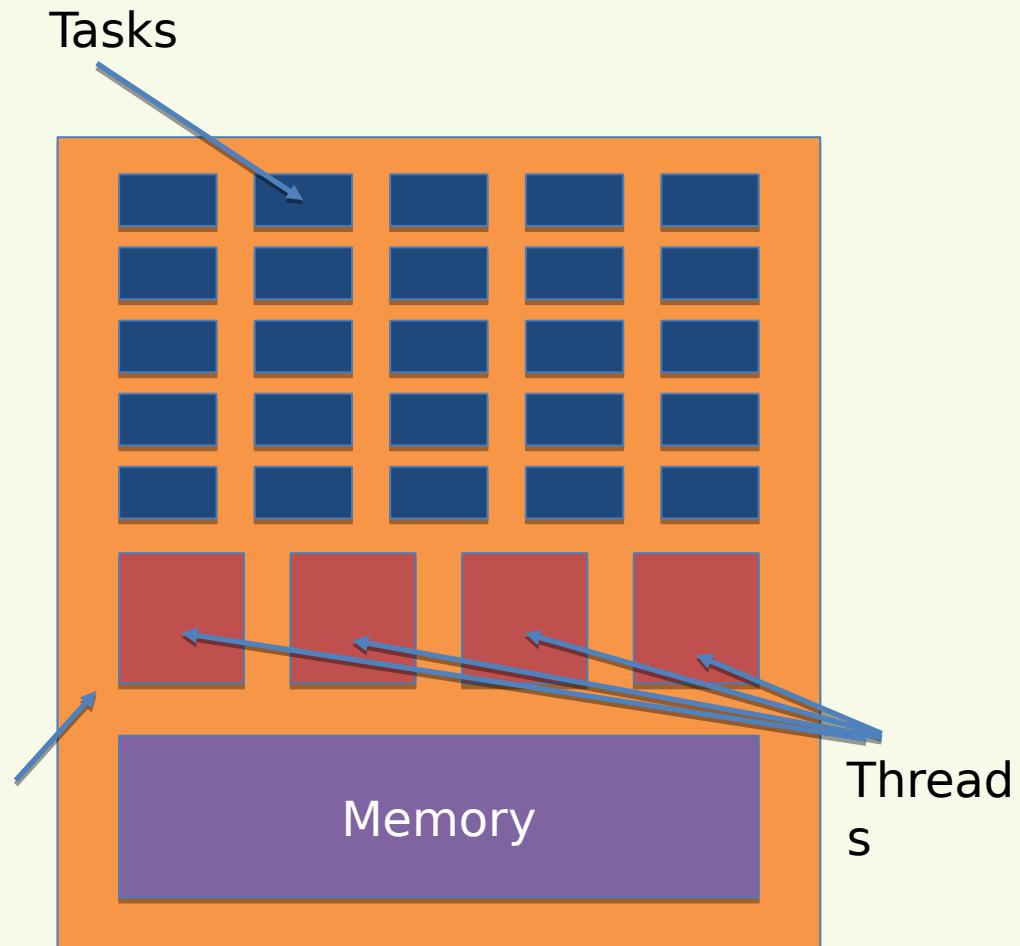
# A typical HPC system



Fast Interconnect



# Shared Memory parallelization (multithreading)



Process

- Process: Independent program with its own allocated memory running on a physical cluster node.
- Thread: Independent unit of execution running within a thread.
- Task: Logical unit of execution.
- Different processes can not see each others memory.
- All threads within a process see the same memory.

# Shared Memory Frameworks

- OpenMP: Industry standard. Supported by most compilers directly. Works with Fortran, C/C++. Cumbersome for complex task dependencies.
- Intel Threading Building Blocks (TBB): Versatile C++ library for task based threading. Parallel Container constructs. Needs to be linked to explicitly.
- Intel Cilk: Simple tasking framework for C/C++. Supported by more and more compilers.

# OpenMP

```
1 #include <iostream>
2
3 int main() {
4
5 #pragma omp parallel
6     { std::cout << "Hello World" << std::endl; }
7 }
8
```

Define a parallel region

- Everything inside the `#pragma omp parallel` block is executed in parallel.
- By default as many parallel threads are created as there are virtual CPU cores.
- The number of threads can be controlled by the environment variable `OMP_NUM_THREADS`.

# Parallel for loops

```
1 #include <iostream>
2 #include <vector>
3
4 int main() {
5
6     int count = 100;
7     std::vector<int> vec(100);
8
9 #pragma omp parallel for
10    for (std::size_t i = 0; i < count; i++) {
11        vec[i] = i * i;
12    }
13    return 0;
14 }
15
```

- The statements within the for loop are executed in parallel.
- Most frequent use form for OpenMP.

# Race conditions

```
1 #include <iostream>
2 #include <vector>
3
4 int main() {
5
6     int count = 100;
7     std::vector<int> vec(100);
8     int sum = 0;
9
10    #pragma omp parallel for
11        for (std::size_t i = 0; i < count; i++) {
12            vec[i] = i * i;
13            // The following code is wrong.
14            // The result of the operation is completely random.
15            sum += vec[i];
16        }
17        std::cout << sum << std::endl;
18        return 0;
19    }
20 }
```

sum is shared among all the threads. The result of the addition becomes random if more than one thread is running.

# What can happen?

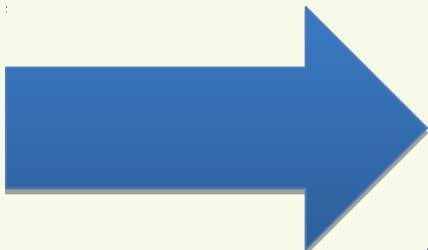
Example:

```
int i = 0;  
  
#pragma omp parallel  
{  
    i += 1;  
}
```

	Thread 1	Thread 2	Value of i
	Read i into register		0
	Add 1 to register	Read i into register	
	Write result back into i	Add 1 to register	1
		Write result back into i	1

The result becomes indetermined.

We need to ensure that only one thread at a time writes into a shared variable.



Fix with omp critical region

# Critical Regions

```
1 #include <iostream>
2 #include <vector>
3
4 int main() {
5
6     int count = 100;
7     std::vector<int> vec(100);
8     int sum = 0;
9
10    #pragma omp parallel for
11        for (std::size_t i = 0; i < count; i++) {
12            vec[i] = i * i;
13            // Only one thread at a time enters the critical region
14            #pragma omp critical
15                { sum += vec[i]; }
16            }
17            std::cout << sum << std::endl;
18            return 0;
19        }
```

- Only one thread at a time is allowed to enter the critical region.
- The race condition is resolved.
- Slow down as threads have to wait.

# General Concept: Mutex

- A Mutex is the general concept underlying the critical region. The following abstract code shows its use.

```
Mutex mutex;
```

```
#Enter parallel region  
lock(mutex)  
i += 1  
release(mutex)  
#Exit parallel region
```

If mutex is locked a thread has to wait. It can only enter if it is not locked already.

Release the mutex again when thread is finished.

# Private and shared variables

- By default variables declared outside a parallel region are shared.
- Variables declared inside a parallel region are private copies for each thread.
- Can control the behavior with the statements: private, shared, firstprivate, lastprivate

# Nested parallelism

```
1 #include <cstdio>
2 #include <iostream>
3 #include <omp.h>
4 #include <vector>
5
6 int main() {
7
8     const int count = 10;
9     double a[count][count];
10
11 #pragma omp parallel
12 {
13     if (omp_get_thread_num() == 0)
14         printf("Outer number of threads: %i\n", omp_get_num_threads());
15 #pragma omp parallel
16 {
17     if (omp_get_thread_num() == 0)
18         printf("Inner number of threads: %i\n", omp_get_num_threads());
19 }
20 }
21 return 0;
22 }
```

- The outer region runs with 8 threads
  - The inner region runs with 1 thread
  - Prevents over-utilisation
  - Can be controlled with ‘schedule’ directive

## Output:

Inner number of  
threads: 1

Inner number of  
threads: 1

Outer number of  
threads: 8

Inner number of  
threads: 1

# Tasking

- Direct mapping of loop iterations onto threads is cumbersome for complex scenarios
- Nested parallelisation is difficult to control
- Better idea: Separate threads and tasks
- OpenMP supports tasking since OpenMP 3

# Fibonacci Sequence

- $\text{Fib}(n) = \text{Fib}(n-1) + \text{Fib}(n-2)$
- $\text{Fib}(0) = \text{Fib}(1) = 1$
- 1, 1, 2, 3, 5, 8, 13, 21, ...
- Standard Loop based threading is not suitable to represent recursive computations
- However, with tasks the problem becomes easy...

# Fibonacci with OpenMP tasks

```
1 #include <iostream>
2
3 int fib(int n) {
4
5     if (n == 0 || n == 1)
6         return 1;
7     else {
8         int f1, f2;
9 #pragma omp task shared(f1) firstprivate(n)
10     f1 = fib(n - 1);
11 #pragma omp task shared(f2) firstprivate(n)
12     f2 = fib(n - 2);
13 #pragma omp taskwait
14     return f1 + f2;
15 }
16 }
17
18 int main() {
19
20     int n = 5;
21 #pragma omp parallel shared(n)
22     {
23 #pragma omp single
24         { std::cout << fib(n) << std::endl; }
25     }
26     return 0;
27 }
```

Start the thread pool

Enter the function with one thread

- `#omp task`: Schedule a new task
- `shared`: The variable is shared with the task
- `firstprivate`: A local copy is created from the previous value
- `taskwait`: Wait until all tasks complete

Warning: This code is for demonstration only. The tasks have far too little to do to be efficient!

# Summary of OpenMP

- Simple loop parallelism is trivial to do in OpenMP and can often give significant speed-ups
- OpenMP supports tasking for complicated parallelization strategies
- Almost every compiler supports it
- OpenMP is not a native language construct. Visibility of variables and other controls can be cumbersome and debugging difficult

# Intel TBB: Advanced threading in C++

+

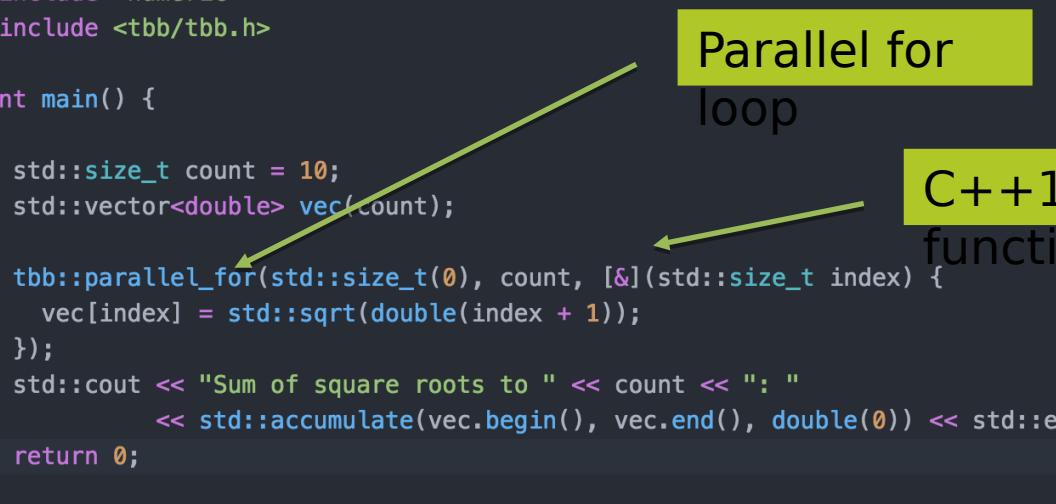
- Advanced threading library written in C++
- Available as open-source and commercially supported version
- Fully task-based design
- Support for advanced C++11 features
- Threadsafe container types
- Straight forward syntax
- Explicit linking of TBB libraries in projects necessary

# Simple parallel for loops in TBB

```
1 #include <cmath>
2 #include <iostream>
3 #include <numeric>
4 #include <tbb/tbb.h>
5
6 int main() {
7
8     std::size_t count = 10;
9     std::vector<double> vec(count);
10
11    tbb::parallel_for(std::size_t(0), count, [&](std::size_t index) {
12        vec[index] = std::sqrt(double(index + 1));
13    });
14    std::cout << "Sum of square roots to " << count << ": "
15          << std::accumulate(vec.begin(), vec.end(), double(0)) << std::endl;
16
17    return 0;
18 }
```

Parallel for loop

C++11 anonymous (lambda) function



- `tbb::parallel_for(start, end, fun)` executes in parallel the function `fun` for the indices `[start, end]`.
- Internally, a task tree is created that is scheduled onto a thread pool.
- The parallel region can be defined as anonymous function, callable object or any other valid callable construct.

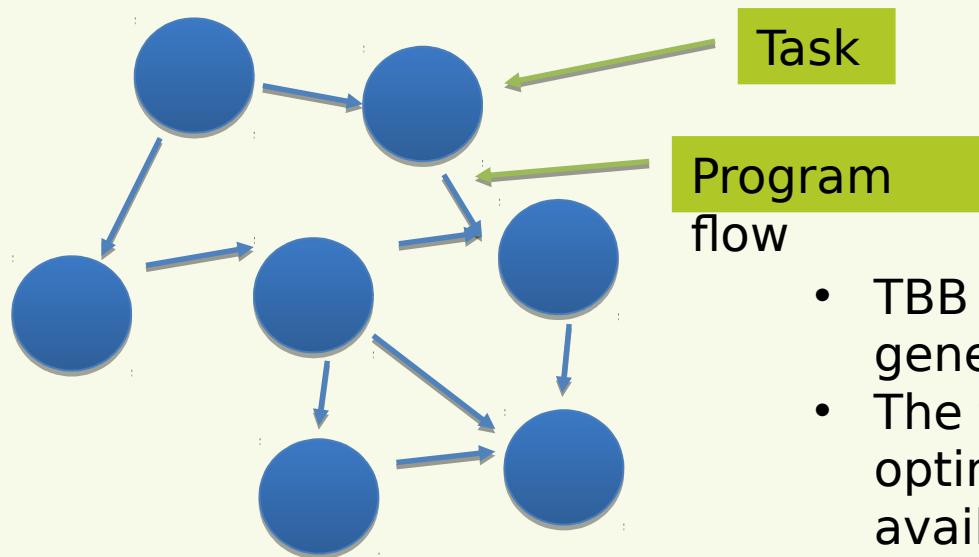
# Fibonacci series in TBB

```
1 #include <iostream>
2 #include <tbb/tbb.h>
3
4 int fib(int n) {
5     if (n == 0 || n == 1)
6         return 1;
7     else {
8         int f1, f2;
9         tbb::task_group group;
10        group.run([&] { f1 = fib(n - 1); });
11        group.run_and_wait([&] { f2 = fib(n - 2); });
12        return f1 + f2;
13    }
14 }
15
16 int main() {
17
18     int n = 5;
19     std::cout << "Fib(" << n << ") = " << fib(n) << "." << std::endl;
20     return 0;
21 }
22 }
```

- Create a task group with two tasks and submit the tasks
- Task scheduler collects the tasks in a tree and schedules them optimally on the threads

# Task graphs

- Complex task dependencies can be represented as a graph



- TBB supports the explicit generation of task trees
- The tree is analyzed and then optimally scheduled on the available threads
- Allows complex execution scenarios

# GPU Computing with OpenCL



28 Cores

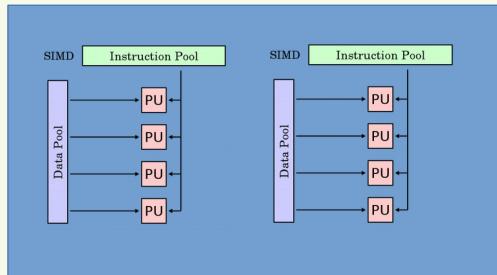
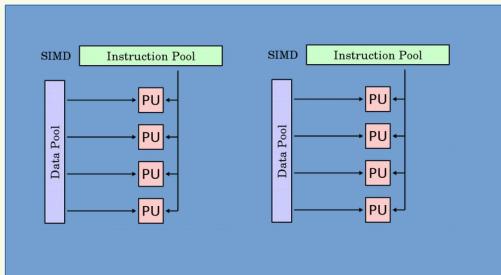
3.2 Ghz

2 AVX 512 FMA Units

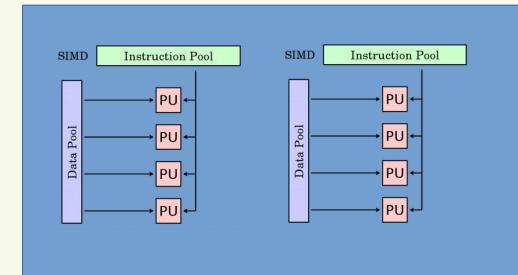
Peak (Single): 2.86 TFlops/s

Peak (Double): 1.43 TFlops/s

How to compute the peak performance?



—



- Each core has two SIMD Units (AVX 512) With Fused Multiply Add (FMA) each SIMD unit can perform up to 32 single precision or 16 double precision operations simultaneously.

# GPU Peak Performance



NVIDIA Tesla V100

5120 CUDA Cores

80 Streaming Multiprocessors

Peak (Single): 14 TFLOPS

Peak (Double): 7 TFLOPS



AMD Radeon Pro WX 9100

4096 Stream Processors

64 Compute Units

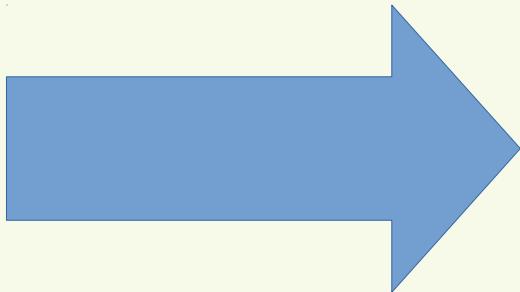
Peak (Single): 12.29 TFLOPS

Peak (Double): 0.77 TFLOPS

- CPU Core roughly similar to Streaming Multiprocessor (NVIDIA) or Compute Unit (AMD/OpenCL). CUDA Core, Stream Processors correspond to one execution in a SIMD Unit.
- Xeon Platinum has  $28 * 2 * 16 = 896$  Compute Units in single precision without FMA (twice that with FMA)
- Inside a Compute Unit (Streaming Multiprocessor) all execution paths must be identical.

# CPU / GPU Differences

- CPUs optimised for general purpose computing tasks. Great for complex, irregular data structures, codes with a lot of branching.
- CPU have complex caching architecture and branch prediction.
- GPUs much more simple designed. Optimised for parallel processing of large blocks of structured data with the same operations (SIMD Execution).
- Data transfer to GPU is expensive via PCI bus. Algorithms need to be adapted for it.



AMD APU: Combine CPU and GPU unit on a single die. Already used in Xbox One and PS4.

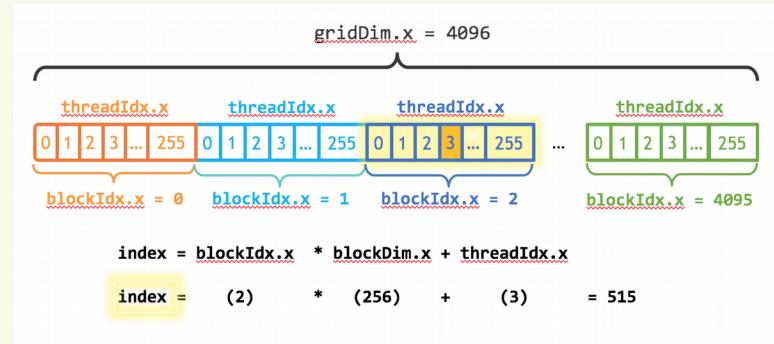
Don't believe all marketing hype. The choice between CPU and GPU is not clear cut. In the medium term we will see a range of hybrid solutions.

# GPU Programming Models

- NVIDIA Cuda: Started in 2007, based on C++ with proprietary extensions. NVIDIA CUDA SDK required. Only runs on NVIDIA GPUs.
- OpenCL: Started in 2009, C based open standards. Open-Source implementations exist. Runs on a range of hardware devices.
- DirectCompute: Windows only library as part of Microsoft DirectX.

# A simple Cuda C example

```
#include <iostream>
#include <math.h>
// Kernel function to add the elements of two arrays
__global__
void add(int n, float *x, float *y)
{
    int index = blockIdx.x * blockDim.x + threadIdx.x;
    int stride = blockDim.x * gridDim.x;
    for (int i = index; i < n; i += stride)
        y[i] = x[i] + y[i];
}
int main(void)
{
    int N = 1<<20;
    float *x, *y;
    // Allocate Unified Memory - accessible from CPU or GPU
    cudaMallocManaged(&x, N*sizeof(float));
    cudaMallocManaged(&y, N*sizeof(float));
    // initialize x and y arrays on the host
    for (int i = 0; i < N; i++) {
        x[i] = 1.0f;
        y[i] = 2.0f;
    }
    // Run kernel on 1M elements on the GPU
    int blockSize = 256;
    int numBlocks = (N + blockSize - 1) / blockSize;
    add<<<numBlocks, blockSize>>>(N, x, y); // Wait for GPU to finish before accessing on host
    cudaDeviceSynchronize();
    // Check for errors (all values should be 3.0f)
    float maxError = 0.0f;
    for (int i = 0; i < N; i++)
        maxError = fmax(maxError, fabs(y[i]-3.0f));
    std::cout << "Max error: " << maxError << std::endl;
    // Free memory
    cudaFree(x);
    return 0;
}
```



# An OpenCL example

<https://www.olcf.ornl.gov/tutorials/opencl-vector-addition/>

(Code too big to show on slide)

# OpenCL remarks

- OpenCL is very verbose. It can be painful to setup a simple kernel in low-level OpenCL.
- Reason: OpenCL needs to work on a range of hardware devices. Cuda just needs to run on NVIDIA.
- In this module we will bypass all the pain by running everything through PyOpenCL, which does all the book keeping automatically.

# A PyOpenCL Example

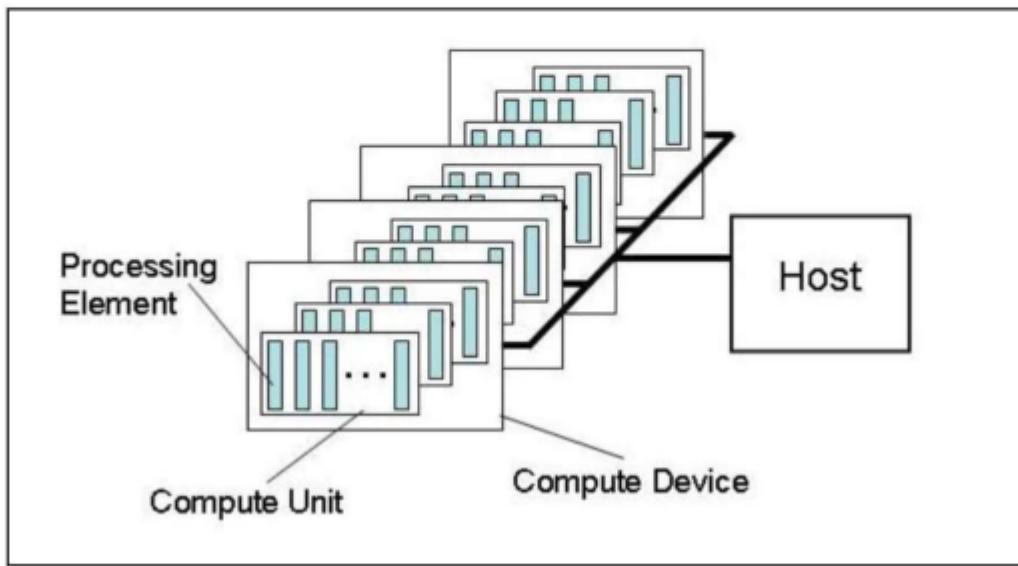
```
#!/usr/bin/env python
# -*- coding: utf-8 -*-
from __future__ import absolute_import, print_function
import numpy as np
import pyopencl as cl
a_np = np.random.rand(50000).astype(np.float32)
b_np = np.random.rand(50000).astype(np.float32)
ctx = cl.create_some_context()
queue = cl.CommandQueue(ctx)
mf = cl.mem_flags
a_g = cl.Buffer(ctx, mf.READ_ONLY | mf.COPY_HOST_PTR, hostbuf=a_np)
b_g = cl.Buffer(ctx, mf.READ_ONLY | mf.COPY_HOST_PTR, hostbuf=b_np)
prg = cl.Program(ctx, """
__kernel void sum(
    __global const float *a_g, __global const float *b_g, __global float *res_g)
{
    int gid = get_global_id(0);
    res_g[gid] = a_g[gid] + b_g[gid];
}
""").build()
res_g = cl.Buffer(ctx, mf.WRITE_ONLY, a_np.nbytes)
prg.sum(queue, a_np.shape, None, a_g, b_g, res_g)
res_np = np.empty_like(a_np)
cl.enqueue_copy(queue, res_np, res_g)
# Check on CPU with Numpy:
print(res_np - (a_np + b_np))
print(np.linalg.norm(res_np - (a_np + b_np)))
```

<https://document.tician.de/pyopencl/>

# OpenCL Platform Model

- **One Host + one or more Compute Devices**

- Each Compute Device is composed of one or more Compute Units
  - Each Compute Unit is further divided into one or more Processing Elements



© Copyright Khronos Group, 2012 - Page 1

# OpenCL Memory Model

- **Private Memory**

- Per work-item

- **Local Memory**

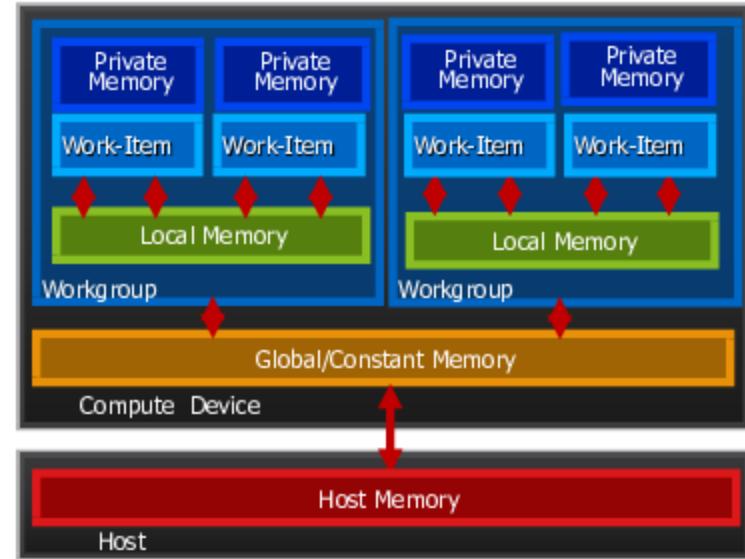
- Shared within a workgroup

- **Global/Constant Memory**

- Visible to all workgroups

- **Host Memory**

- On the CPU

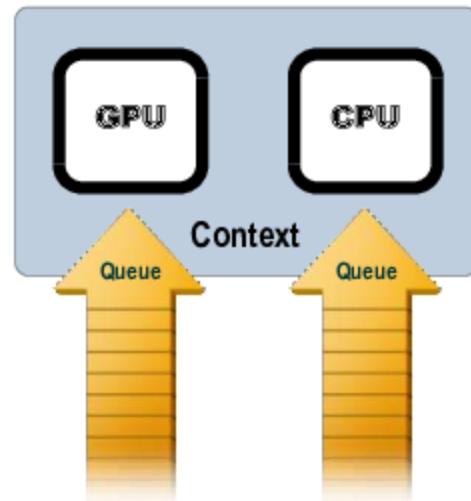


**Memory management is Explicit**

**You must move data from host -> global -> local ... and back**

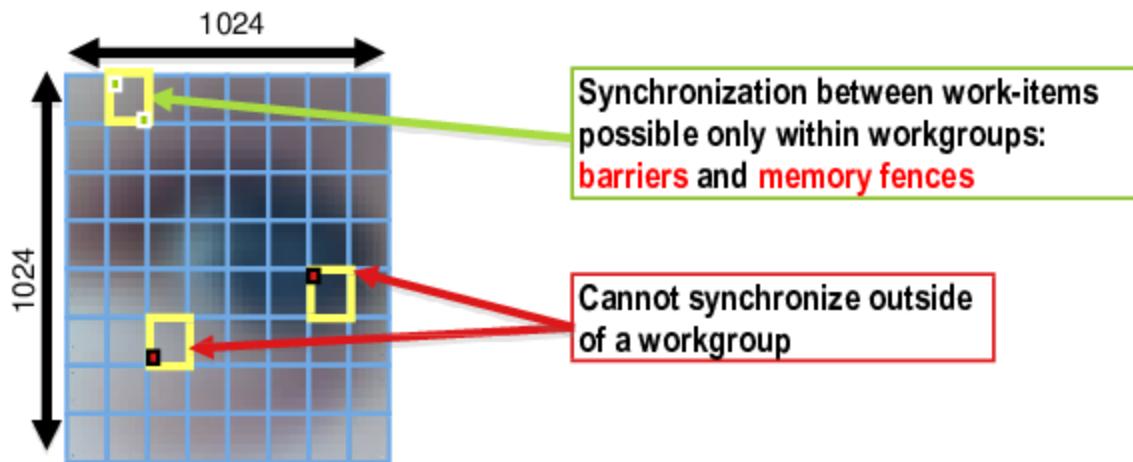
# OpenCL Execution Model

- OpenCL application runs on a host which submits work to the compute devices
  - **Context:** The environment within which work-items executes ... includes devices and their memories and command queues
  - **Program:** Collection of kernels and other functions (Analogous to a dynamic library)
  - **Kernel:** the code for a work item.  
Basically a C function
  - **Work item:** the basic unit of work on an OpenCL device
- Applications queue kernel execution
  - Executed in-order or out-of-order



# An N-dimension domain of work-items

- Kernels executed across a global domain of *work-items*
- Work-items grouped into local *workgroups*
- Define the “best” N-dimensioned index space for your algorithm
  - Global Dimensions: 1024 x 1024 (whole problem space)
  - Local Dimensions: 128 x 128 (work group ... executes together)



# Programming Kernels: OpenCL C

- **Derived from ISO C99**

- But without some C99 features such as standard C99 headers, function pointers, recursion, variable length arrays, and bit fields

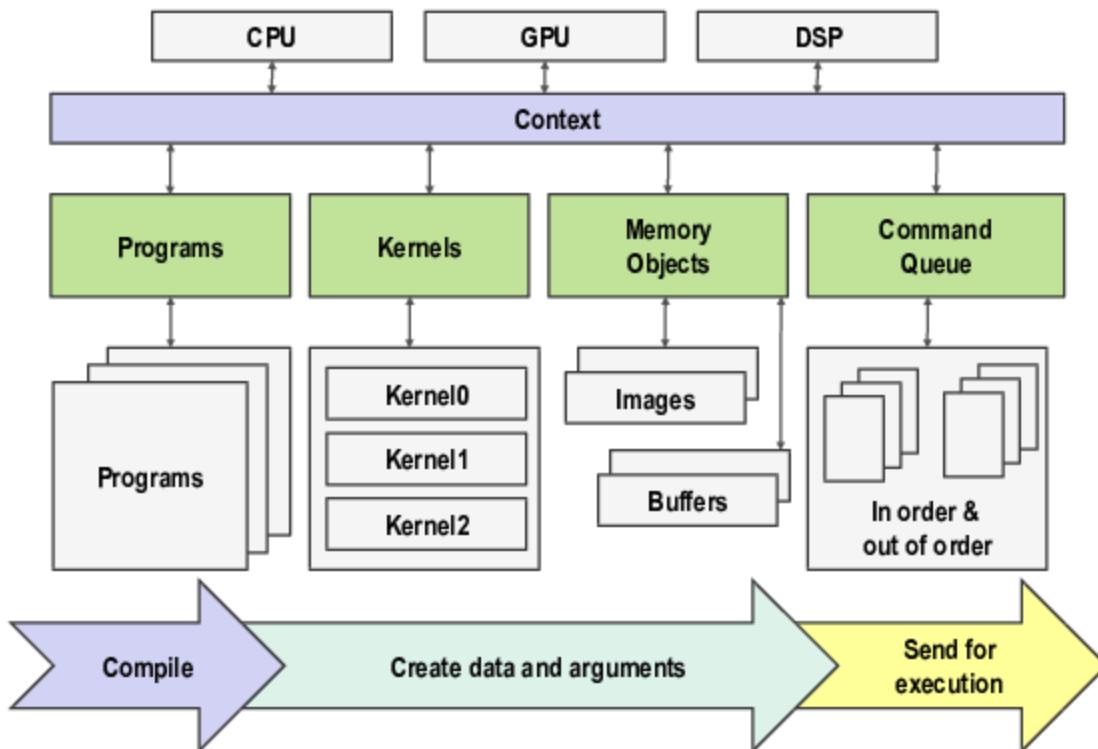
- **Language Features Added**

- Work-items and workgroups
  - Vector types
  - Synchronization
  - Address space qualifiers

- **Also includes a large set of built-in functions**

- Image manipulation
  - Work-item manipulation,
  - Math functions, etc.

# Creating an OpenCL Program

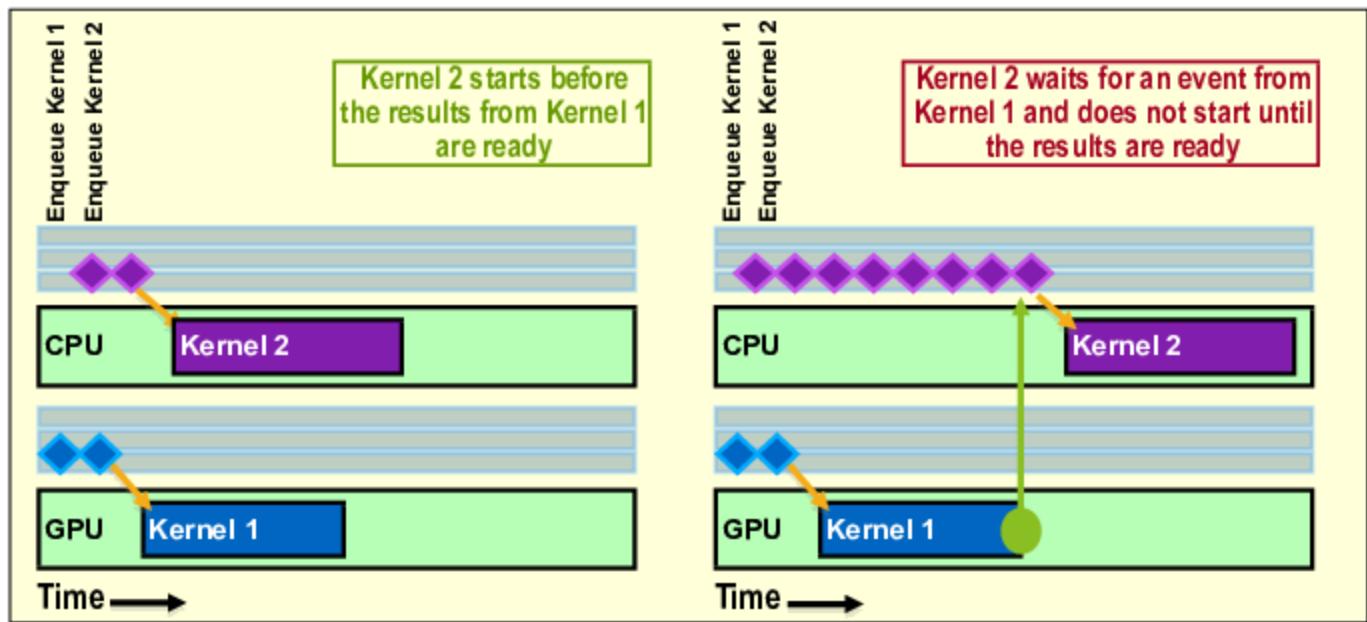


© Copyright Khronos Group, 2012 - Page 6

[https://www.khronos.org/assets/uploads/developers/library/2012-pan-pacific-road-show-June/OpenCL-Details-Taiwan\\_June-2012.pdf](https://www.khronos.org/assets/uploads/developers/library/2012-pan-pacific-road-show-June/OpenCL-Details-Taiwan_June-2012.pdf)

# Synchronization: Queues & Events

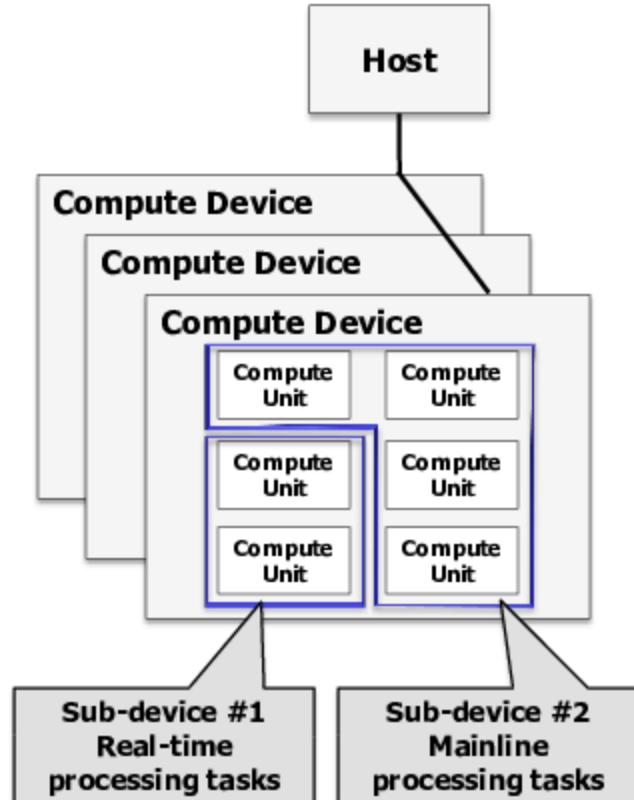
- Events can be used to synchronize kernel executions between queues
- Example: 2 queues with 2 devices



© Copyright Khronos Group, 2012 - Page 7

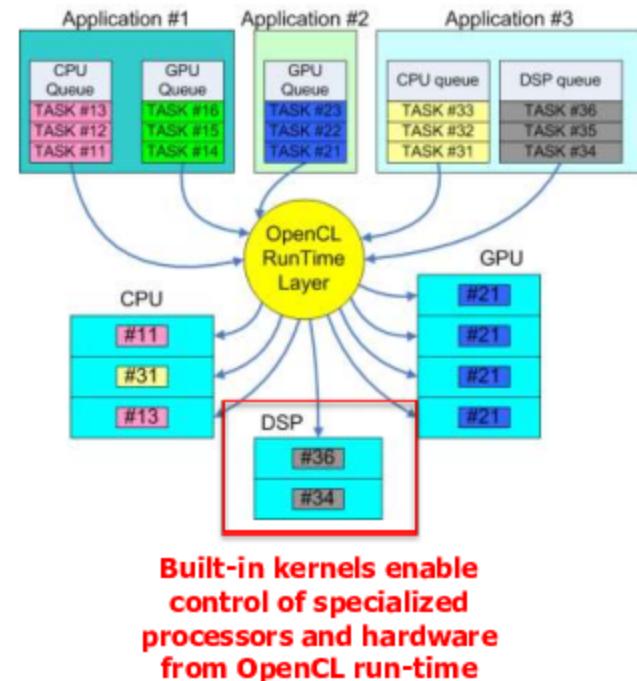
# Partitioning Devices

- Devices can be partitioned into sub-devices
  - More control over how computation is assigned to compute units
- Sub-devices may be used just like a normal device
  - Create contexts, building programs, further partitioning and creating command-queues
- Three ways to partition a device
  - Split into equal-size groups
  - Provide list of group sizes
  - Group devices sharing a part of a cache hierarchy



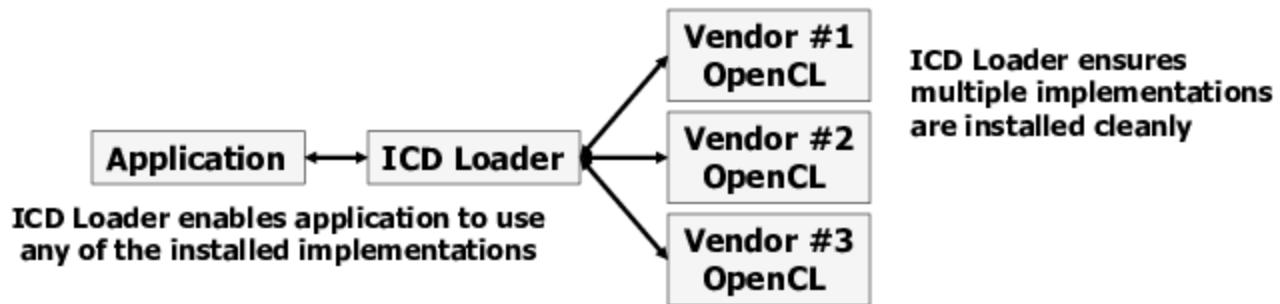
# Custom Devices and Built-in Kernels

- Embedded platforms often contain specialized hardware and firmware
  - That cannot support OpenCL C
- Built-in kernels can represent these hardware and firmware capabilities
  - Such as video encode/decode
- Hardware can be integrated and controlled from the OpenCL framework
  - Can enqueue built-in kernels to custom devices alongside OpenCL kernels
- FPGAs are one example of device that can expose built-in kernels
  - Latest FPGAs can support full OpenCL C as well
- OpenCL becomes a powerful coordinating framework for diverse resources
  - Programmable and non-programmable devices controlled by one run-time



# Installable Client Driver

- **Analogous to OpenGL ICDs in use for many years**
  - Used to handle multiple OpenGL implementations installed on a system
- **Optional extension**
  - Platform vendor will choose whether to use ICD mechanisms
- **Khronos OpenCL installable client driver loader**
  - Exposes multiple separate vendor installable client drivers (Vendor ICDs)
- **Application can access all vendor implementations**
  - The ICD Loader acts as a de-multiplexor



# OpenCL Desktop Implementations

- <http://developer.amd.com/zones/OpenCLZone/>
- <http://software.intel.com/en-us/articles/opencl-sdk/>
- <http://developer.nvidia.com/opencl>

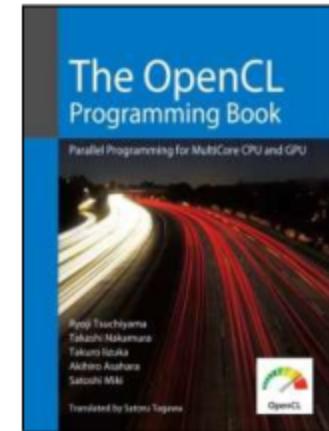
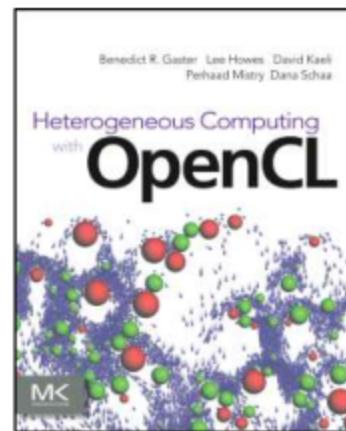
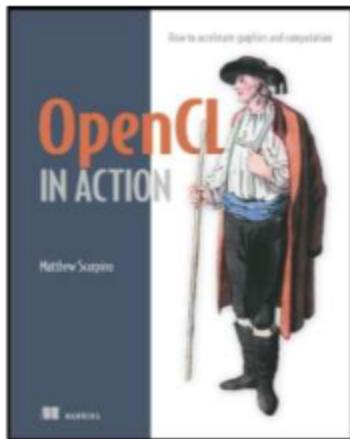
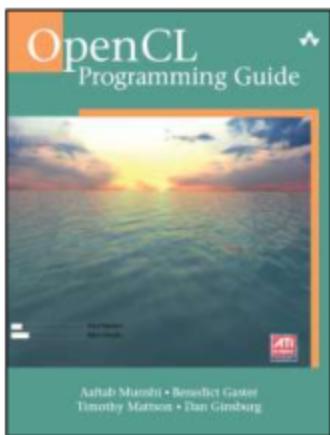


© Copyright Khronos Group, 2012 - Page 13

[https://www.khronos.org/assets/uploads/developers/library/2012-pan-pacific-road-show-June/OpenCL-Details-Taiwan\\_June-2012.pdf](https://www.khronos.org/assets/uploads/developers/library/2012-pan-pacific-road-show-June/OpenCL-Details-Taiwan_June-2012.pdf)

# OpenCL Books – Available Now!

- **OpenCL Programming Guide - The "Red Book" of OpenCL**
  - <http://www.amazon.com/OpenCL-Programming-Guide-Aaftab-Munshi/dp/0321749642>
- **OpenCL in Action**
  - <http://www.amazon.com/OpenCL-Action-Accelerate-Graphics-Computations/dp/1617290173/>
- **Heterogeneous Computing with OpenCL**
  - <http://www.amazon.com/Heterogeneous-Computing-with-OpenCL-ebook/dp/B005JRHYUS>
- **The OpenCL Programming Book**
  - <http://www.fixstars.com/en/opencl/book/>



© Copyright Khronos Group, 2012 - Page 14

[https://www.khronos.org/assets/uploads/developers/library/2012-pan-pacific-road-show-June/OpenCL-Details-Taiwan\\_June-2012.pdf](https://www.khronos.org/assets/uploads/developers/library/2012-pan-pacific-road-show-June/OpenCL-Details-Taiwan_June-2012.pdf)

# OpenCL: Definitions

**Devices:** Collection of OpenCL devices used by the host

**Kernels:** The OpenCL functions that run on OpenCL devices

**Program Objects:** The program source and executable that implements the kernel

**Memory Objects:** A set of memory objects visible to the host and the OpenCL devices. Memory objects contain values that can be operated on by instances of a kernel.

# OpenCL: Definitions...

**Global Memory:** This memory region permits read/write access to all work-items in all work-groups. Work-items can read from or write to any element of a memory object. Reads and writes to global memory may be cached depending on the capabilities of the device.

**Constant Memory:** A region of global memory that remains constant during the execution of a kernel. The host allocates and initializes memory objects placed into constant memory.

**Local Memory:** Local Memory: A memory region local to a work-group. This memory region can be used to allocate variables that are shared by all work-items in that work-group. It may be implemented as dedicated regions of memory on the OpenCL device. Alternatively, the local memory region may be mapped onto sections of the global memory.

**Private Memory:** A region of memory private to a work-item. Variables defined in one work-item's private memory are not visible to another work-item.

# Memory Consistency

OpenCL uses a relaxed consistency memory model; i.e. the state of memory visible to a workitem is not guaranteed to be consistent across the collection of work-items at all times.

Within a work-item memory has load / store consistency. Local memory is consistent across work-items in a single work-group at a work-group barrier. Global memory is consistent across work-items in a single work-group at a work-group barrier, but there are no guarantees of memory consistency between different work-groups executing a kernel.

Memory consistency for memory objects shared between enqueued commands is enforced at a synchronization point.

For more information see the very readable OpenCL Standard document by the Khronos Group.

# More remarks

- The current standard version is 2.2. However, implementations are not yet widely spread.
- The most widely used standard is 1.2 (NVIDIA only supports 1.2). Intel and AMD also support OpenCL 2.
- In this module we focus on OpenCL 1.2.
- Cuda is very dominant in the HPC market. In this module we use OpenCL since it runs on a variety of devices and its abstract device model makes a transition to Cuda easy (use PyCuda for Cuda from Python).
- In the future OpenCL will merge with the Vulkan API for fast computer graphics (already used by many games). This will provide a unified graphics and compute API.

# Distributed Memory Parallelization

- Thread based parallelization only works for shared memory
- All threads run within the same process on the same physical node
- Different approach needed for communication between processes on the same and different nodes
- Standard framework is MPI (Message Passing Interface)

# A brief overview of MPI

- MPI-1 standard published in 1994
- MPI-2 came in 1996. Added parallel I/O, one-sided communication, dynamic process management
- Most MPI programs only require a handful of MPI commands
- MPI is mainly a distributed memory concept. Some shared memory features introduced in MPI-3
- MPI is tremendously successful. Virtually every distributed HPC code uses it for communication
- It scales well from desktops to peta-scale architectures
- The main used implementations are MPICH, OpenMPI and Intel MPI
- MPI is directly callable from Fortran, C/C++. But has bindings for other languages (including Python)

# MPI Hello World

```
1 #include <mpi.h>
2 #include <stdio.h>
3
4 int main() {
5     // Initialize the MPI environment
6     MPI_Init(NULL, NULL);
7
8     // Get the number of processes
9     int world_size;
10    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
11
12    // Get the rank of the process
13    int world_rank;
14    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
15
16    // Get the name of the processor
17    char processor_name[MPI_MAX_PROCESSOR_NAME];
18    int name_len;
19    MPI_Get_processor_name(processor_name, &name_len);
20
21    // Print off a hello world message
22    printf("Hello world from processor %s, rank %d"
23          " out of %d processors\n",
24          processor_name, world_rank, world_size);
25
26    // Finalize the MPI environment.
27    MPI_Finalize();
28 }
|
```

- Compile with  
mpic++ ./hello.cpp
- Run with
- mpirun -n 4 ./a.out

Hello world from processor Timos-MBP-2,  
rank 0 out of 4 processors

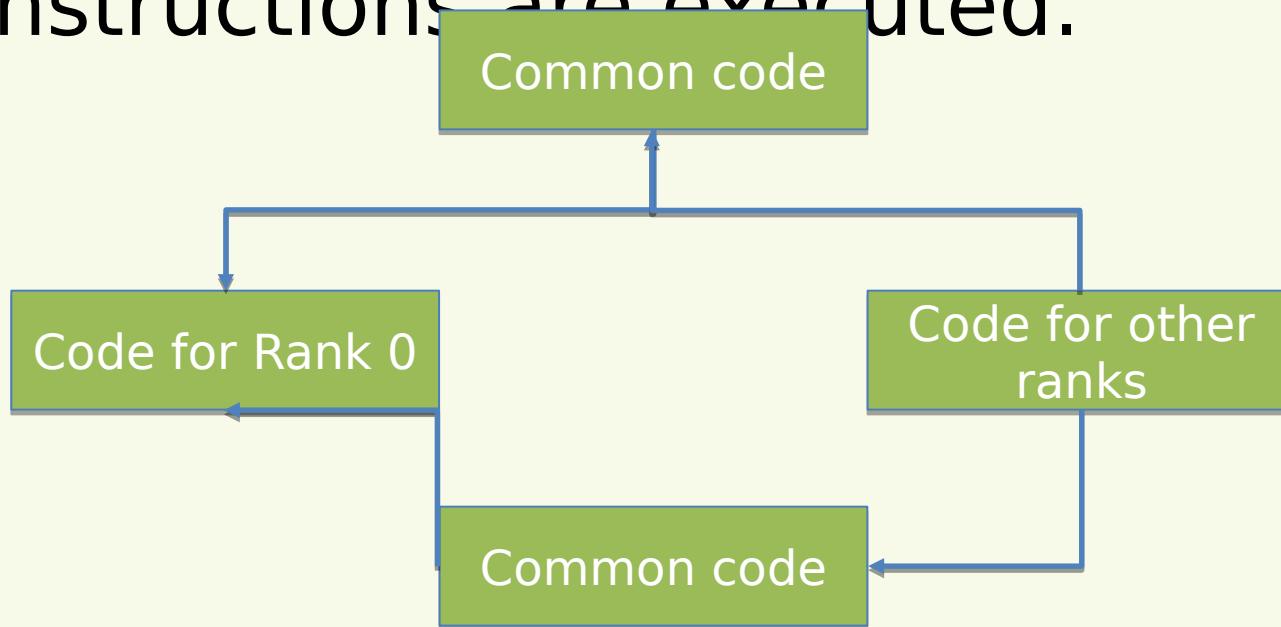
Hello world from processor Timos-MBP-2,  
rank 1 out of 4 processors

Hello world from processor Timos-MBP-2,  
rank 2 out of 4 processors

Hello world from processor Timos-MBP-2,  
rank 3 out of 4 processors

# The MPI Model

- Typically one program for all MPI processes
- Depending on the MPI rank different instructions are executed.



# Sending and receiving data

```
1 #include <iostream>
2 #include <mpi.h>
3
4 int main() {
5
6     int my_number;
7     int other_number;
8
9     MPI_Init(NULL, NULL);
10
11    // Get the number of processes
12    int world_size;
13    MPI_Comm_size(MPI_COMM_WORLD, &world_size);
14
15    // Get the rank of the process
16    int world_rank;
17    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
18
19    int error;
20
21    if (world_rank == 0) {
22        my_number = 42;
23        error = MPI_Send(&my_number, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
24    }
25
26    if (world_rank == 1) {
27        MPI_Status status;
28        error = MPI_Recv(&other_number, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
29        std::cout << "I have received from rank 0 the number "
30              << other_number
31              << "." << std::endl;
32    }
33 }
```

- Run with  
mpirun -n 2 ./send\_receive



More on MPI soon via the mpi4py interface

# Introducing mpi4py

```
1  #!/usr/bin/env python
2  """
3  Parallel Hello World
4  """
5
6  from mpi4py import MPI
7  import sys
8
9  ● size = MPI.COMM_WORLD.Get_size()
10  ● rank = MPI.COMM_WORLD.Get_rank()
11  ● name = MPI.Get_processor_name()
12
13  sys.stdout.write(
14      "Hello, World! I am process %d of %d on %s.\n"
15      % (rank, size, name))
16
```

mpirun -n 2 python hello.py

# Sending data

```
1  from mpi4py import MPI
2  import numpy as np
3  import sys
4
5  comm = MPI.COMM_WORLD
6
7  size = comm.Get_size()
8  rank = comm.Get_rank()
9
10 sendbuf = np.empty(1, dtype='i')
11 recvbuf = np.empty(1, dtype='i')
12
13 sendbuf[:] = rank
14
15 if rank == 0:
16     comm.Send(sendbuf, rank + 1)
17 elif rank > 0 and rank < size-1:
18     comm.Recv(recvbuf, rank - 1)
19     comm.Send(sendbuf, rank + 1)
20 else:
21     comm.Recv(recvbuf, rank - 1)
22
23 if rank > 0:
24     sys.stdout.write(
25         "Process %d received data %d.\n"
26         % (rank, recvbuf[0]))
27 |
```

- Each rank sends data to its successor
- Need to make sure that send and receive operations do not block each other
- Better alternative: Use non-blocking send/receive with Isend/Irecv

# Non-blocking send

```
1  from mpi4py import MPI
2  import numpy as np
3  import sys
4
5  comm = MPI.COMM_WORLD
6
7  size = comm.Get_size()
8  rank = comm.Get_rank()
9
10 sendbuf = np.empty(1, dtype='i')
11 recvbuf = np.empty(1, dtype='i')
12
13 sendbuf[:] = rank
14
15 if rank < size - 1:
16     comm.Isend(sendbuf, rank + 1)
17
18 if rank > 0:
19     comm.Recv(recvbuf, rank - 1)
20
21 comm.Barrier()
22
23 if rank > 0:
24     sys.stdout.write(
25         "Process %d received data %d.\n"
26         % (rank, recvbuf[0]))
27
```

- Each rank sends data to its successor
- Send is non-blocking. The process does not wait for the actual send to take place.
- We enforce all process to finish the send/receive operations with a barrier

# A note on types

- Communication commands in mpi4py can start with a large letter (e.g. Send) or a small letter (send)
- If we use a large letter native data types are used. Compatible Numpy types are automatically recognized
- With small letter commands can send any objects, but execution slow due to necessary serialization of object.

# Scatter data

```
1     """Scatter data across all processes."""
2
3     from mpi4py import MPI
4     import numpy as np
5
6     comm = MPI.COMM_WORLD
7     size = comm.Get_size()
8     rank = comm.Get_rank()
9
10    sendbuf = None
11    if rank == 0:
12        sendbuf = np.empty([size, 10], dtype='i')
13        sendbuf.T[:, :] = range(size)
14    recvbuf = np.empty(10, dtype='i')
15    comm.Scatter(sendbuf, recvbuf, root=0)
16    assert np.allclose(recvbuf, rank)
17
```

```
>>> sendbuf
array([[0, 0, 0, 0, 0],
       [1, 1, 1, 1, 1],
       [2, 2, 2, 2, 2]],
      dtype=int32)
```

On rank 1 have

```
>>> recvbuf
array([1, 1, 1, 1, 1],
      dtype=int32)
```

- Given an input array sendbuf. Distribute the array elements onto all processes.
- For unequal chunk sizes uses Scatterv
- Note that memory buffers need already exist!

# Reduce data

- Sum data across all processes.
- Every process gets the sum of the data.
- Different types of reducers possible.

```
1 """Scatter data across all processes."""
2
3 from mpi4py import MPI
4 import numpy as np
5 import sys
6
7 comm = MPI.COMM_WORLD
8 size = comm.Get_size()
9 rank = comm.Get_rank()
10
11 sendbuf = np.empty(1, dtype='i')
12 sendbuf[:] = 1
13
14 recvbuf = np.empty(1, dtype='i')
15
16 comm.Allreduce(sendbuf, recvbuf, MPI.SUM)
17
18 sys.stdout.write(
19     "Process %d has data %d.\n"
20     % (rank, recvbuf[0]))
```

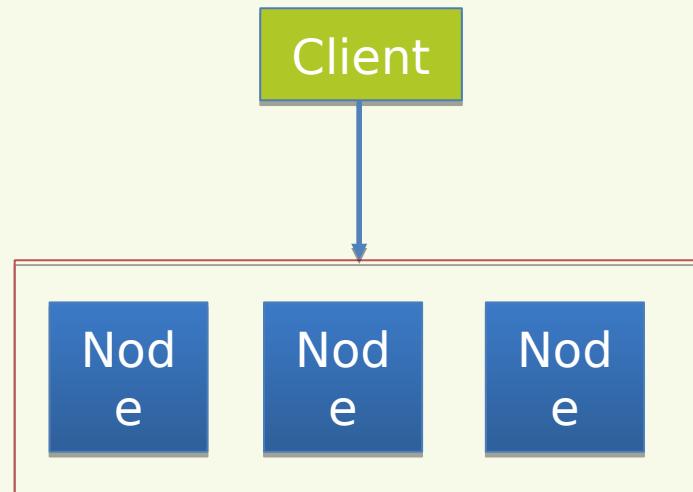
# More remarks on MPI

- Have given a small sample of MPI operations
- A range of sophisticated MPI communication and process management operations available
- Many are supported by mpi4py
- In practice we often only need a few (send/receive and global communication)

# IPython clusters

- Supercomputing on your fingertips using Ipyparallel
- Have a local Python client that communicates with an HPC resource
- More info on

<https://ipyparallel.readthedocs.io/en/latest/>



# Modeling a large-scale PDE

## Diffusion in a complex medium

$$-\nabla \cdot \sigma(\mathbf{x}) \nabla u(\mathbf{x}) = f(\mathbf{x}), \text{ in } \Omega$$

$u(\mathbf{x}) = 0, \text{ on } \partial\Omega$

Diffusion coefficients in material

Density of material

Source term

- Diffusion of a pollutant in soil
- Often diffusion coefficient a random parameter

Ellipticity conditon:

$$\langle \sigma(\mathbf{x})\mathbf{y}, \mathbf{y} \rangle \geq C\|\mathbf{y}\|_2^2, \quad \forall \mathbf{y} \in \mathbf{R}^3, \quad C > 0$$

# Deriving the weak form

Multiply with test  
function in  $H_0^1(\Omega)$

$$-\nabla \cdot \sigma(\mathbf{x}) \nabla u(\mathbf{x}) v(\mathbf{x}) = f(\mathbf{x}) v(\mathbf{x})$$

Integration by parts  
gives

$$\int_{\Omega} [\sigma(\mathbf{x}) \nabla u(\mathbf{x})] \cdot \nabla v(\mathbf{x}) d\mathbf{x} - \int_{\partial\Omega} \sigma(\mathbf{x}) \frac{\partial u(\mathbf{x})}{\partial n(\mathbf{x})} v(\mathbf{x}) ds(\mathbf{x}) = \int_{\Omega} f(\mathbf{x}) v(\mathbf{x})$$

Apply the boundary condition to  
arrive at

$$\int_{\Omega} [\sigma(\mathbf{x}) \nabla u(\mathbf{x})] \cdot \nabla v(\mathbf{x}) d\mathbf{x} = \int_{\Omega} f(\mathbf{x}) v(\mathbf{x})$$

Weak form of the partial differential equation

# Deriving the weak form...

Find  $u \in H_0^1(\Omega)$  such that

$$\int_{\Omega} [\sigma(\mathbf{x}) \nabla u(\mathbf{x})] \cdot \nabla v(\mathbf{x}) d\mathbf{x} = \int_{\Omega} f(\mathbf{x}) v(\mathbf{x})$$

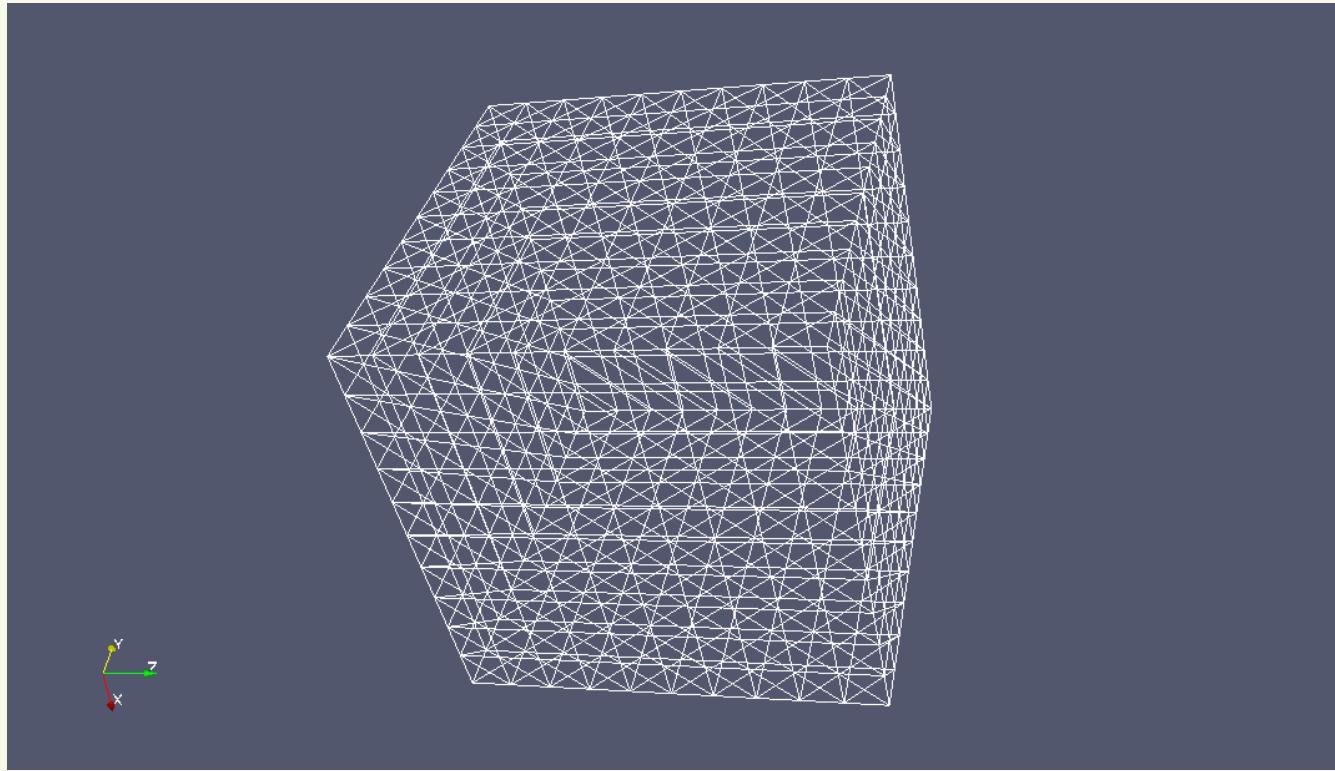
for all  $v \in H_0^1(\Omega)$ .

Weak form of the partial differential equation

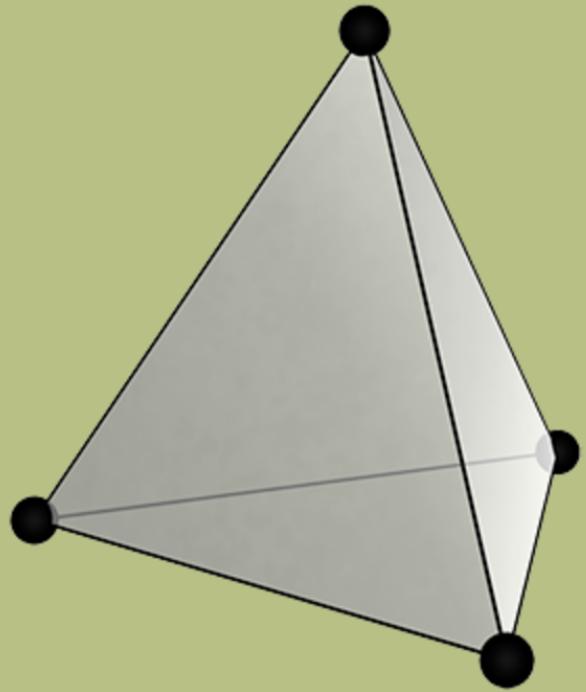
- Create a computational grid of the domain
- Define a discrete function space over the domain
- Assemble the equation on the discrete function space
- Solve the resulting linear system of equations

# Creating a computational grid

```
from dolfin import *
Mesh = UnitCubeMesh(10,
10, 10)
```



# Standard P1 Elements



A single P1 basis function has the value 1 at one node and the value zero at all other nodes.

The set of all P1 basis functions over a grid span a space of piecewise linear, continuous basis functions.

Define a P1 function space in Fenics:

```
V = FunctionSpace(mesh, "Lagrange"
```

# The discrete weak form

- Denote by  $H_{0,h}^1(\Omega)$  the finite dimensional space spanned by P1 basis functions over a given grid with element diameter  $h$ .
- Discretize the variational problem by restricting it into this finite dimensional space.

Find  $u \in H_{0,h}^1(\Omega)$  such that

$$\int_{\Omega} [\sigma(\mathbf{x}) \nabla u(\mathbf{x})] \cdot \nabla v(\mathbf{x}) d\mathbf{x} = \int_{\Omega} f(\mathbf{x}) v(\mathbf{x})$$

for all  $v \in H_{0,h}^1(\Omega)$ .

Define

$$a(u, v) := \int_{\Omega} [\sigma(\mathbf{x}) \nabla u(\mathbf{x})] \cdot \nabla v(\mathbf{x}) d\mathbf{x}$$

# Obtaining a matrix

Let  $\psi_j$  be the P1 basis function associated with the jth node in the grid.

We can therefore write

$$H_{0,h}^1(\Omega) := \text{span}\{\psi_1, \dots, \psi_N\}$$

Any function  $\psi$  in  $H_{0,h}^1(\Omega)$  can be written as

$$\psi = \sum_{j=1}^N x_j \Psi_j$$

for some coefficient vector  $x$ .

Define  $A \in \mathbb{R}^{N \times N}$  by  $A_{i,j} = a(\psi_j, \psi_i)$  and  $\bar{f}_j = \int_{\Omega} f(\mathbf{x}) \psi_j(\mathbf{x}) d\mathbf{x}$   
The variational problem is equivalent to

Find  $x \in \mathbb{R}^N$  such that

$$y^T A x = y^T \bar{f}$$

for all  $y \in \mathbb{R}^N$ .

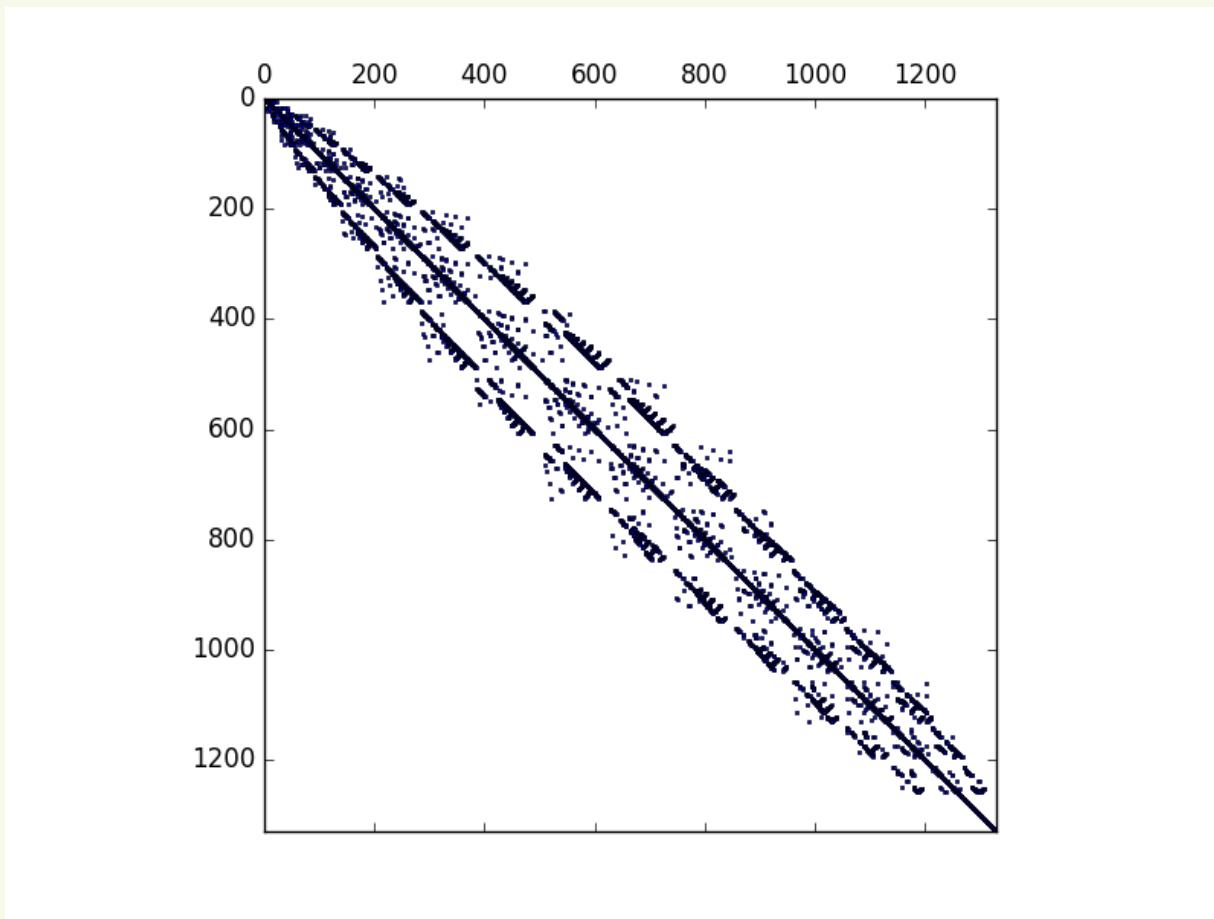
# Obtaining a matrix...

This is equivalent to the linear system of equations

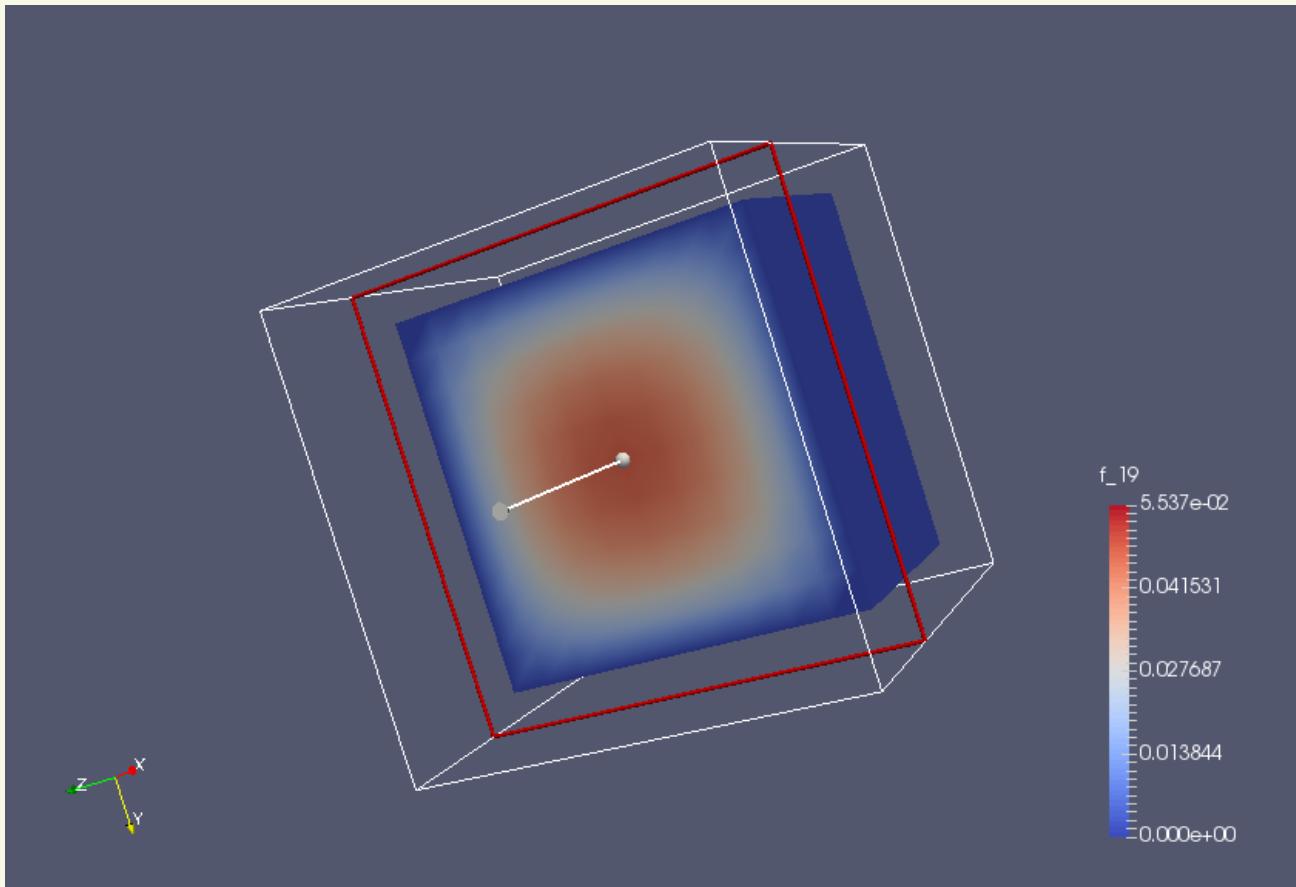
$$Ax = \bar{f}$$

- The system matrix is sparse, that is only a small percentage of the elements is nonzero.
- For smaller problem sizes a direct solution is possible. For larger problems iterative solvers need to be used.
- Very large-scale practical applications can have billions of unknowns.

# Obtaining a matrix...



# The complete solution



# An example FeniCS Script

```
1 from dolfin import *
2
3 mesh = UnitCubeMesh(10, 10, 10)
4
5 V = FunctionSpace(mesh, "Lagrange", 1)
6 def boundary(x):
7     return (x[0] < DOLFIN_EPS or x[0] > 1 - DOLFIN_EPS
8         or x[1] < DOLFIN_EPS or x[1] > 1 - DOLFIN_EPS
9         or x[2] < DOLFIN_EPS or x[2] > 1 - DOLFIN_EPS)
10
11 u0 = Constant(0)
12
13 bc = DirichletBC(V, u0, boundary)
14
15 u = TrialFunction(V)
16 v = TestFunction(V)
17 f = Expression('1', degree=1)
18
19 a = inner(grad(u), grad(v)) * dx
20 rhs = f * v * dx
21
22 A = assemble(a)
23 b = assemble(rhs)
24 bc.apply(A, b)
25 sol = Vector()
26 solve(A, sol, b, 'lu')
27 output_file = File('poisson.pvd')
28 output_file << Function(V, sol)
29
30
```

Define the mesh

Specify boundary region (for complex domains other ways are possible)

Definition of the weak form

Assemble the operators and apply boundary condition

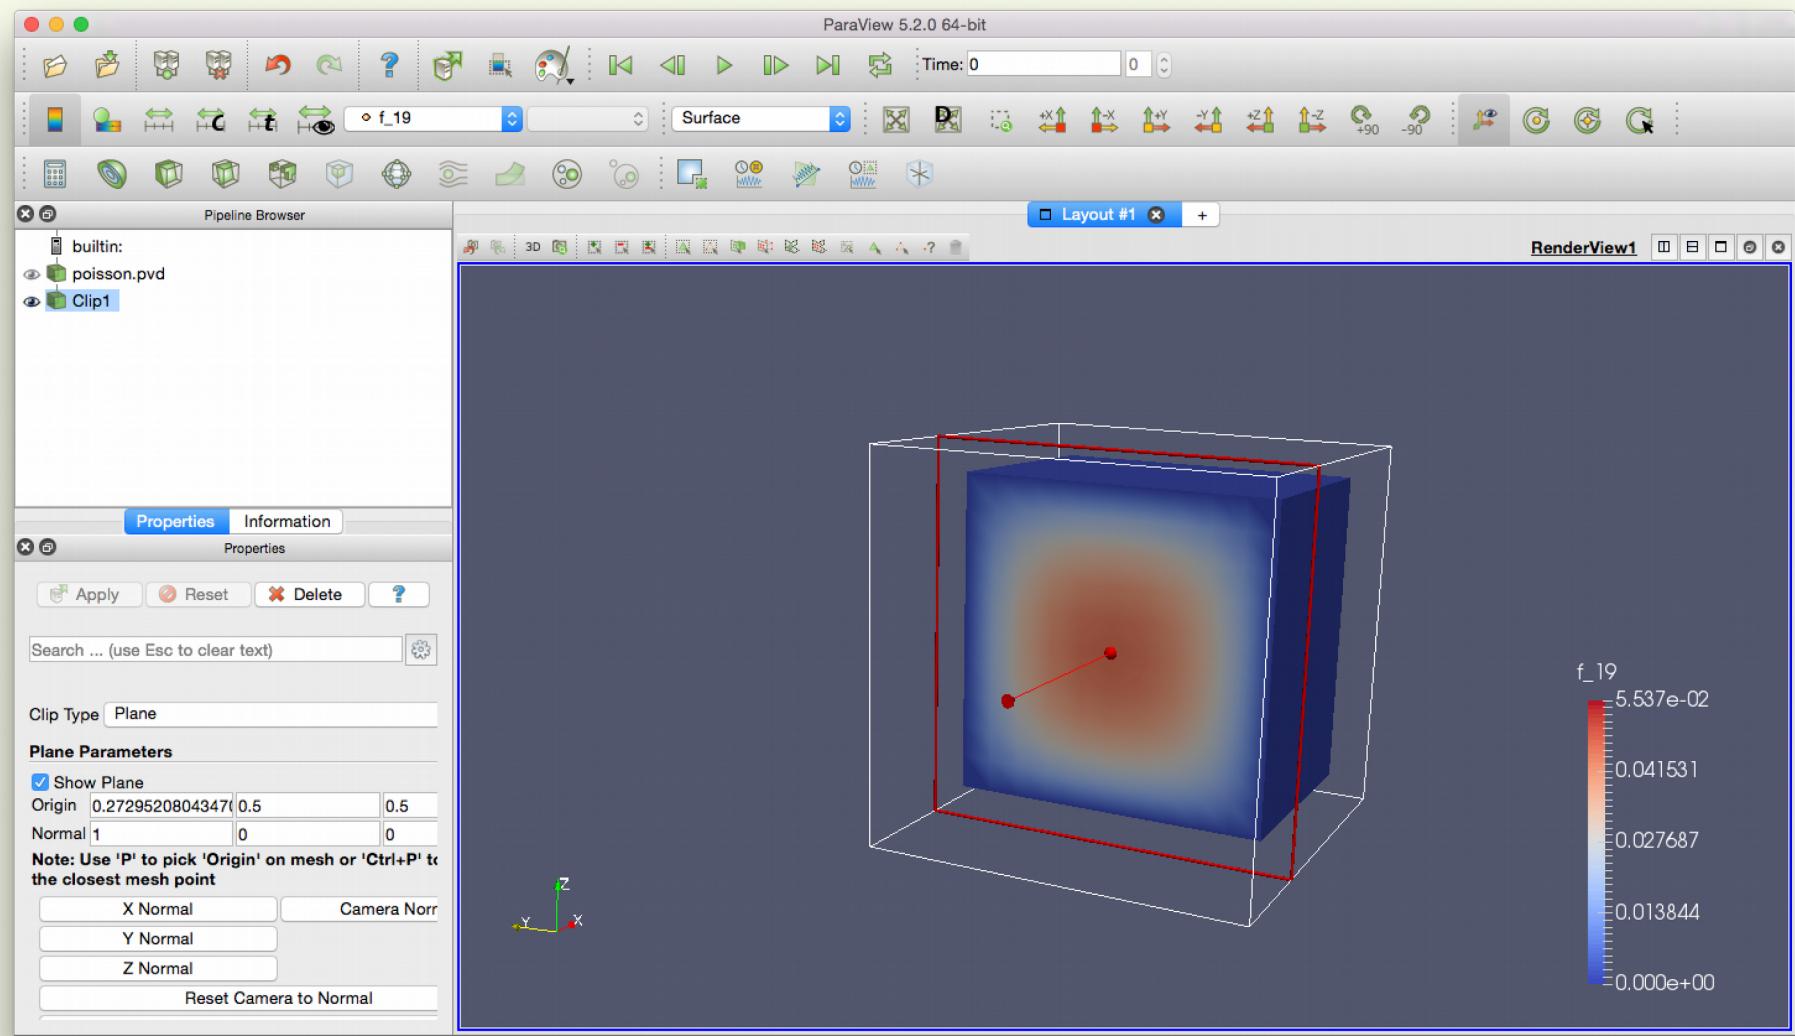
Solve

Write out the solution

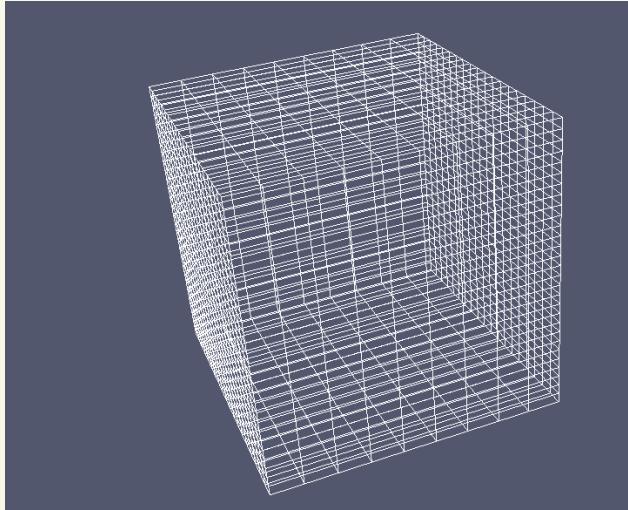
# Visualising with VTK

- VTK is widely used data format for storing meshes and data over meshes
- Parallel operations possible
- Developed by Kitware
- Bindings available for all major programming languages (including Python)
- Example viewers: Paraview, VisIt, Mayavi2, etc.

# Paraview

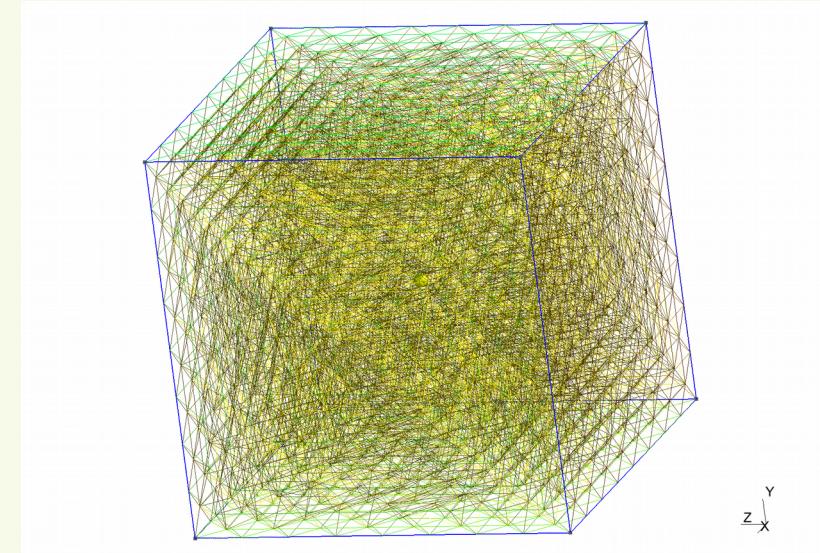


# Computational Grids



- Each point uniquely identified by index triplet (i, j, k)
- Fast algorithms due to regular structure (in particular good for GPUs)
- Used in many finite difference applications
- Only simple geometries possible

- Need to store points and connectivity explicitly
- Perfect for generic unstructured geometries
- Allows refinement towards complex geometric features



# Unstructured grid data structures

- Typically two sets of data:
  - A point structure that stores the nodal information
  - An element structure that stores the connectivity between points

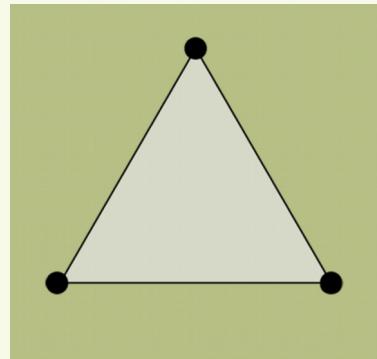
```
66    4 13 12 10 8
67    4 10 12 11 8
68    4 9 10 11 8
69    4 13 10 9 8
70    4 5 13 6 9
71    4 13 7 6 10
72    4 7 4 12 13
73    4 5 9 1 8
74    4 0 1 11 8
75    4 5 4 13 8
76    4 3 7 12 10
77    4 6 10 2 9
78    4 4 0 12 8
79    4 10 3 2 11
```

```
1 # vtk DataFile Version 2.0
2 untitled, Created by Gmsh
3 ASCII
4 DATASET UNSTRUCTURED_GRID
5 POINTS 14 double
6 0 0 0
7 1 0 0
8 1 1 0
9 0 1 0
10 0 0 1
11 1 0 1
12 1 1 1
13 0 1 1
14 0.5 0 0.5
15 1 0.5 0.5
16 0.5 1 0.5
17 0.5 0.5 0
18 0 0.5 0.5
19 0.5 0.5 1
20
21 CELLS 68 268
22 1 0
23 1 1
24 1 2
25 1 3
26 1 4
27 1 5
28 1 6
29 1 7
30 2 0 4
31 2 4 5
32 2 5 1
33 2 1 0
```

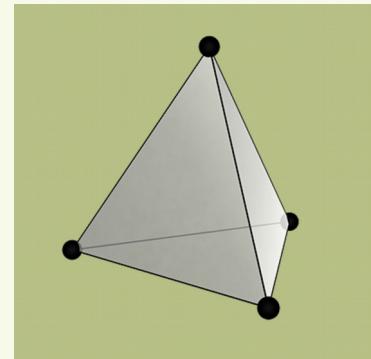
# Frequent element types



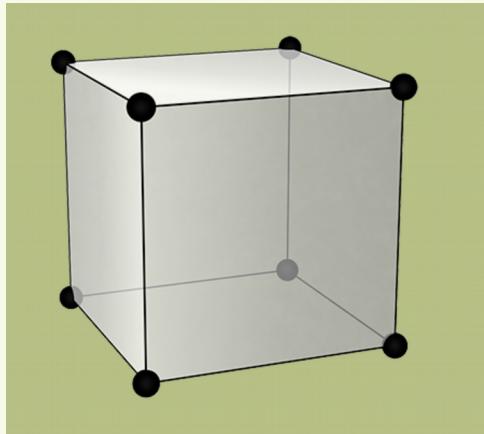
Line



Triangle



Tetrahedron



Cub  
e

- These are lowest order elements
- Higher order variants available
- Many more element types exist
- See [www.femtable.org](http://www.femtable.org)

# Periodic Table of the Finite Elements

$k=0$	$k=1$	$k=2$	$k=3$	$k=0$	$k=1$	$k=2$	$k=3$	$k=0$	$k=1$	$k=2$	$k=3$																																																						
<b><math>n=1</math></b>	$r=1$	<b><math>P_1</math></b>  $\mathcal{P}_1 \Lambda^k$ The shape function space for $\mathcal{P}_1 \Lambda^k$ is $\mathcal{P}_1 \Lambda^k = \mathcal{P}_1 \Lambda^0 \times \mathcal{P}_1 \Lambda^{k-1}$ , where $\mathcal{P}_1 \Lambda^0$ is the Raviart-Thomas 1 finite element. It includes all the full polynomial space $\mathcal{P}_1$ included in $\mathcal{P}_1 \Lambda^k$ , and has dimension $\dim \mathcal{P}_1 \Lambda^k(\Delta) = \binom{r+k}{r} + \binom{r+k-1}{r}$ . The degrees of freedom are given on faces $f$ of dimension $d \geq k$ by moments of the trace weighted by a full polynomial space: $u \mapsto \int_{\partial f} (tr u) \wedge q, \quad q \in \mathcal{P}_{r-k-1} \Lambda^{k-1}(f).$ The spaces with constant degree $r$ form a complex: $\mathcal{P}_1 \Lambda^0 \xrightarrow{\Delta} \mathcal{P}_1 \Lambda^1 \xrightarrow{\Delta} \dots \xrightarrow{\Delta} \mathcal{P}_1 \Lambda^k.$	<b><math>dP_1</math></b>  $\mathcal{P}_1 \Lambda^k$	<b><math>P_2</math></b>  $\mathcal{P}_2 \Lambda^k$ The shape function space for $\mathcal{P}_2 \Lambda^k$ consists of all differential $k$ -forms with polynomial degrees of at most $r$ . The space has dimension $\dim \mathcal{P}_2 \Lambda^k(\Delta) = (r+k) \binom{r+k}{k}$ . The degrees of freedom are given on faces $f$ of dimension $d \geq k$ by moments of the trace weighted by a $\mathcal{P}_1$ space: $u \mapsto \int_f (tr u) \wedge q, \quad q \in \mathcal{P}_{r-k-1} \Lambda^{k-1}(f).$ The spaces with decreasing degree $r$ form a complex: $\mathcal{P}_2 \Lambda^0 \xrightarrow{\Delta} \mathcal{P}_2 \Lambda^1 \xrightarrow{\Delta} \dots \xrightarrow{\Delta} \mathcal{P}_2 \Lambda^k.$	<b><math>P_1</math></b>  $\mathcal{P}_1 \Lambda^k$ The shape function space for $\mathcal{P}_1 \Lambda^k$ consists of all differential $k$ -forms with polynomial degrees of at most $r$ . The space has dimension $\dim \mathcal{P}_1 \Lambda^k(\Delta) = (r+k) \binom{r+k}{k}$ . The degrees of freedom are given on faces $f$ of dimension $d \geq k$ by moments of the trace weighted by a $\mathcal{P}_1$ space: $u \mapsto \int_f (tr u) \wedge q, \quad q \in \mathcal{P}_{r-k-1} \Lambda^{k-1}(f).$ The spaces with decreasing degree $r$ form a complex: $\mathcal{P}_1 \Lambda^0 \xrightarrow{\Delta} \mathcal{P}_1 \Lambda^1 \xrightarrow{\Delta} \dots \xrightarrow{\Delta} \mathcal{P}_1 \Lambda^k.$	<b><math>dP_2</math></b>  $\mathcal{P}_1 \Lambda^k$	<b><math>P_3</math></b>  $\mathcal{P}_3 \Lambda^k$ The shape function space for $\mathcal{P}_3 \Lambda^k$ consists of all differential $k$ -forms with polynomial degrees of at most $r$ . The space has dimension $\dim \mathcal{P}_3 \Lambda^k(\Delta) = (r+k) \binom{r+k}{k}$ . The degrees of freedom are given on faces $f$ of dimension $d \geq k$ by moments of the trace weighted by a $\mathcal{P}_2$ space: $u \mapsto \int_f (tr u) \wedge q, \quad q \in \mathcal{P}_{r-k-1} \Lambda^{k-1}(f).$ The spaces with decreasing degree $r$ form a complex: $\mathcal{P}_3 \Lambda^0 \xrightarrow{\Delta} \mathcal{P}_3 \Lambda^1 \xrightarrow{\Delta} \dots \xrightarrow{\Delta} \mathcal{P}_3 \Lambda^k.$	<b><math>dP_3</math></b>  $\mathcal{P}_2 \Lambda^k$	<b><math>S_1</math></b>  $\mathcal{S}_1 \Lambda^k$ The shape function space for $\mathcal{S}_1 \Lambda^k$ is given by the complex of 1-dimensional finite elements under a tensor product construction. The $i$ -th tensor product space on the unit cube $\square_1$ is given by $\bigoplus_{I \in \Sigma^k} \mathcal{P}_{n+1}(\square_1)^{ I },$ where $\Sigma^k$ denotes the increasing maps $I: \{1, \dots, n\} \rightarrow \{1, \dots, n\}$ . Its dimension is $\dim \mathcal{S}_1 \Lambda^k(\square_1) = \binom{n+1}{k} (n+1)^{k+1}$ . The degrees of freedom are $u \mapsto \int_{\square_1} (tr u) \wedge q, \quad q \in \mathcal{Q}_{n+1} \Lambda^{k-1}(\square_1).$ The spaces with constant degree $r$ form a complex: $\mathcal{S}_1 \Lambda^0 \xrightarrow{\Delta} \mathcal{S}_1 \Lambda^1 \xrightarrow{\Delta} \dots \xrightarrow{\Delta} \mathcal{S}_1 \Lambda^k.$	<b><math>dQ_1</math></b>  $\mathcal{S}_1 \Lambda^k$	<b><math>S_2</math></b>  $\mathcal{S}_2 \Lambda^k$ The shape function space for $\mathcal{S}_2 \Lambda^k$ is given by the complex of 1-dimensional finite elements under a tensor product construction. The $i$ -th tensor product space on the unit cube $\square_1$ is given by $\bigoplus_{I \in \Sigma^k} \mathcal{P}_{n+1}(\square_1)^{ I },$ where $\Sigma^k$ denotes the increasing maps $I: \{1, \dots, n\} \rightarrow \{1, \dots, n\}$ . Its dimension is $\dim \mathcal{S}_2 \Lambda^k(\square_1) = \binom{n+1}{k} (n+1)^{k+1}$ . The degrees of freedom are $u \mapsto \int_{\square_1} (tr u) \wedge q, \quad q \in \mathcal{Q}_{n+1} \Lambda^{k-1}(\square_1).$ The spaces with decreasing degree $r$ form a complex: $\mathcal{S}_2 \Lambda^0 \xrightarrow{\Delta} \mathcal{S}_2 \Lambda^1 \xrightarrow{\Delta} \dots \xrightarrow{\Delta} \mathcal{S}_2 \Lambda^k.$	<b><math>dQ_2</math></b>  $\mathcal{S}_2 \Lambda^k$	<b><math>S_3</math></b>  $\mathcal{S}_3 \Lambda^k$ The shape function space for $\mathcal{S}_3 \Lambda^k$ is given by the complex of 1-dimensional finite elements under a tensor product construction. The $i$ -th tensor product space on the unit cube $\square_1$ is given by $\bigoplus_{I \in \Sigma^k} \mathcal{P}_{n+1}(\square_1)^{ I },$ where $\Sigma^k$ denotes the increasing maps $I: \{1, \dots, n\} \rightarrow \{1, \dots, n\}$ . Its dimension is $\dim \mathcal{S}_3 \Lambda^k(\square_1) = \binom{n+1}{k} (n+1)^{k+1}$ . The degrees of freedom are $u \mapsto \int_{\square_1} (tr u) \wedge q, \quad q \in \mathcal{Q}_{n+1} \Lambda^{k-1}(\square_1).$ The spaces with decreasing degree $r$ form a complex: $\mathcal{S}_3 \Lambda^0 \xrightarrow{\Delta} \mathcal{S}_3 \Lambda^1 \xrightarrow{\Delta} \dots \xrightarrow{\Delta} \mathcal{S}_3 \Lambda^k.$	<b><math>dQ_3</math></b>  $\mathcal{S}_3 \Lambda^k$																																																			
<b><math>n=2</math></b>	$r=1$	<b><math>P_1</math></b>  $\mathcal{P}_1 \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dP_1</math></b>  $\mathcal{P}_1 \Lambda^k$	<b><math>P_2</math></b>  $\mathcal{P}_2 \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dP_2</math></b>  $\mathcal{P}_2 \Lambda^k$	<b><math>P_3</math></b>  $\mathcal{P}_3 \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dP_3</math></b>  $\mathcal{P}_3 \Lambda^k$	<b><math>P_4</math></b>  $\mathcal{P}_4 \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dP_4</math></b>  $\mathcal{P}_4 \Lambda^k$	<b><math>P_5</math></b>  $\mathcal{P}_5 \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dP_5</math></b>  $\mathcal{P}_5 \Lambda^k$	<b><math>P_6</math></b>  $\mathcal{P}_6 \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dP_6</math></b>  $\mathcal{P}_6 \Lambda^k$	<b><math>P_7</math></b>  $\mathcal{P}_7 \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dP_7</math></b>  $\mathcal{P}_7 \Lambda^k$	<b><math>P_8</math></b>  $\mathcal{P}_8 \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dP_8</math></b>  $\mathcal{P}_8 \Lambda^k$	<b><math>P_9</math></b>  $\mathcal{P}_9 \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dP_9</math></b>  $\mathcal{P}_9 \Lambda^k$	<b><math>P_{10}</math></b>  $\mathcal{P}_{10} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dP_{10}</math></b>  $\mathcal{P}_{10} \Lambda^k$	<b><math>P_{11}</math></b>  $\mathcal{P}_{11} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dP_{11}</math></b>  $\mathcal{P}_{11} \Lambda^k$	<b><math>P_{12}</math></b>  $\mathcal{P}_{12} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dP_{12}</math></b>  $\mathcal{P}_{12} \Lambda^k$	<b><math>P_{13}</math></b>  $\mathcal{P}_{13} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dP_{13}</math></b>  $\mathcal{P}_{13} \Lambda^k$	<b><math>P_{14}</math></b>  $\mathcal{P}_{14} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dP_{14}</math></b>  $\mathcal{P}_{14} \Lambda^k$	<b><math>P_{15}</math></b>  $\mathcal{P}_{15} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dP_{15}</math></b>  $\mathcal{P}_{15} \Lambda^k$	<b><math>P_{16}</math></b>  $\mathcal{P}_{16} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dP_{16}</math></b>  $\mathcal{P}_{16} \Lambda^k$	<b><math>P_{17}</math></b>  $\mathcal{P}_{17} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dP_{17}</math></b>  $\mathcal{P}_{17} \Lambda^k$	<b><math>P_{18}</math></b>  $\mathcal{P}_{18} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dP_{18}</math></b>  $\mathcal{P}_{18} \Lambda^k$	<b><math>P_{19}</math></b>  $\mathcal{P}_{19} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dP_{19}</math></b>  $\mathcal{P}_{19} \Lambda^k$	<b><math>P_{20}</math></b>  $\mathcal{P}_{20} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dP_{20}</math></b>  $\mathcal{P}_{20} \Lambda^k$	<b><math>N1</math></b>  $\mathcal{N1} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN1</math></b>  $\mathcal{N1} \Lambda^k$	<b><math>N1'</math></b>  $\mathcal{N1'} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN1'</math></b>  $\mathcal{N1'} \Lambda^k$	<b><math>N2</math></b>  $\mathcal{N2} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN2</math></b>  $\mathcal{N2} \Lambda^k$	<b><math>N2'</math></b>  $\mathcal{N2'} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN2'</math></b>  $\mathcal{N2'} \Lambda^k$	<b><math>N3</math></b>  $\mathcal{N3} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN3</math></b>  $\mathcal{N3} \Lambda^k$	<b><math>N4</math></b>  $\mathcal{N4} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN4</math></b>  $\mathcal{N4} \Lambda^k$	<b><math>N5</math></b>  $\mathcal{N5} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN5</math></b>  $\mathcal{N5} \Lambda^k$	<b><math>N6</math></b>  $\mathcal{N6} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN6</math></b>  $\mathcal{N6} \Lambda^k$	<b><math>N7</math></b>  $\mathcal{N7} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN7</math></b>  $\mathcal{N7} \Lambda^k$	<b><math>N8</math></b>  $\mathcal{N8} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN8</math></b>  $\mathcal{N8} \Lambda^k$	<b><math>N9</math></b>  $\mathcal{N9} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN9</math></b>  $\mathcal{N9} \Lambda^k$	<b><math>N10</math></b>  $\mathcal{N10} \Lambda^k$ <b><math>RT1^{(0)}</math></b>  $\mathcal{RT1}^{(0)} \Lambda^k$ <b><math>dN10</math></b>  $\mathcal{N10} \Lambda^k$	<b><math>N11</math></b>  $\mathcal{N11} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN11</math></b>  $\mathcal{N11} \Lambda^k$	<b><math>N11'</math></b>  $\mathcal{N11'} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN11'</math></b>  $\mathcal{N11'} \Lambda^k$	<b><math>N12</math></b>  $\mathcal{N12} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN12</math></b>  $\mathcal{N12} \Lambda^k$	<b><math>N12'</math></b>  $\mathcal{N12'} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN12'</math></b>  $\mathcal{N12'} \Lambda^k$	<b><math>N13</math></b>  $\mathcal{N13} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN13</math></b>  $\mathcal{N13} \Lambda^k$	<b><math>N14</math></b>  $\mathcal{N14} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN14</math></b>  $\mathcal{N14} \Lambda^k$	<b><math>N15</math></b>  $\mathcal{N15} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN15</math></b>  $\mathcal{N15} \Lambda^k$	<b><math>N16</math></b>  $\mathcal{N16} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN16</math></b>  $\mathcal{N16} \Lambda^k$	<b><math>N17</math></b>  $\mathcal{N17} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN17</math></b>  $\mathcal{N17} \Lambda^k$	<b><math>N18</math></b>  $\mathcal{N18} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN18</math></b>  $\mathcal{N18} \Lambda^k$	<b><math>N19</math></b>  $\mathcal{N19} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN19</math></b>  $\mathcal{N19} \Lambda^k$	<b><math>N20</math></b>  $\mathcal{N20} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN20</math></b>  $\mathcal{N20} \Lambda^k$	<b><math>N21</math></b>  $\mathcal{N21} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN21</math></b>  $\mathcal{N21} \Lambda^k$	<b><math>N22</math></b>  $\mathcal{N22} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN22</math></b>  $\mathcal{N22} \Lambda^k$	<b><math>N23</math></b>  $\mathcal{N23} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN23</math></b>  $\mathcal{N23} \Lambda^k$	<b><math>N24</math></b>  $\mathcal{N24} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN24</math></b>  $\mathcal{N24} \Lambda^k$	<b><math>N25</math></b>  $\mathcal{N25} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN25</math></b>  $\mathcal{N25} \Lambda^k$	<b><math>N26</math></b>  $\mathcal{N26} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN26</math></b>  $\mathcal{N26} \Lambda^k$	<b><math>N27</math></b>  $\mathcal{N27} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN27</math></b>  $\mathcal{N27} \Lambda^k$	<b><math>N28</math></b>  $\mathcal{N28} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN28</math></b>  $\mathcal{N28} \Lambda^k$	<b><math>N29</math></b>  $\mathcal{N29} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN29</math></b>  $\mathcal{N29} \Lambda^k$	<b><math>N30</math></b>  $\mathcal{N30} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN30</math></b>  $\mathcal{N30} \Lambda^k$	<b><math>N31</math></b>  $\mathcal{N31} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN31</math></b>  $\mathcal{N31} \Lambda^k$	<b><math>N32</math></b>  $\mathcal{N32} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN32</math></b>  $\mathcal{N32} \Lambda^k$	<b><math>N33</math></b>  $\mathcal{N33} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN33</math></b>  $\mathcal{N33} \Lambda^k$	<b><math>N34</math></b>  $\mathcal{N34} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN34</math></b>  $\mathcal{N34} \Lambda^k$	<b><math>N35</math></b>  $\mathcal{N35} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN35</math></b>  $\mathcal{N35} \Lambda^k$	<b><math>N36</math></b>  $\mathcal{N36} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN36</math></b>  $\mathcal{N36} \Lambda^k$	<b><math>N37</math></b>  $\mathcal{N37} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN37</math></b>  $\mathcal{N37} \Lambda^k$	<b><math>N38</math></b>  $\mathcal{N38} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN38</math></b>  $\mathcal{N38} \Lambda^k$	<b><math>N39</math></b>  $\mathcal{N39} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN39</math></b>  $\mathcal{N39} \Lambda^k$	<b><math>N40</math></b>  $\mathcal{N40} \Lambda^k$ <b><math>RT1^{(1)}</math></b>  $\mathcal{RT1}^{(1)} \Lambda^k$ <b><math>dN40</math></b>  $\mathcal{N40} \Lambda^k$

		$\dim \mathcal{P}_r(\mathcal{H}(L))$
$k$	$r=1$	
$k=1$	0	1
	1	1
$k=2$	0	1
	1	3
	2	1
$k=3$	0	1
	1	0
	2	0
	3	1
$k=4$	0	1
	1	10
	2	10
	3	10

		dim $\mathbb{P}_k A^k(\Delta_n)$		
$k$	$n=1$	2	3	4
$n=1$	0	2	3	4
		1	2	4
$n=2$	0	3	8	10
	1	6	12	20
	2	3	5	10
$n=3$	0	4	10	20
	1	12	20	30
	2	12	30	50
	3	4	10	20
$n=4$	0	5	15	35
	1	20	40	80
	2	20	50	120
	3	5	10	20

	1	5	8	3	
1	6	3	1	6	c=1
2	8	5	2	8	
3	21	26	36	21	c=2
4	43	26	36	43	
5	21	26	36	21	c=3
6	56	64	128	56	
7	108	203	308	108	c=4
8	168	203	308	168	
9	56	64	128	56	c=5
10	126	216	300	126	
11	564	946	1226	564	c=6
12	796	1296	1696	796	

$n$	$Q_n^r A^k(\square_n)$
1	$\binom{0}{0} r^0 (r+1)^{n-k}$
2	$\binom{2}{0} r^2 (r+1)^{n-4}$
3	$\binom{3}{0} r^3 (r+1)^{n-6}$
4	$\binom{4}{0} r^4 (r+1)^{n-8}$
5	$\binom{5}{0} r^5 (r+1)^{n-10}$
6	$\binom{6}{0} r^6 (r+1)^{n-12}$
7	$\binom{7}{0} r^7 (r+1)^{n-14}$
8	$\binom{8}{0} r^8 (r+1)^{n-16}$
9	$\binom{9}{0} r^9 (r+1)^{n-18}$
10	$\binom{10}{0} r^{10} (r+1)^{n-20}$
11	$\binom{11}{0} r^{11} (r+1)^{n-22}$
12	$\binom{12}{0} r^{12} (r+1)^{n-24}$
13	$\binom{13}{0} r^{13} (r+1)^{n-26}$
14	$\binom{14}{0} r^{14} (r+1)^{n-28}$
15	$\binom{15}{0} r^{15} (r+1)^{n-30}$
16	$\binom{16}{0} r^{16} (r+1)^{n-32}$
17	$\binom{17}{0} r^{17} (r+1)^{n-34}$
18	$\binom{18}{0} r^{18} (r+1)^{n-36}$
19	$\binom{19}{0} r^{19} (r+1)^{n-38}$
20	$\binom{20}{0} r^{20} (r+1)^{n-40}$
21	$\binom{21}{0} r^{21} (r+1)^{n-42}$
22	$\binom{22}{0} r^{22} (r+1)^{n-44}$
23	$\binom{23}{0} r^{23} (r+1)^{n-46}$
24	$\binom{24}{0} r^{24} (r+1)^{n-48}$
25	$\binom{25}{0} r^{25} (r+1)^{n-50}$
26	$\binom{26}{0} r^{26} (r+1)^{n-52}$
27	$\binom{27}{0} r^{27} (r+1)^{n-54}$
28	$\binom{28}{0} r^{28} (r+1)^{n-56}$
29	$\binom{29}{0} r^{29} (r+1)^{n-58}$
30	$\binom{30}{0} r^{30} (r+1)^{n-60}$
31	$\binom{31}{0} r^{31} (r+1)^{n-62}$
32	$\binom{32}{0} r^{32} (r+1)^{n-64}$
33	$\binom{33}{0} r^{33} (r+1)^{n-66}$
34	$\binom{34}{0} r^{34} (r+1)^{n-68}$
35	$\binom{35}{0} r^{35} (r+1)^{n-70}$
36	$\binom{36}{0} r^{36} (r+1)^{n-72}$
37	$\binom{37}{0} r^{37} (r+1)^{n-74}$
38	$\binom{38}{0} r^{38} (r+1)^{n-76}$
39	$\binom{39}{0} r^{39} (r+1)^{n-78}$
40	$\binom{40}{0} r^{40} (r+1)^{n-80}$
41	$\binom{41}{0} r^{41} (r+1)^{n-82}$
42	$\binom{42}{0} r^{42} (r+1)^{n-84}$
43	$\binom{43}{0} r^{43} (r+1)^{n-86}$
44	$\binom{44}{0} r^{44} (r+1)^{n-88}$
45	$\binom{45}{0} r^{45} (r+1)^{n-90}$
46	$\binom{46}{0} r^{46} (r+1)^{n-92}$
47	$\binom{47}{0} r^{47} (r+1)^{n-94}$
48	$\binom{48}{0} r^{48} (r+1)^{n-96}$
49	$\binom{49}{0} r^{49} (r+1)^{n-98}$
50	$\binom{50}{0} r^{50} (r+1)^{n-100}$

		dim $S_n H^k(\mathbb{C}_n)$	
$n$	$k$	$n=1$	$n=2$
1	2	6	22
2	3	6	22
3	2	7	24
4	5	6	54
5	112	7	96
6	13	7	33
7	223	6	0
8	1396	7	24
9	1779	2	96
10	363	3	24
11	4896	6	36
12	14208	7	64
13	10559	2	72

$m$	$n(m+2)/k$	$2^{n-d}$	$\binom{n}{d}$	$\binom{r-d+2k}{d}$	$\binom{s}{k}$
3	4	8	1	2	3
4	5	16	4	7	10
5	6	32	10	20	35
6	7	64	20	50	90
7	8	128	40	100	190
8	9	256	80	200	400
9	10	512	160	400	800
10	11	1024	320	800	1600
11	12	2048	640	1600	3200
12	13	4096	1280	3200	6400
13	14	8192	2560	6400	12800
14	15	16384	5120	12800	25600
15	16	32768	10240	25600	51200
16	17	65536	20480	51200	102400
17	18	131072	40960	102400	204800
18	19	262144	81920	204800	409600
19	20	524288	163840	409600	819200
20	21	1048576	327680	819200	1638400
21	22	2097152	655360	1638400	3276800
22	23	4194304	1310720	3276800	6553600
23	24	8388608	2621440	6553600	13107200
24	25	16777216	5242880	13107200	26214400
25	26	33554432	10485760	26214400	52428800
26	27	67108864	20971520	52428800	104857600
27	28	134217728	41943040	104857600	209715200
28	29	268435456	83886080	209715200	419430400
29	30	536870912	167772160	419430400	838860800
30	31	1073741824	335544320	838860800	1677721600
31	32	2147483648	671088640	1677721600	3355443200
32	33	4294967296	1342177280	3355443200	6710886400
33	34	8589934592	2684354560	6710886400	13421772800
34	35	17179869184	5368709120	13421772800	26843545600
35	36	34359738368	10737418240	26843545600	53687091200
36	37	68719476736	21474836480	53687091200	107374182400
37	38	137438953472	42949672960	107374182400	214748364800
38	39	274877906944	85899345920	214748364800	429496729600
39	40	549755813888	171798691840	429496729600	858993459200
40	41	1099511627760	343597383680	858993459200	1717986918400
41	42	2199023255520	687194767360	1717986918400	3435973836800
42	43	4398046511040	1374389534720	3435973836800	6871947673600
43	44	8796093022080	2748779069440	6871947673600	13743895347200
44	45	17592186044160	5497558138880	13743895347200	27487790694400
45	46	35184372088320	10995116277600	27487790694400	54975581388800
46	47	70368744176640	21990232555200	54975581388800	109951162776000
47	48	140737488353280	43980465110400	109951162776000	219902325552000
48	49	281474976706560	87960930220800	219902325552000	439804651104000
49	50	562949953413120	175921860441600	439804651104000	879609302208000
50	51	1125899906826240	351843720883200	879609302208000	1759218604416000
51	52	2251799813652480	703687441766400	1759218604416000	3518437208832000
52	53	4503599627304960	1407374883532800	3518437208832000	7036874417664000
53	54	9007199254609920	2814749767065600	7036874417664000	14073748835328000
54	55	18014398509219840	5629499534131200	14073748835328000	28147497670656000
55	56	36028797018439680	11258999068262400	28147497670656000	56294995341312000
56	57	72057594036879360	22517998136524800	56294995341312000	112589990682624000
57	58	144115188073758720	45035996273049600	112589990682624000	225179981365248000
58	59	288230376147517440	90071992546099200	225179981365248000	450359962730496000
59	60	576460752295034880	180143985092198400	450359962730496000	900719925460992000
60	61	1152921504590069600	2882303761475174400	900719925460992000	1801439850921984000
61	62	2305843009180138400	5764607522950348800	1801439850921984000	3602879701843968000
62	63	4611686018360276800	11529215045900696000	3602879701843968000	7205759403687936000
63	64	9223372036720553600	23058430091801384000	7205759403687936000	14411518807375872000
64	65	18446744073441107200	46116860183602768000	14411518807375872000	28823037614751744000
65	66	36893488146882214400	92233720367205536000	28823037614751744000	57646075229503488000
66	67	73786976293764428800	184467440734411072000	57646075229503488000	11529215045900696000
67	68	147573952587528857600	368934881468822144000	11529215045900696000	23058430091801384000
68	69	295147905175057715200	737869762937644288000	23058430091801384000	46116860183602768000
69	70	590295810350115430400	1475739525875288576000	46116860183602768000	92233720367205536000
70	71	1180591620700230860800	2951479051750577152000	92233720367205536000	184467440734411072000
71	72	2361183241400461721600	5902958103501154304000	184467440734411072000	368934881468822144000
72	73	4722366482800923443200	11805916207002308608000	368934881468822144000	737869762937644288000
73	74	9444732965601846886400	23611832414004617216000	737869762937644288000	1475739525875288576000
74	75	18889465931203693772000	47223664828009234432000	1475739525875288576000	2951479051750577152000
75	76	37778931862407387544000	94447329656018468864000	2951479051750577152000	5902958103501154304000
76	77	75557863724814775088000	188894659312036937720000	5902958103501154304000	11805916207002308608000
77	78	151115727449629550160000	377789318624073875440000	11805916207002308608000	236118324140046172160000
78	79	302231454899259100320000	755578637248147750880000	236118324140046172160000	472236648280092344320000
79	80	604462909798518200640000	1511157274496295501600000	4722366482800923443200000	9444732965601846886400000
80	81	1208925819597036401280000	3022314548992591003200000	9444732965601846886400000	18889465931203693772000000
81	82	2417851639194072802560000	60446290979851820064000000	18889465931203693772000000	377789318624073875440000000
82	83	4835703278388145605120000	120892581959703640128000000	188894659312036937720000000	755578637248147750880000000
83	84	9671406556776291210240000	241785163919407280256000000	1888946593120369377200000000	1511157274496295501600000000
84	85	19342813113552582420480000	483570327838814560512000000	18889465931203693772000000000	30223145489925910032000000000
85	86	38685626227105164840960000	967140655677629121024000000	188894659312036937720000000000	604462909798518200640000000000
86	87	77371252454210329681920000	1934281311355258242048000000	1888946593120369377200000000000	12089258195970364012800000000000
87	88	154742504908420659363840000	3868562622710516484096000000	18889465931203693772000000000000	241785163919407280256000000000000
88	89	309485009816841318727680000	7737125245421032968192000000	188894659312036937720000000000000	4835703278388145605120000000000000
89	90	618970019633682637455360000	15474250490842065936384000000	1888946593120369377200000000000000	9671406556776291210240000000000000
90	91	1237940039267365278910720000	30948500981684131872768000000	1888946593120369377200000000000000	19342813113552582420480000000000000
91	92	2475880078534730557821440000	61897001963368263745536000000	1888946593120369377200000000000000	38685626227105164840960000000000000
92	93	4951760017069461115642880000	123794003926736527891072000000	1888946593120369377200000000000000	77371252454210329681920000000000000
93	94	9903520034138922231285760000	247588007853473055782144000000	1888946593120369377200000000000000	154742504908420659363840000000000000
94	95	1980704006827784446257520000	495176001706946111564288000000	1888946593120369377200000000000000	3094850098168413187276800000000000000
95	96	3961408013655568892515040000	990352003413892223128576000000	1888946593120369377200000000000000	61897001963368263745536000000000000000
96	97	7922816027311137785030080000	198070400682778444625752000000	1888946593120369377200000000000000	123794003926736527891072000000000000000
97	98	15845632054622275570060160000	396140801365556889251504000000	1888946593120369377200000000000000	247588007853473055782144000000000000000
98	99	31691264109244551140120320000	792281602731113778503008000000	1888946593120369377200000000000000	495176001706946111564288000000000000000
99	100	63382528218489102280240640000	1584563205462227557006016000000	1888946593120369377200000000000000	990352003413892223128576000000000000000

The table presents a summary of the student, cut and analysis of variance results of the two-dimensional IMA study. The above functions, shape function of each ICF being assessed and specifying a cut showing the lack of association between the two variables.

elements

The primary purpose of finite elements for the numerical solution of boundary value problems is to find a discrete space or a set of functions,  $\mathcal{V}$ , which approximate the continuous function space defined by (2). A typical element of  $\mathcal{V}$  is called a shape function. The set of all shape functions on an element is called a basis. A collection of basis functions on an element is called a shape function space. If  $\mathcal{V}$  is a generalized space of the element, then each basis function has a single value at every point in the element. Diagrams depicting (1)–(3) and their

reading. Thus for  $k=1$ , the approach describes the first iteration of the gradient operation (see Fig. 1). Then for  $k>1$  the  $\mathbf{C}_k$  and  $\mathbf{M}_k$  are  $\mathbf{I} - \mathbf{A}_k^{-1}$  they characterize the  $k$ -th iteration in  $S$ -resonance, and for  $k < 0$  they characterize  $S^*$ -resonance. The  $\mathbf{P}_k$  and  $\mathbf{A}_k$  which occur here are the iterates, going back in the case  $k = -1$  to the initial state, and so on. We can implement in an algorithm  $\mathbf{P}_k$  and  $\mathbf{A}_k$ , which  $k$  includes, as a sequence of successive, connected, pieces of programs. This is what we have done, first in  $S$ -resonance.

In space H<sup>1</sup>,  
like H<sup>1</sup>(Ω),  
the domain  
is measurable  
with respect to  
Carathéodory's measure. The  
distance function  
is continuous with  
respect to the metric by Stoyan's (P)<sup>1</sup>  
and Rogers' (P)<sup>2</sup>  
approaches.  
In the theory of  
the p.v. A  
concept of weak  
convergence for the

and can be ordered from the Sales Department, University of Alberta, Box 23500, Edmonton, Alberta T6G 2G4, Canada. Price and postage for the first issue are \$10.00 and for the second issue \$10.00.

Journal of Numerical Mathematics 49, 1942.  
J. Math. Mech. 10, 1961.  
Mathematics Education Series, Springer, 1977.  
Wolfram Mathematica, 17, 1990.  
Wolfram Alpha, 2008.  
Wolfram Demonstrations Project, 2011.  
Wolfram Alpha, 2011.  
Automated Collision of Differential Equations, 2012.  
Guttmann, 2004.  
Mathematical Software, 2010.  
en.Julie.Wolfram, ACM Transactions

— 10 —

The table presents the primary aspects of finite elements for the approximation of the fundamental properties of electric fields: the gradient, curl and divergence. These primary spaces are in space of physical quantities, and therefore they are called primary. The second row of the table shows the primary spaces of each element (20) while the third row of the table shows the basis functions of each element called the shape functions. The last two rows of the table show the mapping of each element to a reference element, and the mapping of each shape function of the element to its corresponding shape function of the reference element.

Then for  $i = C$  the second derivative in the Sobolev space  $H^1(\Omega)$ ,  $\|u_i\|_{H^1(\Omega)} \leq C \|f_i\|_{L^2(\Omega)}$ . In the case  $i = 1 - C$  we have  $u_i \in H^1(\Omega)$ , the domain of  $A_1$ , and for  $i < 1 - C$  this is so too.

It follows from the above that the solution  $u$  is a weak solution going back in the case  $i = 1$  to the elements of  $C_c^\infty(\Omega)$  and in the case  $i = 0$  to the elements of  $C_0(\Omega)$ . Theorem 1 is proved.

**Remark 1.** The proof presented here is based on the method of successive approximations, consisting of iterations with respect to the parameter  $\lambda$ . It is also possible to introduce by short steps the parameter  $\lambda$  in the Schauder and spaces  $H^1(\Omega)$  and  $L^2(\Omega)$  and to obtain the corresponding results for the case  $i = 0$  in 2 dimensions. As a generalization to 3 dimensions, one can consider the case  $i = 1$  in the Sobolev and spaces  $H^1(\Omega)$  and  $L^2(\Omega)$ , the uniform regularity and continuity of the  $y$ -derivative of the solution  $u$  being established by the same methods as in the case of the "classical" elliptic boundary value problems. A detailed analysis of the problem in the case  $i = 1$  is given in [1].

be derived from the [US Environmental Protection Agency's "Intergovernmental Cooperation and Decisions" page](#) ([http://www.epa.gov/igc/decisions.html](#)). The second table is taken from [EPA's website](#).

The information is part of the [EPA's "Intergovernmental Cooperation and Decisions" page](#) ([http://www.epa.gov/igc/decisions.html](#)).

30. [EPA's website](#)

31. [EPA's website](#)

32. [EPA's website](#)

33. [EPA's website](#)

34. [EPA's website](#)

35. [EPA's website](#)

36. [EPA's website](#)

37. [EPA's website](#)

38. [EPA's website](#)

39. [EPA's website](#)

40. [EPA's website](#)

41. [EPA's website](#)

42. [EPA's website](#)

43. [EPA's website](#)

44. [EPA's website](#)

45. [EPA's website](#)

46. [EPA's website](#)

47. [EPA's website](#)

48. [EPA's website](#)

49. [EPA's website](#)

50. [EPA's website](#)

51. [EPA's website](#)

52. [EPA's website](#)

53. [EPA's website](#)

54. [EPA's website](#)

55. [EPA's website](#)

56. [EPA's website](#)

57. [EPA's website](#)

58. [EPA's website](#)

59. [EPA's website](#)

60. [EPA's website](#)

61. [EPA's website](#)

62. [EPA's website](#)

63. [EPA's website](#)

64. [EPA's website](#)

65. [EPA's website](#)

66. [EPA's website](#)

67. [EPA's website](#)

68. [EPA's website](#)

69. [EPA's website](#)

70. [EPA's website](#)

71. [EPA's website](#)

72. [EPA's website](#)

73. [EPA's website](#)

74. [EPA's website](#)

75. [EPA's website](#)

76. [EPA's website](#)

77. [EPA's website](#)

78. [EPA's website](#)

79. [EPA's website](#)

80. [EPA's website](#)

81. [EPA's website](#)

82. [EPA's website](#)

83. [EPA's website](#)

84. [EPA's website](#)

85. [EPA's website](#)

86. [EPA's website](#)

87. [EPA's website](#)

88. [EPA's website](#)

89. [EPA's website](#)

90. [EPA's website](#)

91. [EPA's website](#)

92. [EPA's website](#)

93. [EPA's website](#)

94. [EPA's website](#)

95. [EPA's website](#)

96. [EPA's website](#)

97. [EPA's website](#)

98. [EPA's website](#)

99. [EPA's website](#)

100. [EPA's website](#)

101. [EPA's website](#)

102. [EPA's website](#)

103. [EPA's website](#)

104. [EPA's website](#)

105. [EPA's website](#)

106. [EPA's website](#)

107. [EPA's website](#)

108. [EPA's website](#)

109. [EPA's website](#)

110. [EPA's website](#)

111. [EPA's website](#)

112. [EPA's website](#)

113. [EPA's website](#)

114. [EPA's website](#)

115. [EPA's website](#)

116. [EPA's website](#)

117. [EPA's website](#)

118. [EPA's website](#)

119. [EPA's website](#)

120. [EPA's website](#)

121. [EPA's website](#)

122. [EPA's website](#)

123. [EPA's website](#)

124. [EPA's website](#)

125. [EPA's website](#)

126. [EPA's website](#)

127. [EPA's website](#)

128. [EPA's website](#)

129. [EPA's website](#)

130. [EPA's website](#)

131. [EPA's website](#)

132. [EPA's website](#)

133. [EPA's website](#)

134. [EPA's website](#)

135. [EPA's website](#)

136. [EPA's website](#)

137. [EPA's website](#)

138. [EPA's website](#)

139. [EPA's website](#)

140. [EPA's website](#)

141. [EPA's website](#)

142. [EPA's website](#)

143. [EPA's website](#)

144. [EPA's website](#)

145. [EPA's website](#)

146. [EPA's website](#)

147. [EPA's website](#)

148. [EPA's website](#)

149. [EPA's website](#)

150. [EPA's website](#)

151. [EPA's website](#)

152. [EPA's website](#)

153. [EPA's website](#)

154. [EPA's website](#)

155. [EPA's website](#)

156. [EPA's website](#)

157. [EPA's website](#)

158. [EPA's website](#)

159. [EPA's website](#)

160. [EPA's website](#)

161. [EPA's website](#)

162. [EPA's website](#)

163. [EPA's website](#)

164. [EPA's website](#)

165. [EPA's website](#)

166. [EPA's website](#)

167. [EPA's website](#)

168. [EPA's website](#)

169. [EPA's website](#)

170. [EPA's website](#)

171. [EPA's website](#)

172. [EPA's website](#)

173. [EPA's website](#)

174. [EPA's website](#)

175. [EPA's website](#)

176. [EPA's website](#)

177. [EPA's website](#)

178. [EPA's website](#)

179. [EPA's website](#)

180. [EPA's website](#)

181. [EPA's website](#)

182. [EPA's website](#)

183. [EPA's website](#)

184. [EPA's website](#)

185. [EPA's website](#)

186. [EPA's website](#)

187. [EPA's website](#)

188. [EPA's website](#)

189. [EPA's website](#)

190. [EPA's website](#)

191. [EPA's website](#)

192. [EPA's website](#)

193. [EPA's website](#)

194. [EPA's website](#)

195. [EPA's website](#)

196. [EPA's website](#)

197. [EPA's website](#)

198. [EPA's website](#)

199. [EPA's website](#)

200. [EPA's website](#)

201. [EPA's website](#)

202. [EPA's website](#)

203. [EPA's website](#)

204. [EPA's website](#)

205. [EPA's website](#)

206. [EPA's website](#)

207. [EPA's website](#)

208. [EPA's website](#)

209. [EPA's website](#)

210. [EPA's website](#)

211. [EPA's website](#)

212. [EPA's website](#)

213. [EPA's website](#)

214. [EPA's website](#)

215. [EPA's website](#)

216. [EPA's website](#)

217. [EPA's website](#)

218. [EPA's website](#)

219. [EPA's website](#)

220. [EPA's website](#)

221. [EPA's website](#)

222. [EPA's website](#)

223. [EPA's website](#)

224. [EPA's website](#)

225. [EPA's website](#)

226. [EPA's website](#)

227. [EPA's website](#)

228. [EPA's website](#)

229. [EPA's website](#)

230. [EPA's website](#)

231. [EPA's website](#)

232. [EPA's website](#)

233. [EPA's website](#)

234. [EPA's website](#)

235. [EPA's website](#)

236. [EPA's website](#)

237. [EPA's website](#)

238. [EPA's website](#)

239. [EPA's website](#)

240. [EPA's website](#)

241. [EPA's website](#)

242. [EPA's website](#)

243. [EPA's website](#)

244. [EPA's website](#)

245. [EPA's website](#)

246. [EPA's website](#)

247. [EPA's website](#)

248. [EPA's website](#)

249. [EPA's website](#)

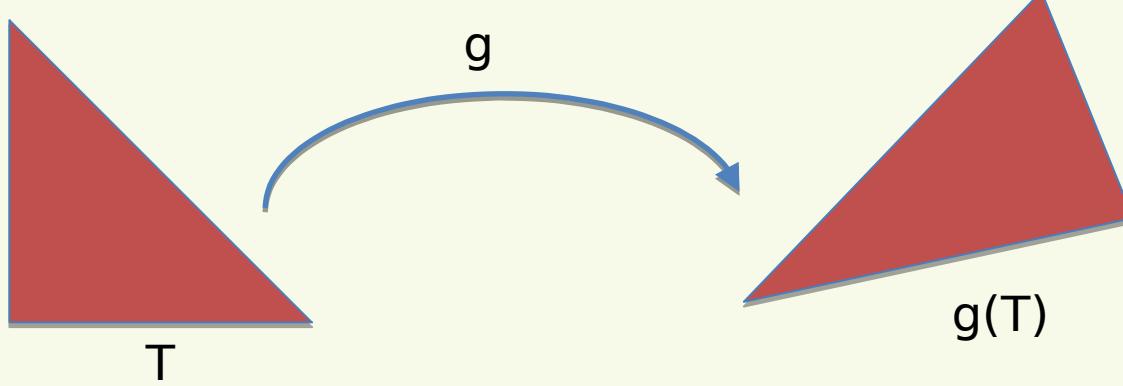
250. [EPA's website](#)

73

# Evaluating integrals over grids

How can we efficiently evaluate  $\int_{\Omega} [\sigma(\mathbf{x}) \nabla u(\mathbf{x})] \cdot \nabla v(\mathbf{x}) d\mathbf{x}$

over a given grid?



Reference  
Element

For simple flat triangles  
have:

$$g(x) = Jx + b$$

Physical Element

# Evaluating integrals over grids...

Integration by parts formula:

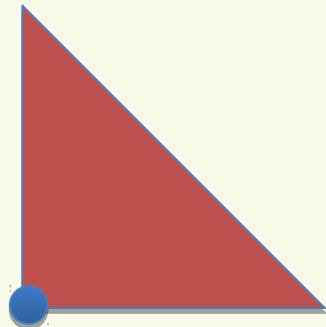
$$\int_{g(T)} f(\mathbf{y}) d\mathbf{y} = \int_T f(g(\mathbf{x})) |\det(J(\mathbf{x}))| d\mathbf{x}$$

↗  
Jacobain (derivative)  
of  $g$

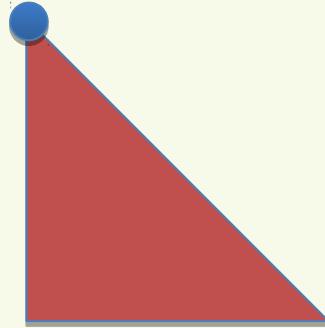
$$\int_{\Omega} [\sigma(\mathbf{y}) \nabla u(\mathbf{y})] \cdot \nabla v(\mathbf{y}) d\mathbf{y} = \sum_T \int_T [\sigma(g_T(\mathbf{x})) J_T^{-T}(\mathbf{x}) \nabla \hat{u}_T(\mathbf{x})] \cdot J_T^{-T} \nabla \hat{u}_T(\mathbf{x}) |\det(J_T(\mathbf{x}))| d\mathbf{x}$$

- Only integrals over triangles
- All quantities expressed in coordinates with respect to reference element
- Only need quadrature rule for reference element

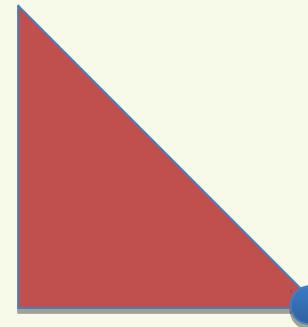
# Linear shape functions



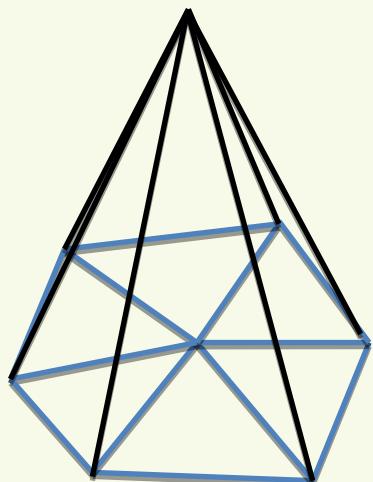
$$u_0(x, y) = 1 - x - y$$



$$u_1(x, y) = y$$



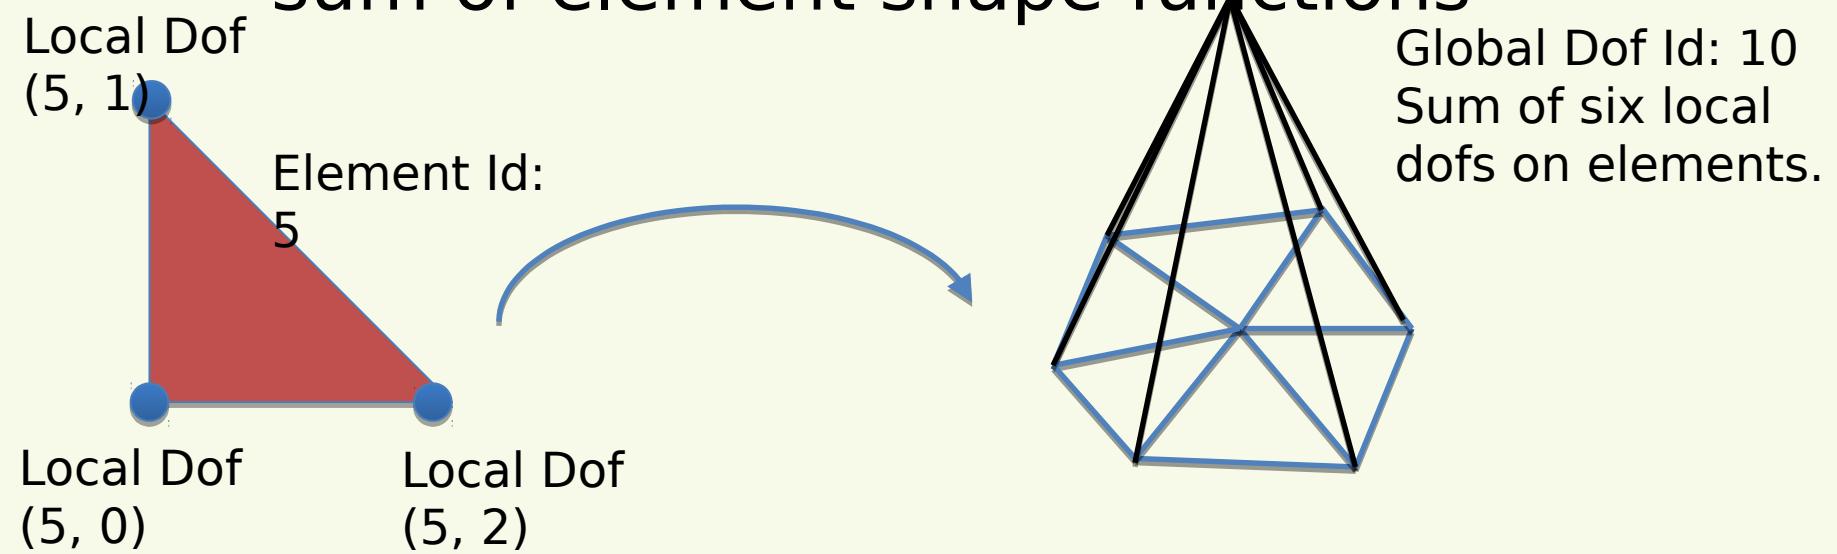
$$u_2(x, y) = x$$



- Each linear basis function is a linear combination of shape functions on each element.
- There are three different shape functions on the reference element.

# Dof maps

- Two different types of degrees of freedom
  - Local element shape functions indexed by element id and local id
  - Global P1 basis function consisting of a sum of element shape functions



# Dof Maps...

- Local-To-Global Map:
  - Element id -> A list of global indices that the local shape functions sum into
- Global-To-Local Map:
  - Global Id -> local shape functions. A list of local shape functions whose sum is the global ID.

```
V = FunctionSpace(mesh,  
"Lagrange", 1)  
dofmap = V.dofmap()
```

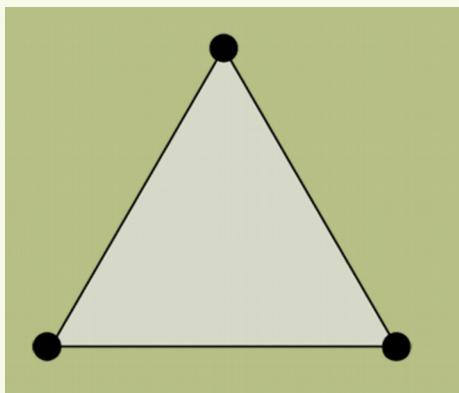
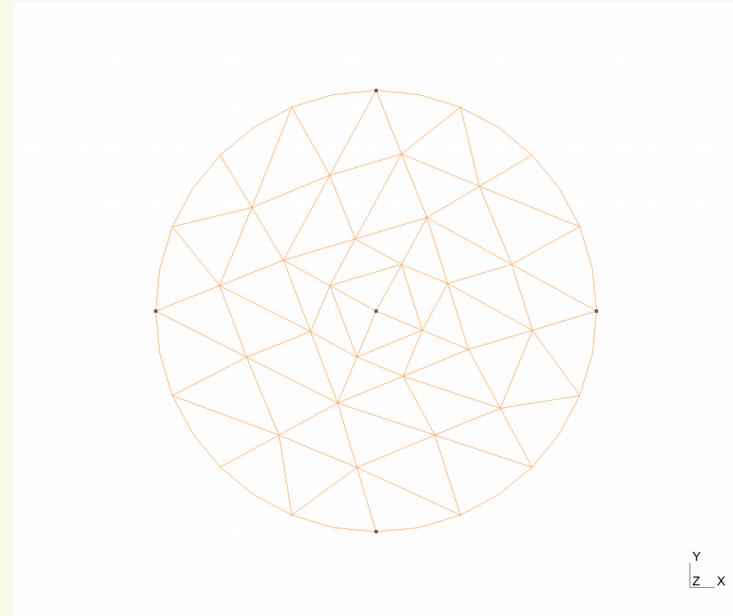
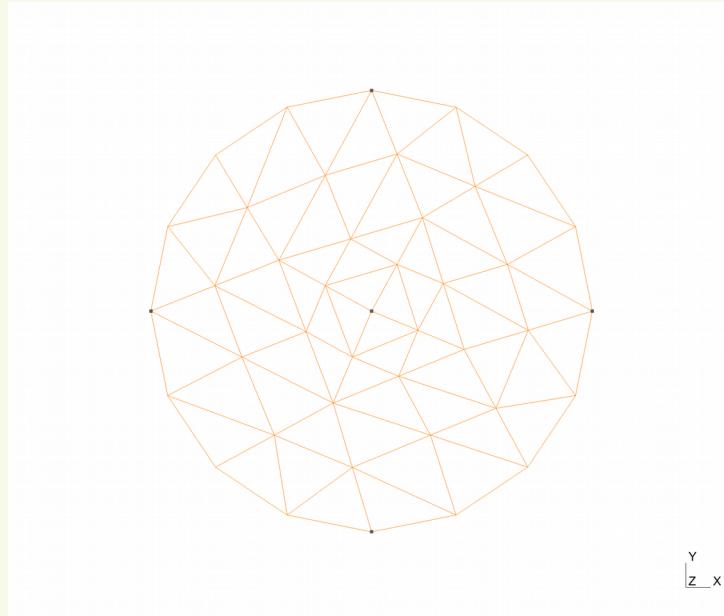
```
# Local to Global Dofs for element  
0:  
dofmap.cell_dofs(0)
```

# Assembling the global matrix

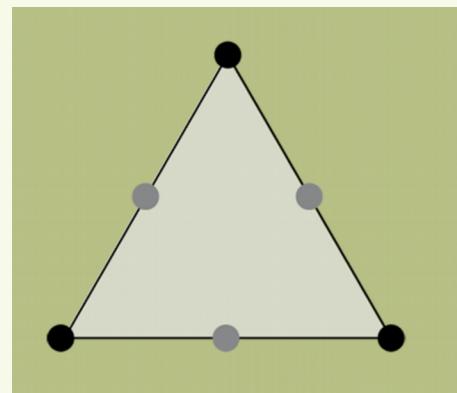
```
for cell in list of cells:  
    for i in local_dofs:  
        for j in local_dofs:  
            Evaluate a(local_dof_j, local_dof_i)  
            A[local_to_global(i), local_to_global(j)] +=  
a(local_dof_j, local_dof_i)
```

- In practice all weak forms on an element are assembled first and then summed into the global matrix to cache geometric information from the reference map
- Difficult to implement efficiently using Python as fast for loops needed
- Use Numba/PyOpenCL or similar packages to speed up assembly or code directly in C/C++
- Have not yet talked about parallel assembly. Will do this later.

# From low to high order

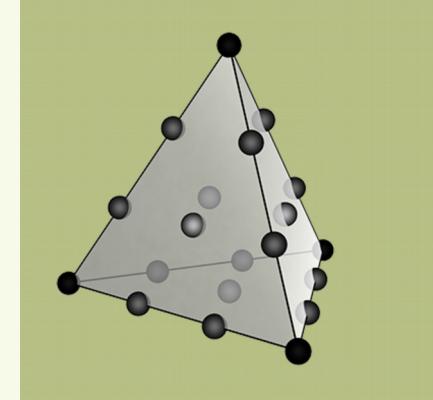
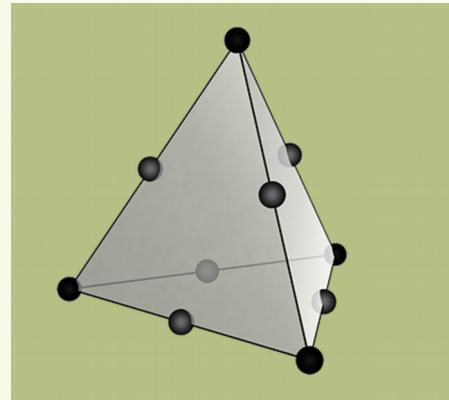
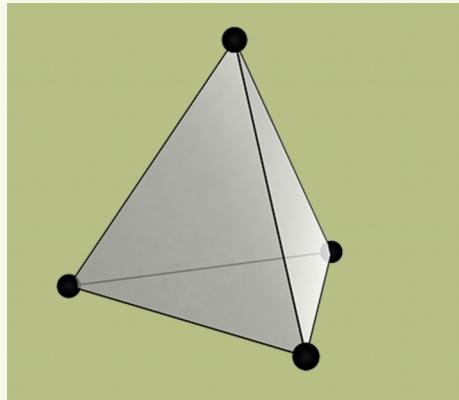
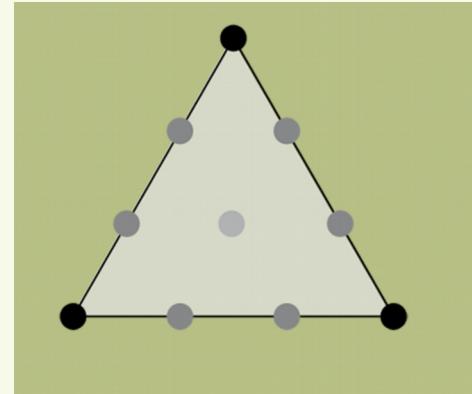
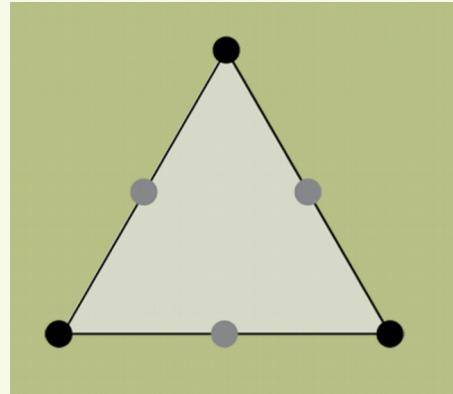
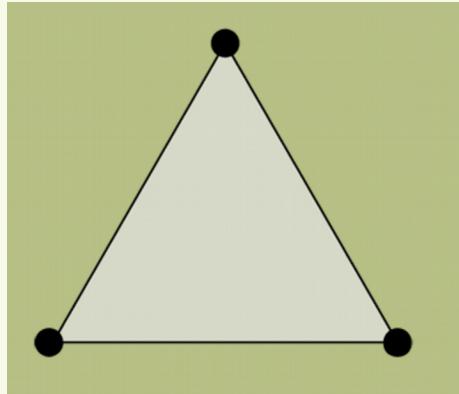


First order  
elements



Second order  
elements

# From low to high order...



Linear

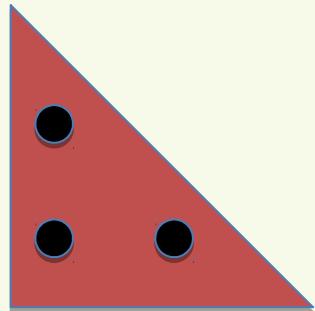
Quadrat  
ic

Cubi  
c

# Quadrature on the reference element

Approximate integrals over reference element by quadrature rule

$$\int_T f(\mathbf{x}) d\mathbf{x} \approx \sum_{j=1}^N f(\mathbf{x}_j) \omega_j$$



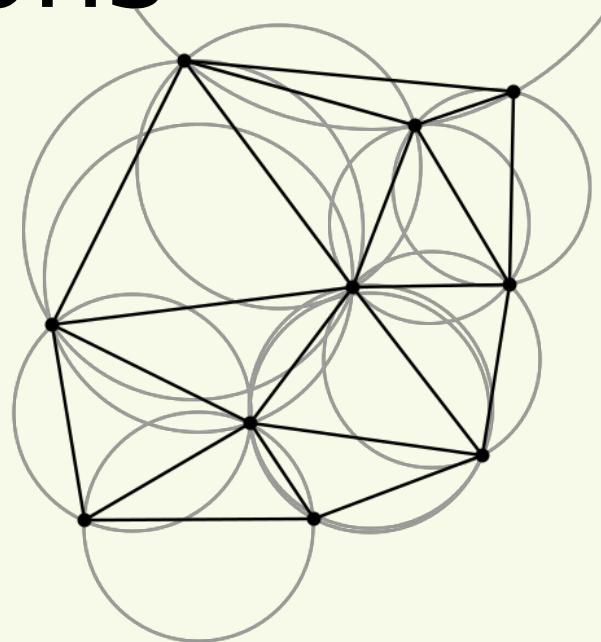
Symmetric Gauss Triangle quadrature rule:

- Find nodes and weights on the reference triangle such that as high polynomials as possible are integrated exactly
- Nodes and weights must be symmetric
- Detailed descriptions and node definitions for example contained in
- [http://math2.uncc.edu/~shaodeng/TEACHING/math5172/Lectures/Lec\\_15.PDF](http://math2.uncc.edu/~shaodeng/TEACHING/math5172/Lectures/Lec_15.PDF)

# Triangulation in two dimensions

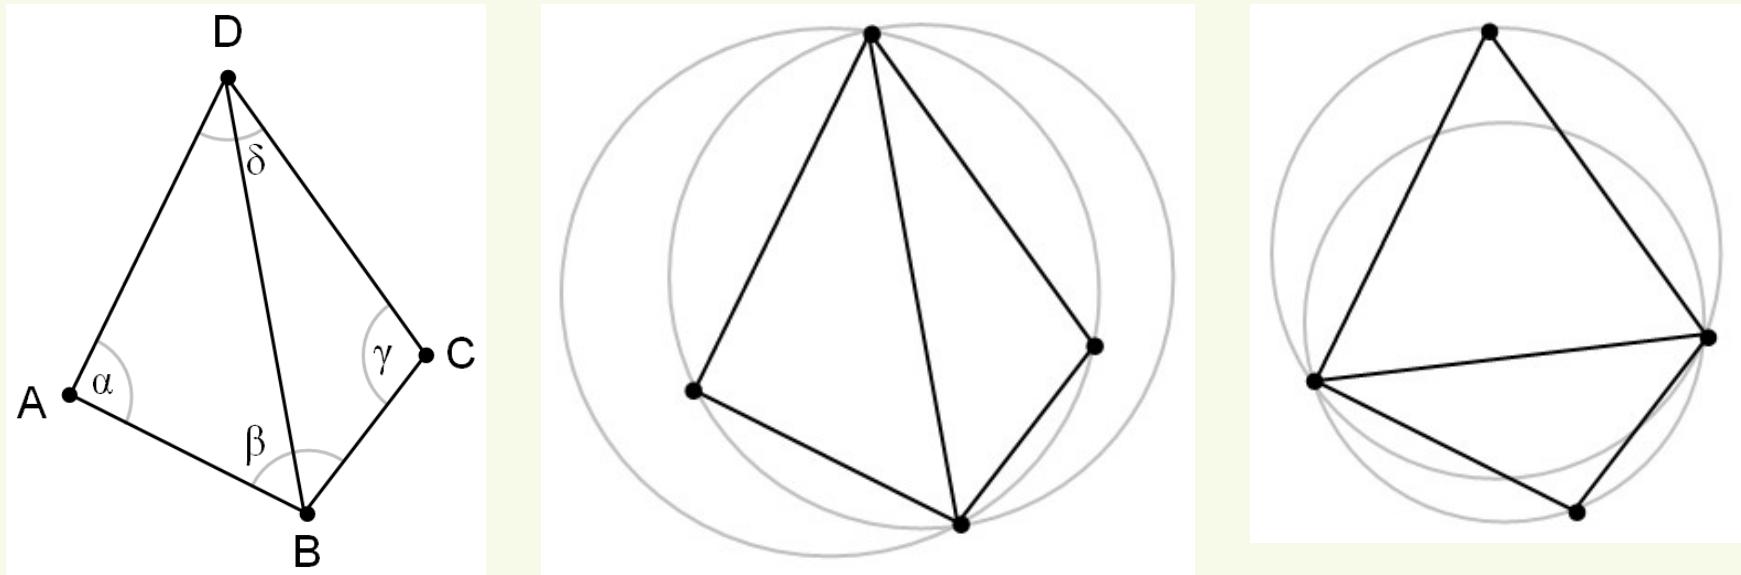
A set of triangles over a point set is called a triangulation if:

1. The union of the triangles equals the convex hull of the point set.
2. The union of all triangle vertices equals the point set.
3. Two distinct triangles share either a vertex, an edge or nothing. In particular, they do not overlap.



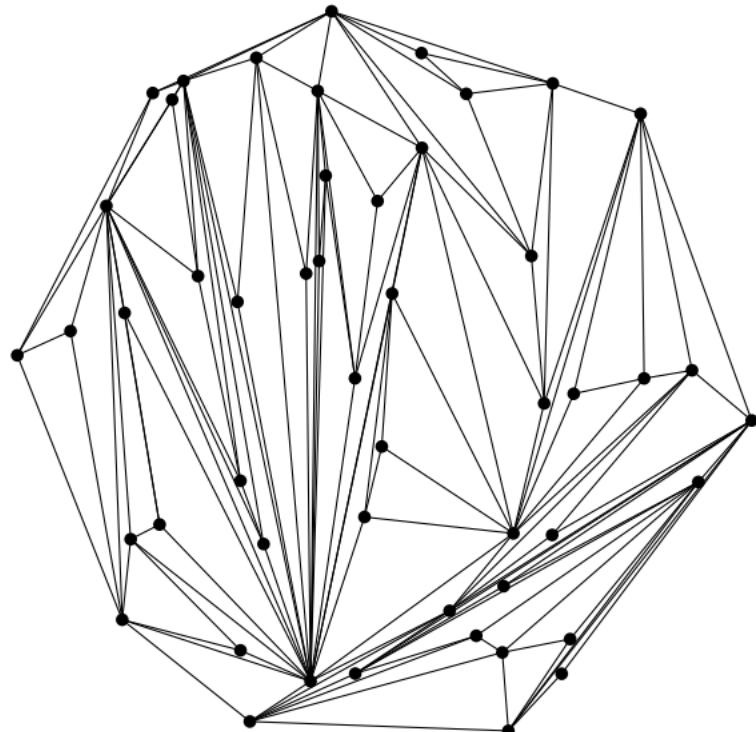
**Delaunay Triangulation:** A triangulation is called a Delaunay triangulation if the circumcircle of every triangle is empty, that is no point is contained inside the circumcircle.

# Delaunay Property

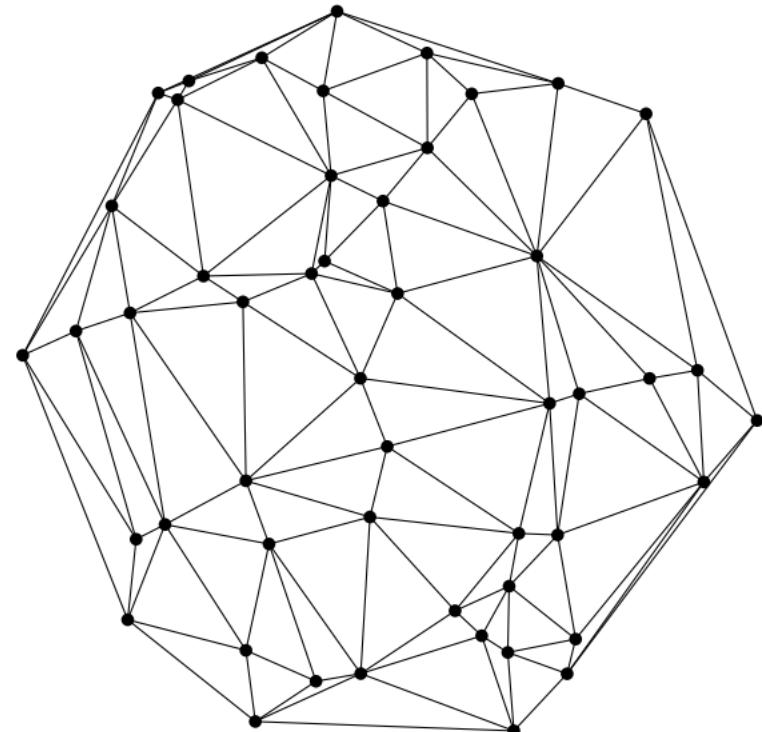


- An edge is called locally Delaunay if it has an open circumdisk containing no vertex of its adjacent triangles.
- The edge BD is not locally Delaunay.
- Flip BD into a new edge AC to become locally Delaunay
- Flip Algorithm: Given any triangulation. Perform flips until all edges are locally Delaunay. The resulting triangulation is Delaunay.

# Quality of Delaunay triangulation



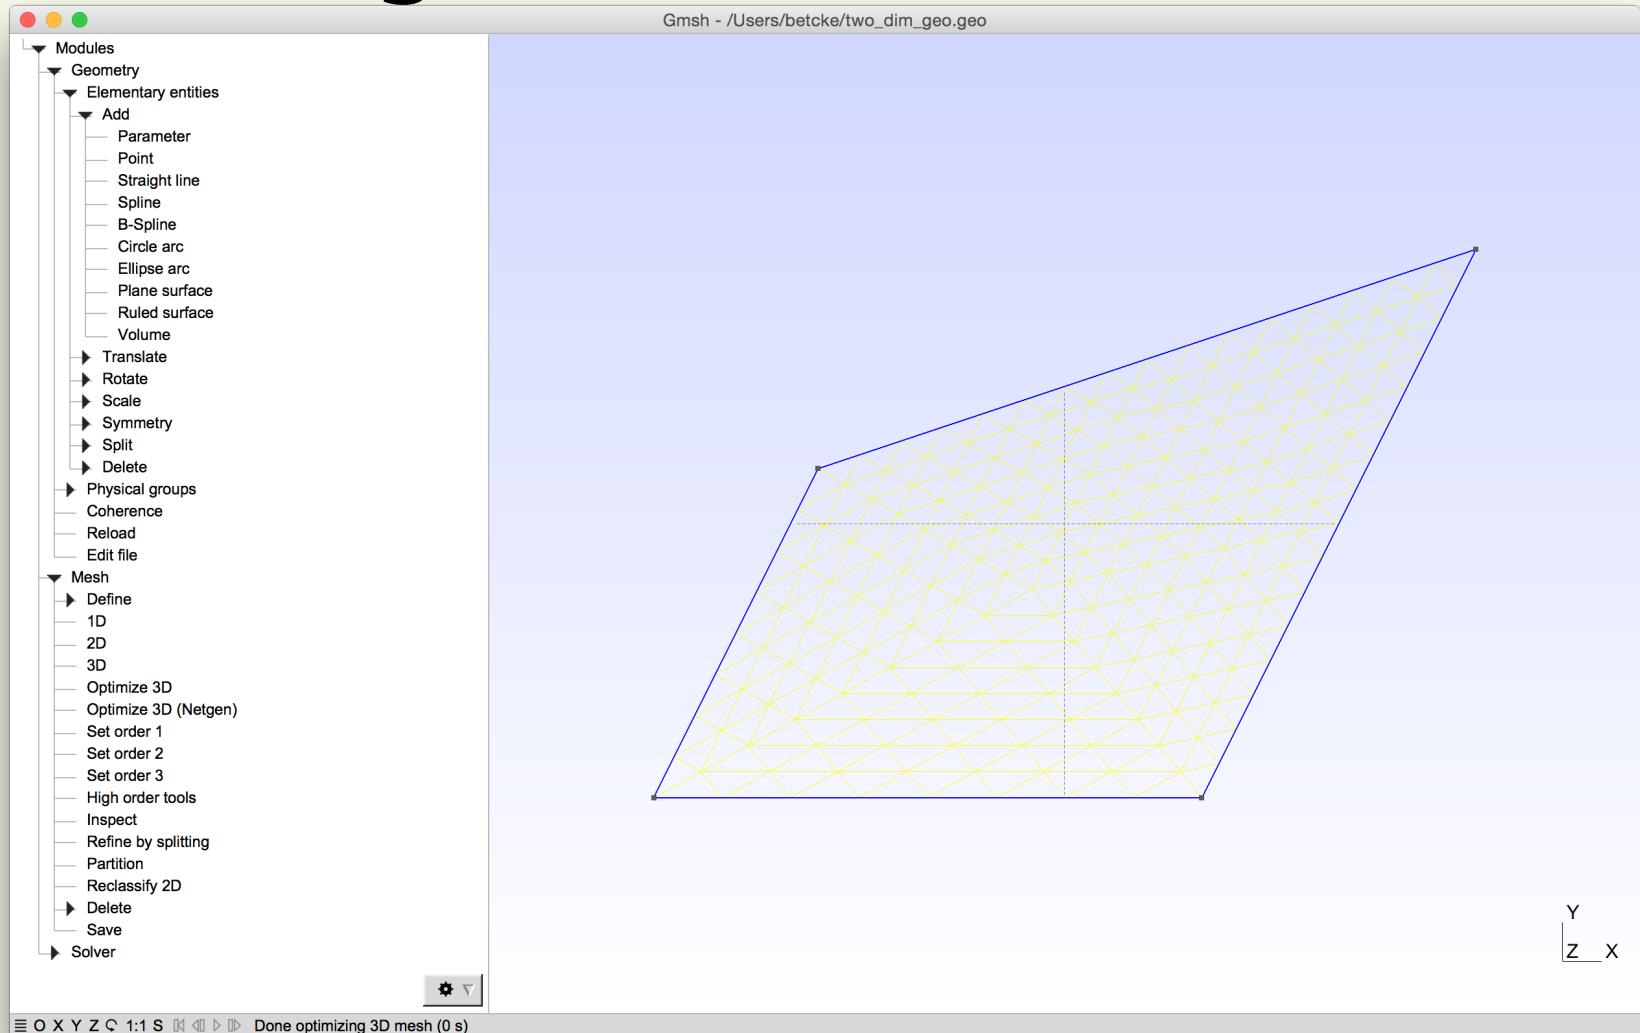
(a) Scan triangulation.



(b) Delaunay triangulation.

- Every Delaunay triangulation maximises the minimum angle over all possible triangulations.

# Triangulation with Gmsh



<http://gmsh.info>

# Triangulation in 3d

- Can extend Delaunay triangulation to tetrahedra (now working with spheres instead of circles)
- Good overview contained in:
- M. Bern, D. Eppstein, *Mesh generation and optimal triangulation*, 1995,
- <https://www.ics.uci.edu/~eppstein/pubs/BerEpp-CEG-95.pdf>
- Gmsh is great tool for triangulation in three dimensions.

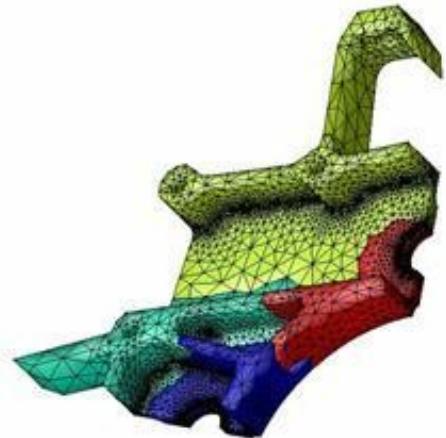
# Creation of regular grids

```
1 import vtk
2 import numpy as np
3
4 grid = vtk.vtkStructuredGrid()
5 points = vtk.vtkPoints()
6 writer = vtk.vtkStructuredGridWriter()
7
8 dims = [10, 20, 30]
9 data   = np.mgrid[0:1:dims[0] * 1j,
10                  0:1:dims[1] * 1j,
11                  0:1:dims[2] * 1j]
12
13 grid.SetDimensions(dims)
14
15 for k in range(dims[2]):
16     for j in range(dims[1]):
17         for i in range(dims[0]):
18             points.InsertNextPoint(
19                 data[0, i, j, k], data[1, i, j, k], data[2, i, j, k])
20
21
22 grid.SetPoints(points)
23
24 writer.SetFileName('structured_grid.vtk')
25 writer.SetInputData(grid)
26 writer.Update()
27
```

- Can create regular grids directly with VTK
- Applications in finite differences and finite volume methods (see examples towards the end of term)

# Mesh partitioning

- Large-Scale Grids do not fit onto a single compute node
- Need to distribute grids across multiple node
- Want to minimize the interfaces between partitions to avoid communication



Consider equivalent graph partitioning problem. Each element corresponds to a vertex. If two elements share a common edge (or face in 3d) then the corresponding vertices in the graph are connected by an edge.

# Graph partitioning problem

Definition (Graph): A graph is a tuple  
 $(V, E)$

consisting of a set of vertices  $V$  and a  
set  $E$  of edges  $(v_i, v_j)$  between  
vertices.

Given a subset  $S$  of the vertices  $V$ .

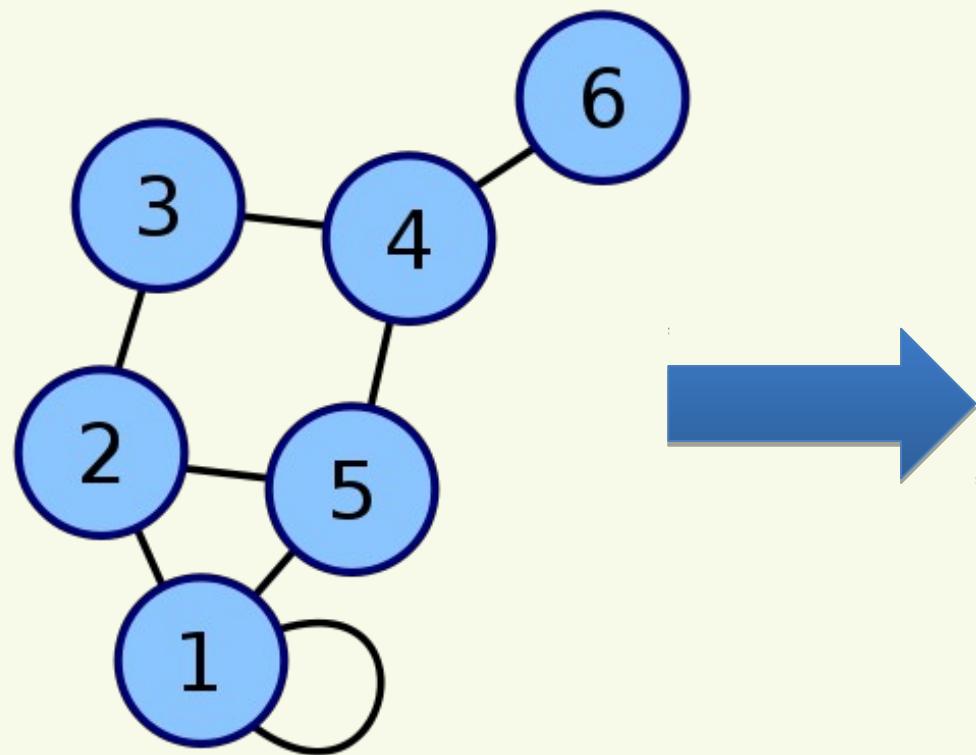
Define

$$\partial S = \{(i, j) \in E : i \in S \wedge j \notin S\}$$

Minimum balanced cut:

$$\rho(S) = \min_S \frac{|\partial S|}{|S|}$$

# Graph Adjacency matrix



Various conventions  
how to deal with loops

$$\begin{pmatrix} 2 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 \end{pmatrix}$$

# Graph Laplacian

Definition (Graph Laplacian):

$$L = D - A$$

Incidence Matrix      Adjacency Matrix

$$D_{ij} = \begin{cases} \sum_{\ell} A_{i\ell}, & i = j \\ 0, & i \neq j \end{cases}$$

Properties:

$L$  is symmetric and positive semidefinite with

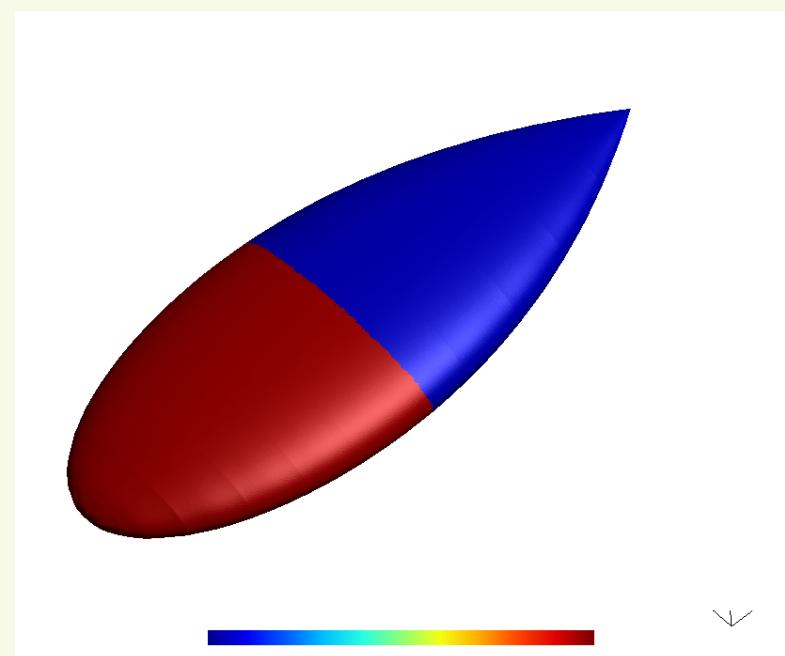
$$0 = \lambda_1 < \lambda_2 \leq \lambda_3 \leq \dots$$

- if  $L$  is connected.
- The number of zero eigenvalues equals the number of connected components
- We can use the first nonzero eigenvalue and corresponding eigenvector for graph partitioning

# Computing a 2-Partition

- Assemble the graph Laplacian  $L$
- Compute the first nonzero eigenvalue of  $L$  and the corresponding eigenvector  $v$ .
- The positive entries of  $v$  correspond to one subdomain and the negative entries to the other subdomain.
- The eigenvector  $v$  is called Fiedler vector

Partitioning of a surface mesh arising from the Fiedler vector.



# FEniCS Example

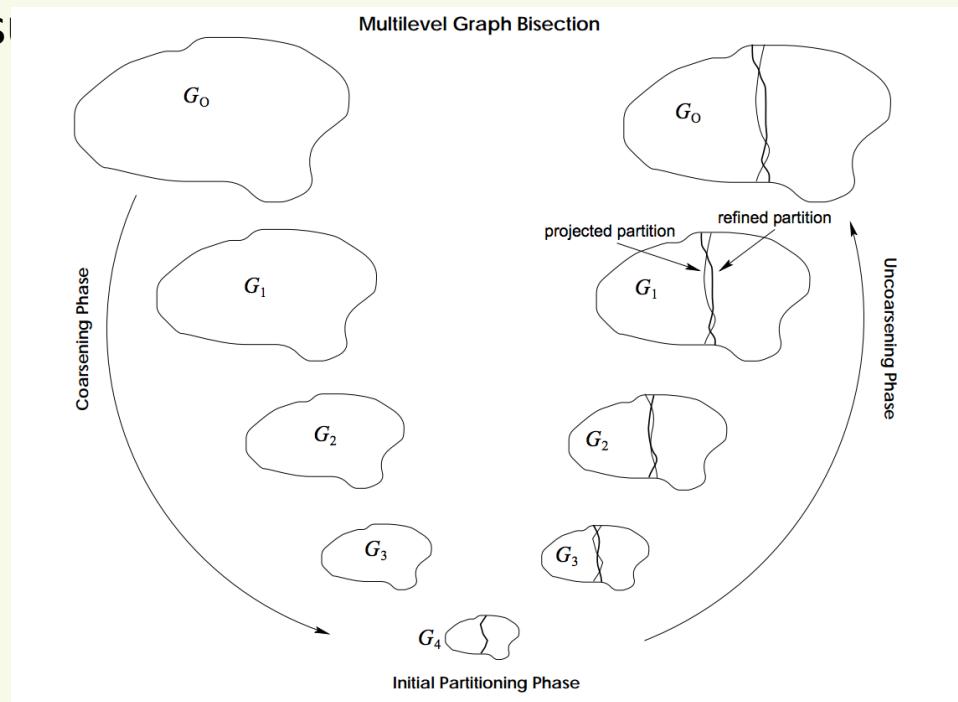
```
1 import numpy as np
2 import scipy.sparse
3 import scipy.sparse.linalg
4
5 from dolfin import *
6 from matplotlib import pyplot as plt
7
8 # Create the mesh
9 mesh = UnitCubeMesh(10, 10, 10)
10
11 # Export the cells
12 cells = mesh.cells()
13
14 # Get the element to vertices map
15 indices = cells.flatten()
16 indptr = 4 * np.arange(cells.shape[0] + 1)
17 e2v = scipy.sparse.csr_matrix((np.ones(len(indices)), indices, indptr))
18
19 # Now compute the adjacency matrix
20 items = (e2v.dot(e2v.T)).todok().items()
21 items = [item for item in items if (2.5 < item[1] and item[1] < 3.5) ]
22 row_ind = [item[0][0] for item in items]
23 col_ind = [item[0][1] for item in items]
24 adjacency = scipy.sparse.coo_matrix((np.ones(len(row_ind)), (row_ind, col_ind))).tocsr()
25
26 # Compute the Graph Laplacian
27 e = np.ones(adjacency.shape[0])
28 d = adjacency.dot(e)
29 D = scipy.sparse.spdiags(d, 0, adjacency.shape[0], adjacency.shape[0])
30 L = D - adjacency
31
32 # Compute the Fiedler vector
33 lam, vec = scipy.sparse.linalg.eigs(L, 2, sigma=0)
34 fiedler = np.real(vec[:, 1])
35
36 # Visualize the resulting partitioning
37 V = FunctionSpace(mesh, "DP", 0)
38 u = Function(V)
39 u.vector()[:] = np.ascontiguousarray(np.sign(fiedler))
40 f = File('sol.pvd')
41 f << u
```

# Heuristic explanation

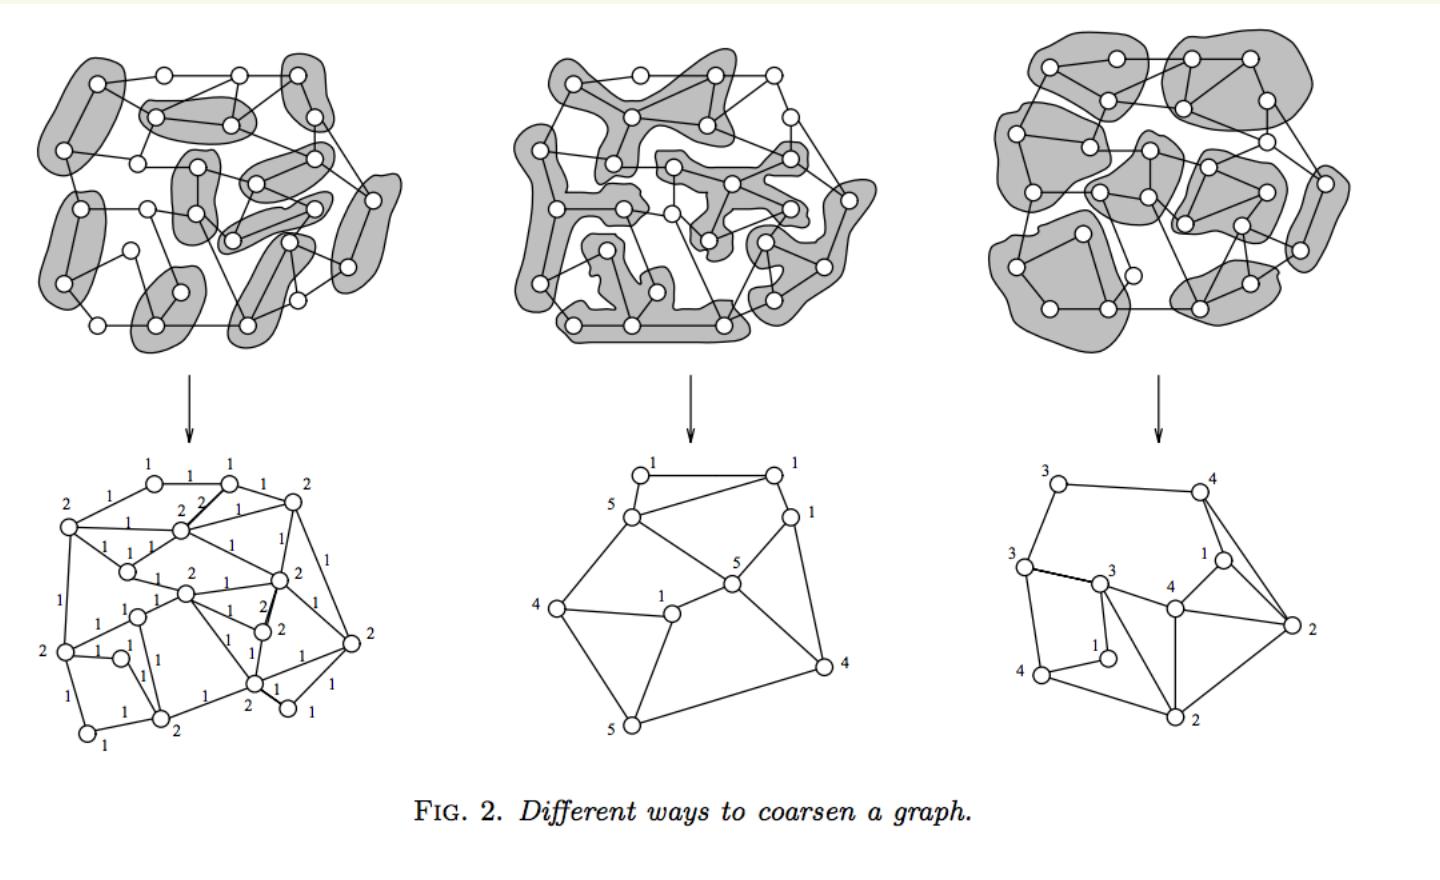
- The graph Laplacian acts like a discrete Laplace operator
- The eigenfunction of the first nonzero eigenvalue has one nodal line that separates a domain into a positive and a negative part
- The smaller the first nonzero eigenvalue the less connected a graph
- For more than one partition consider the first  $k$  eigenvectors.

# Multilevel Partitioning

- Computing the eigenvectors corresponding to a large graph is expensive.
- Idea: Use a multilevel technique by first coarsening the graph and computing the partitioning on a small s



# Multilevel Partitioning...



# Graph partitioning software

- METIS (<http://glaros.dtc.umn.edu/gkhome/metis/metis/overview>),  
pymetis available (<https://mathematician.de/software/pymetis/>)
- 
- Scotch (<https://www.labri.fr/perso/pelegrin/scotch/>)
- 
- Chaco  
(<http://www3.cs.stonybrook.edu/~algorith/implement/chaco/implment.shtml>)

# Grid Management Software

- Dune-Grid ([www.dune-project.org](http://www.dune-project.org))
  - Extensive package for grid based algorithms
  - Supports various grid managers, large-scale parallel grids, refinement, etc.
- STK (<https://trilinos.org/packages/stk/>)
  - Grid manager of Trilinos
  - Large-Scale Parallel grids
- FEniCS ([www.fenicsproject.org](http://www.fenicsproject.org))
  - High-quality grid manager as part of the FEniCS distribution
  - Well integrated into FEniCS

# Sparse Matrices

```
from dolfin import *
mesh = UnitCubeMesh(100, 100,
100)
print(mesh.num_vertices())
```

Creates mesh with over 1m vertices.

- Storage of the resulting matrix in dense format would take around 8 TeraBytes of RAM.
- However, almost all entries are zero.
- We should only store nonzero entries.
- Algebra to deal with matrices in sparse storage format required
- Software Support: Petsc, Trilinos (Epetra/Tpetra), Scipy
- In the following we will use Y. Saad, Iterative Methods for Sparse Linear Systems, 2<sup>nd</sup> Edition (freely available for download)

# Sparse Formats

$$\begin{bmatrix} 1 & 0 & 0 & 2 & 0 \\ 3 & 4 & 0 & 5 & 0 \\ 6 & 0 & 7 & 8 & 9 \\ 0 & 0 & 10 & 11 & 0 \\ 0 & 0 & 0 & 0 & 12 \end{bmatrix}$$

1	2	3	4	5	6	7	8	9	10	11	12
0	0	1	1	1	2	2	2	2	3	3	4
0	3	0	1	3	0	2	3	4	2	3	4

data

## Row indices

# Column indices

- Simple coordinate list
  - Easy to insert elements
  - Not memory efficient

# Compressed Sparse Row Format (CSR)

$$\begin{bmatrix} 1 & 0 & 0 & 2 & 0 \\ 3 & 4 & 0 & 5 & 0 \\ 6 & 0 & 7 & 8 & 9 \\ 0 & 0 & 10 & 11 & 0 \\ 0 & 0 & 0 & 0 & 12 \end{bmatrix}$$

1	2	3	4	5	6	7	8	9	10	11	12
0	3	0	1	3	0	2	3	4	2	3	4

0	2	5	9	11	12
---	---	---	---	----	----

data

Column indices

Index pointers

Last element contains number of nonzeros in matrix

- The nonzero col indices in the  $j$ th row are contained in  $\text{JA}[\text{IA}[j]:\text{IA}[j+1] - 1]$
- Variant is compressed sparse column format (CSC)

# Matrix-Vector product

```
def matvec(x, data, indices, indptr):
    """Matvec of a matrix in CSR format with
vector x"""
    y = numpy.zeros(x.shape[0],
dtype='float64')
    n = len(indptr) - 1

    for i in range(n):
        y[i] = data[ indptr[i]: indptr[i+1] ].dot(
            x[indices[indptr[i]: indptr[i+1]]])
```

- Very easy to parallelize by distributing matrix rows among the processes
- Efficient vectorized operations

# LU Decomposition

Want to  
solve

$$Ax = b$$

for a general  $n \times n$  matrix A and an arbitrary right-hand side vector b.

Idea: Represent A as a product of two simpler matrices.

$$A = LU$$

and solve

$$LUx = b$$

instead by first solving  $Ly = b$  and then  $Ux = y$

.

How do we get L and U?

# LU Decomposition...

Reviewing Gaussian Elimination...

Consider the system  $Ax = b$ . We write it in the form

$$\left[ \begin{array}{cccc|c} a_{11} & a_{12} & \dots & a_{1n} & b_1 \\ a_{21} & a_{22} & \dots & a_{2n} & b_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} & b_n \end{array} \right]$$

Now subtract  $\ell_{21} := a_{21}/a_{11}$  times the first row from the second row to obtain

$$\left[ \begin{array}{cccc|c} a_{11} & a_{12} & \dots & a_{1n} & b_1 \\ 0 & a_{22} - \ell_{21}a_{12} & \dots & a_{2n} - \ell_{21}a_{1n} & b_2 - \ell_{21}b_1 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ a_{n1} & a_{n2} & \dots & a_{nn} & b_n \end{array} \right]$$

# LU Decomposition...

Continuing this way we obtain an upper triangular matrix U of the form

$$\left[ \begin{array}{cccc|c} * & * & \dots & * & b_1 \\ 0 & * & \dots & * & * \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ * & * & \dots & * & * \end{array} \right].$$

The trick is to represent the subtraction of a multiple of the  $i$ th row from the  $j$ th row as a matrix vector product of the form

$$L_{ij}A := [I - \ell_{ij}e_i e_j^T] A$$

# LU Decomposition...

We obtain therefore an upper triangular matrix by the sequence of operations

$$\begin{aligned} U &:= L_{n,n-1} \dots \\ &\dots L_{42} L_{32} L_{n1} \dots L_{31} L_{21} A \end{aligned}$$

The matrix L is therefore obtained as

$$\begin{aligned} L &:= (L_{n,n-1} \dots L_{42} L_{32} L_{n1} \dots L_{31} L_{21})^{-1} \\ &= L_{21}^{-1} L_{31}^{-1} \dots L_{n1}^{-1} L_{32}^{-1} L_{42}^{-1} \dots L_{n,n-1}^{-1}. \end{aligned}$$

A direct computation shows that

$$\begin{aligned} L &= L_{21}^{-1} L_{31}^{-1} \dots L_{n1}^{-1} L_{32}^{-1} L_{42}^{-1} \dots L_{n,n-1}^{-1} \\ &= (I + \ell_{21} e_2 e_1^T) (I + \ell_{31} e_3 e_1^T) \dots \end{aligned}$$

$$= \begin{bmatrix} 1 & 0 & \dots & 0 \\ \ell_{21} & 1 & \dots & 0 \\ \dots & \dots & \ddots & \vdots \\ \ell_{n1} & \ell_{n2} & \dots & 1 \end{bmatrix}$$

# LU Decomposition...

We can now solve with L using simple forward substitution

$$Ly = \begin{bmatrix} 1 & 0 & \dots & 0 \\ \ell_{21} & 1 & \dots & 0 \\ \dots & \dots & \ddots & \vdots \\ \ell_{n1} & \ell_{n2} & \dots & 1 \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{bmatrix}$$

$$y_i = b_i - \sum_{j=1}^{i-1} \ell_{ij} y_j; \quad i = 1, \dots, n$$

Correspondingly can solve for U by using backward substitution.

# Remarks

- LU decomposition has a complexity of  $O(n^3)$  for general matrices of dimension  $n$
- In practice, we use use pivoting strategy to reorder the rows during the computation to improve the stability of the LU decomposition.
- LU has a wonderful property. For certain sparse matrices its running time reduces from  $O(n^3)$  to  $O(n)$ .

# Cholesky Decomposition

A matrix  $A$  is called symmetric positive definite (SPD) if  $A$  is symmetric and

$$x^T A x > 0$$

- for all nonzero  $x$ .
- SPD matrices are very common. Essentially, they can be interpreted as energy functionals. Solving linear systems with SPD matrices amounts to finding minimum energy states of a system.
- For SPD matrices a special variant of Gaussian Elimination is possible that adheres to the structure of SPD matrices, namely we can decompose  $A$  into

$$A = C^T C$$

with  $C$  upper triangular. This can be achieved in half the operations as required for the standard LU decomposition.

# LU for sparse matrices

A matrix is called banded if

$$a_{ij} \neq 0 \text{ if and only if } i - m_L \leq j \leq i + m_U$$

The number

$$m_L + m_U + 1$$

is called the bandwidth of A.

Thm: The lower bandwidth of L in the LU decomposition  $A = LU$  is  $m_L$  and the upper bandwidth in U is  $m_U$ .

Thm: Let A be an  $n \times n$  matrix. If the bandwidth is constant  $O(1)$  for growing  $n$ . Then the complexity of the LU decomposition is  $O(n)$ .

# Reorderings

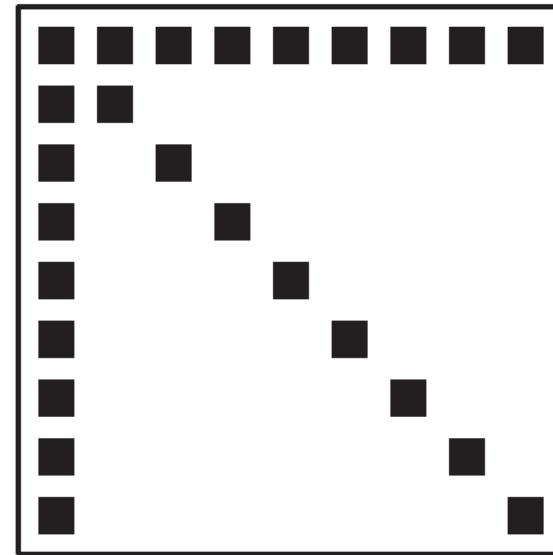
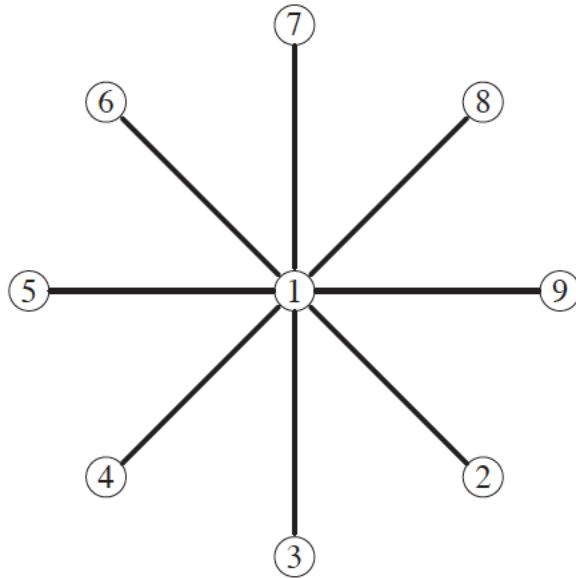


Figure 3.4: Pattern of a  $9 \times 9$  arrow matrix and its adjacency graph.

Performing Gaussian Elimination fills up the matrix.

# Reorderings...

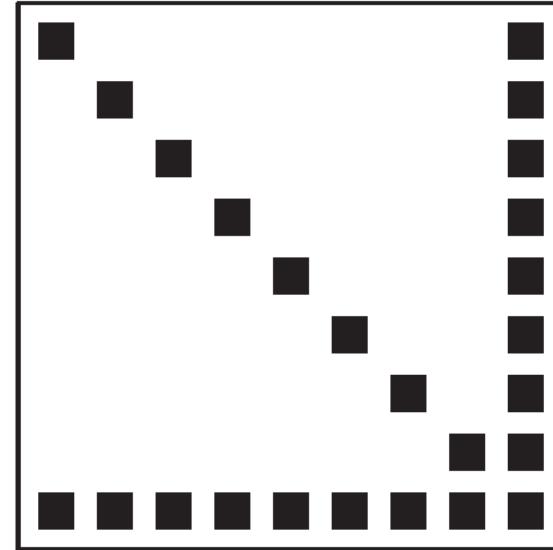
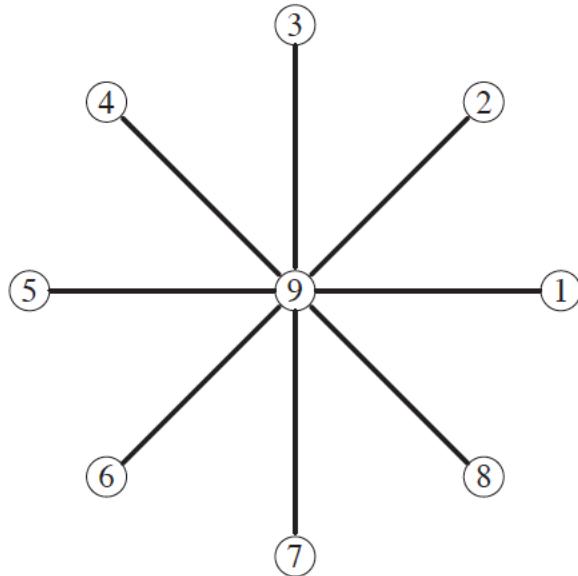


Figure 3.5: Adjacency graph and matrix obtained from above figure after permuting the nodes in reverse order.

A simple reordering produces no fill-in.

# Permutations

Definition: Let  $A$  be a matrix and  $\pi = \{i_1, i_2, \dots, i_n\}$  a given permutation of the set  $\{1, 2, \dots, n\}$ . Then

$$A_{\pi,*} = \{a_{\pi(i),j}\}_{i=1,\dots,n; j=1,\dots,m}$$
$$A_{*,\pi} = \{a_{i,\pi(j)}\}_{i=1,\dots,n; j=1,\dots,m}$$

are called row, respectively, column perturbation of  $A$ .

$$P_\pi = I_{\pi,*}, \quad P_\pi^T = I_{*,\pi} \quad A_{\pi,*} = I_{\pi,*}A = P_\pi A \quad A_{*,\pi} = AI_{*,\pi} = AP_\pi^T$$

# Permutations

**Example 3.1.** Consider, for example, the linear system  $Ax = b$  where

$$A = \begin{pmatrix} a_{11} & 0 & a_{13} & 0 \\ 0 & a_{22} & a_{23} & a_{24} \\ a_{31} & a_{32} & a_{33} & 0 \\ 0 & a_{42} & 0 & a_{44} \end{pmatrix}$$

and  $\pi = \{1, 3, 2, 4\}$ , then the (column-) permuted linear system is

$$\begin{pmatrix} a_{11} & a_{13} & 0 & 0 \\ 0 & a_{23} & a_{22} & a_{24} \\ a_{31} & a_{33} & a_{32} & 0 \\ 0 & 0 & a_{42} & a_{44} \end{pmatrix} \begin{pmatrix} x_1 \\ x_3 \\ x_2 \\ x_4 \end{pmatrix} = \begin{pmatrix} b_1 \\ b_2 \\ b_3 \\ b_4 \end{pmatrix}.$$

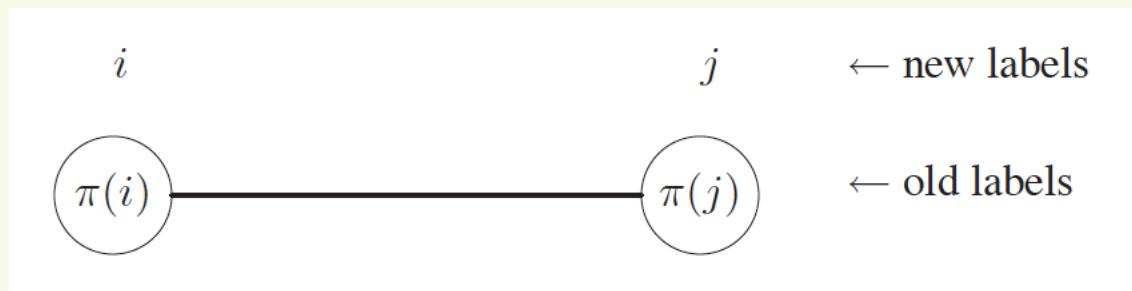
# Symmetric permutation

A symmetric permutation of a matrix  $A$  is defined as

$$A' = A_{\pi, \pi}$$

We have  $A'_{ij} = a_{\pi(i), \pi(j)}$ .

Symmetric permutations do not change the adjacency graph. They just relabel the vertices.



# Reducing bandwidth by Level-Set Reordering

- Reorder a graph by traversing level-sets
- A level set is defined recursively as the set of all neighbors of the previous level-set
- We start with a single node

## Reordering by Breadth-First Search

1.     *Initialize  $S = \{v\}$ ,  $seen = 1$ ,  $\pi(seen) = v$ ; Mark  $v$ ;*
2.     *While  $seen < n$  Do*
3.          $S_{new} = \emptyset$ ;
4.         *For each node  $v$  in  $S$  do*
5.             *For each unmarked  $w$  in  $adj(v)$  do*
6.                 *Add  $w$  to  $S_{new}$ ;*
7.                 *Mark  $w$ ;*
8.                  $\pi(plusplus seen) = w$ ;
9.             *EndDo*
10.           $S := S_{new}$
11.         *EndDo*
12.         *EndWhile*

In what order  
should be traverse  
the adjacent  
vertices ?

# Cuthill-McKee Algorithm

0. *Find an initial node  $v$  for the traversal*
1. *Initialize  $S = \{v\}$ ,  $seen = 1$ ,  $\pi(seen) = v$ ; Mark  $v$ ;*
2. *While  $seen < n$  Do*
3.    $S_{new} = \emptyset$ ;
4.   *For each node  $v$  Do:*
5.      $\pi( + + seen) = v$ ;
6.     *For each unmarked  $w$  in  $adj(v)$ , going from lowest to highest degree Do:*
7.       *Add  $w$  to  $S_{new}$ ;*
8.       *Mark  $w$ ;*
9.     *EndDo*
10.     $S := S_{new}$
11.    *EndDo*
12.   *EndWhile*

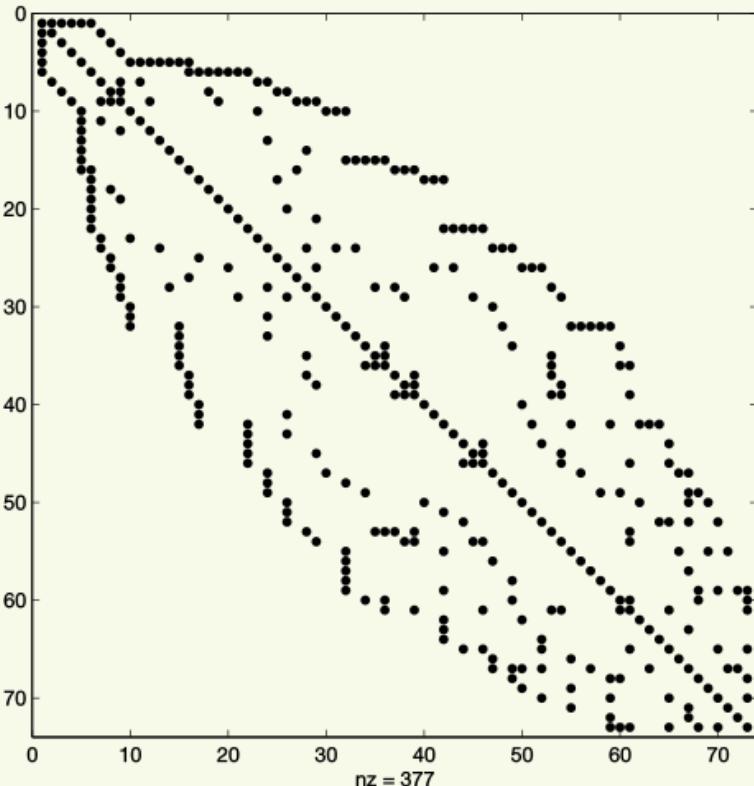


Traverse from  
lowest to highest  
degree

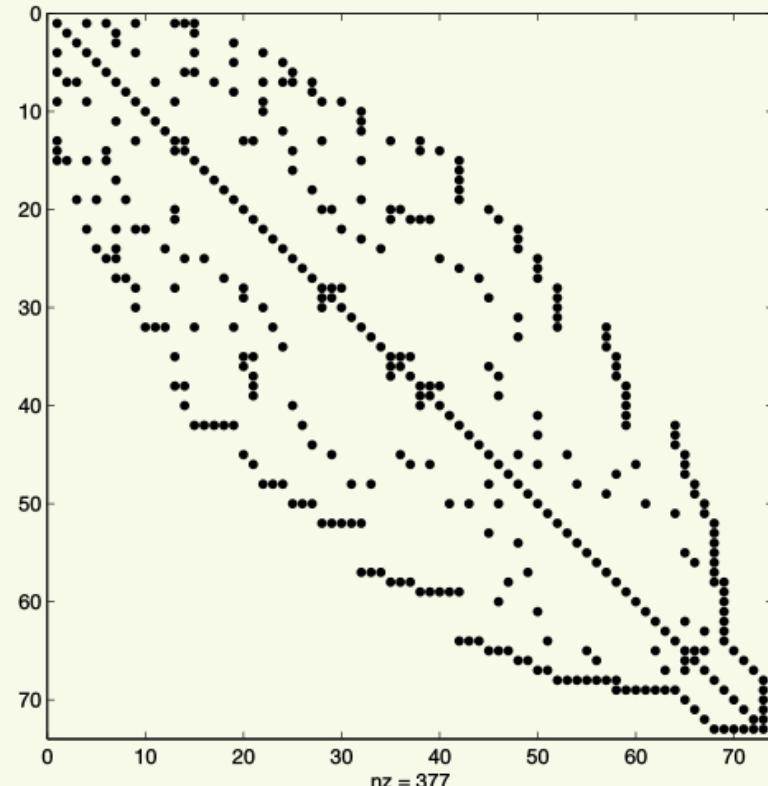
Reverse Cuthill-McKee:

Use Cuthill-McKee and then reverse the ordering. Often reduced the fill-in better than Cuthill-McKee.

# Cuthill-Mckee and Reverse Cuthill McKee



Cuthill-McKee



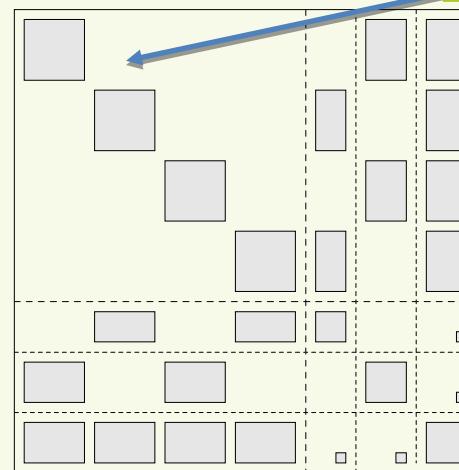
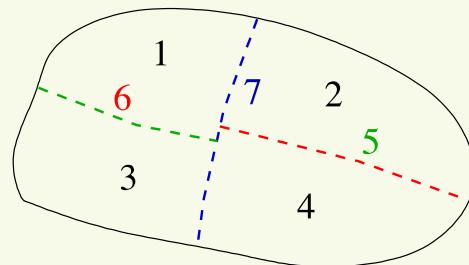
Reverse Cuthill-McKee

By Kxx - Own work, CC BY-SA 3.0,  
<https://commons.wikimedia.org/w/index.php?curid=32481057>

# Nested Dissection

- Reordering for SPD matrices that allows the parallelization of the Cholesky decomposition.
- Idea: Split a graph into two separated subgraphs and a separator layer that connects the subgraphs. Recursively continue for each subgraph.
- Label nodes in the separator after nodes in the subgraphs.

Can perform Cholesky decomposition independently in each diagonal block



# Shopping for matrices in the Matrix Market

- <http://math.nist.gov/MatrixMarket/>
- Huge collection of sparse matrices with search features for a wide range of characteristics
- Ideal to test your own algorithms
- Scipy.io can directly read matrix market format. Routines for Matlab/C/Fortran can be downloaded

# Sparse matrix packages

- Matlab has excellent built-in sparse matrix format, using UMFPack as sparse direct solver
- Scipy has decent sparse-matrix support (good enough for many things, not as extensive as Matlab). Default solver is SuperLU
- Intel MKL can deal with sparse matrices and provides a built-in version of Pardiso, a non-free high-performance solver
- PETSc/Trilinos have outstanding parallel sparse matrix support

# Scipy sparse matrix formats

- `bsr_matrix`: A blockwise sparse format for sparse matrices with dense subblocks. These often appear in FEM with high-order basis functions
- `coo_matrix`: Sparse matrix in triplet format ( $i, j, \text{data}$ ). Multiple entries are summed up when converting to CSR format, making it ideal for FEM assembly
- `csc_matrix`: Compressed sparse column format
- `csr_matrix`: Compressed sparse row format
- `dia_matrix`: Ideal for matrices where data is given along diagonals
- `dok_matrix`: Dictionary of keys. Allows  $O(1)$  assignment operations and then conversion to coo format (note: Does not allow creation of duplicate entries).
- `lil_matrix`: Linked-list format for row-wise construction.

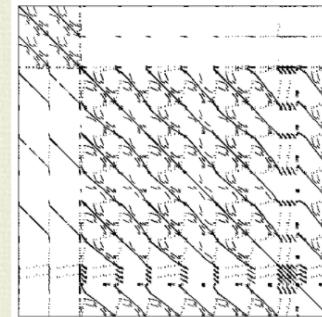
# Sparse Direct Solvers - Mumps

- MPI Parallel Multifrontal Solver
- Widely used in applications
- Open-Source License
- <http://mumps.enseeiht.fr/>

## MUMPS : A PARALLEL SPARSE DIRECT SOLVER

### MUMPS MAIN FEATURES

- Solution of large linear systems with  
**symmetric positive definite** matrices  
**general symmetric** matrices  
**general unsymmetric** matrices
- Version for **complex arithmetic**
- **Parallel** factorization and solve phases  
(uniprocessor version also available)
- **Iterative refinement** and **backward error analysis**
- Various matrix input formats  
**assembled, distributed, elemental** format
- **Partial factorization** and **Schur complement matrix** (centralized or 2D block-cyclic)
- **Interfaces to MUMPS:** Fortran, C, Matlab and Scilab
- Several **reorderings** interfaced: AMD, QMAD, AMF, PORD, METIS, PARMETIS, SCOTCH, PT-SCOTCH



# Sparse Direct Solvers - SuiteSparse

- Multi-Threaded parallel multifrontal solvers
- Integrated into Matlab
- Widely used

**SuiteSparse is a suite of sparse matrix algorithms, including:**

- UMFPACK: multifrontal LU factorization. Appears as LU and  $x=A\backslash b$  in MATLAB.
- CHOLMOD: supernodal Cholesky. Appears as CHOL and  $x=A\backslash b$  in MATLAB. Used in Google Ceres. Now with CUDA acceleration, in collaboration with NVIDIA.
- SPQR: multifrontal QR. Appears as QR and  $x=A\backslash b$  in MATLAB. CUDA acceleration just released (October 10, 2014, SuiteSparse 4.4.0), and submitted to ACM Trans. Math. Software
- KLU and BTF: sparse LU factorization, well-suited for circuit simulation. Appears in Xyce by Sandia, and many commercial circuit simulation packages.
- Ordering methods (AMD, CAMD, COLAMD, and CCOLAMD). AMD and COLAMD appear in MATLAB.
- CSparse and CXSparse: a concise sparse Cholesky factorization package for my SIAM book.
- UFget: a MATLAB interface for the UF Sparse Matrix Collection
- spqr\_rank: a MATLAB package for reliable sparse rank detection, null set bases, pseudoinverse solutions, and basic solutions.
- Factorize: an object-oriented solver for MATLAB (a reusable backslash).
- SSMULT and SFMULT: sparse matrix multiplication. Appears as  $C=A^*B$  in MATLAB.
- ... and many other packages.

# Sparse direct solvers - SuperLU

- General purpose sparse direct multi-threaded and MPI enabled solver
- Single-Core version default sparse solver in Scipy

SuperLU is a general purpose library for the direct solution of large, sparse, nonsymmetric systems of linear equations. The library is written in C and is callable from either C or Fortran program. It uses MPI, OpenMP and CUDA to support various forms of parallelism. It supports both real and complex datatypes, both single and double precision, and 64-bit integer indexing. The library routines performs an LU decomposition with partial pivoting and triangular system solves through forward and back substitution. The LU factorization routines can handle non-square matrices but the triangular solves are performed only for square matrices. The matrix columns may be preordered (before factorization) either through library or user supplied routines. This preordering for sparsity is completely separate from the factorization. Working precision iterative refinement subroutines are provided for improved backward stability. Routines are also provided to equilibrate the system, estimate the condition number, calculate the relative backward error, and estimate error bounds for the refined solutions.

Serial SuperLU package also contains ILU routines, using numerical threshold-based dropping, with partial pivoting (ILUTP).

SuperLU package comes in three different flavors:

- [SuperLU](#) for sequential machines ([code documentation](#) in HTML)
- [SuperLU\\_MT](#) for shared memory parallel machines
- [SuperLU\\_DIST](#) for distributed memory ([code documentation](#) in HTML)

The target machines for SuperLU\_DIST are the highly parallel distributed memory hybrid systems. The numerical factorization routines are already implemented for hybrid systems with multiple GPUs. Further work will be needed to implement the other phases of the algorithms on the hybrid systems and to enhance strong scaling to extreme scale.

# Sparse direct solvers - Pardiso

- Highly efficient sparse direct multithreaded and MPI solver
- Two versions:
  - MKL Pardiso (part of the Intel MKL Library)
  - Pardiso (<http://www.pardiso-project.org>)
- Pardiso is not free software
- Academic licenses from Pardiso
- available for free
- MKL Pardiso less advanced than
- Pardiso



# Iterative Solvers

- Sparse direct solvers become infeasibly expensive for big problems with complicated connectivity that creates a large amount of fill-in for Gaussian elimination
- Idea: Develop robust iterative solvers that compute not the exact solution, but an approximate solution to a certain given accuracy.
- Modern developments: Combine sparse direct and iterative solvers (more later)

# Splitting Methods

- Splitting Methods are among the earliest and most simple iterative methods.
- By themselves not efficient, but very good as preconditioners or as solvers within a multigrid framework.

We want to solve

$$Ax = b$$

Idea: Split A into two matrices M and N such that

$$A = M - N$$

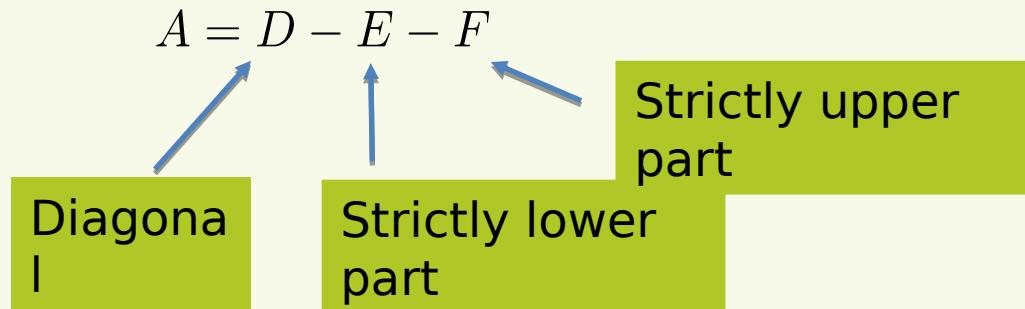
 
$$x = M^{-1}b + M^{-1}Nx$$

We consider this equation as the fixed point iteration

$$x_{k+1} = M^{-1}b + M^{-1}N x_k$$

# Splitting Methods...

Let



Jacobi

:

$$x_{k+1} = D^{-1}(E + F)x_k + D^{-1}b$$

Gauss-Seidel:

$$x_{k+1} = (D - E)^{-1}Fx_k + (D - E)^{-1}b$$

SOR:

$$x_{k+1} = (D - \omega E)^{-1} [\omega F + (1 - \omega)D] x_k + \omega(D - \omega E)^{-1}b$$

SOR Splitting;

$$\omega A = (D - \omega E) - (\omega F + (1 - \omega)D)$$

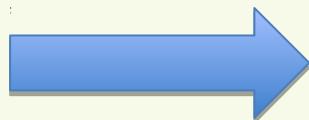
# Convergence of splitting methods

General iteration scheme:

$$x_{k+1} = Gx_k + f$$

Exact solution satisfies

$$x^* = Gx^* + f$$



$$x_{k+1} - x^* = G(x_k - x^*) = G^k(x_0 - x^*)$$

Take norms to obtain

$$\|x_{k+1} - x^*\| \leq \|G\|^k \|(x_0 - x^*)\|$$

Thm: The scheme converges if  $\|G\| < 1$

Since we did not specify the norm a more intrinsic argumentation shows that it is enough for the largest eigenvalue of  $G$  to be of magnitude smaller than 1.

# Convergence of splitting methods...

A splitting method is a special case of this iteration scheme with  $G = M^{-1}N$ .

A matrix  $A$  is **strictly** diagonally dominant if

$$|a_{jj}| > \sum_{i=1, i \neq j}^n |a_{ij}|, \quad j = 1, \dots, n.$$

It is called **irreducibly diagonally dominant** if  $A$  is irreducible and

$$|a_{jj}| \geq \sum_{i=1, i \neq j}^{i=1} |a_{ij}|, \quad j = 1, \dots, n$$

with strict inequality for at least one  $j$ .

Thm: If  $A$  is strictly diagonally dominant or irreducibly diagonally dominant

Jacobi and Gauss-Seidel converge for any given starting vector.

# Summary

- Splitting methods not competitive by themselves.
- However, they can be used as preconditioners (more later)...
- ...or as smoothers in multigrid methods (again, later).
- Even if we don't use splitting methods directly any more they still play an important role.

# Krylov Subspace Methods

Motivation: Consider the simple iteration

$$x_{k+1} = (I - A)x_k + b$$

If the iteration converges, it converges to the solution of  $Ax = b$

We have

$$x_1 = b$$

$$x_2 = (I - A)b + b = 2b - Ab$$

$$x_3 = (I - A)x_2 + b = 3b - 3Ab + A^2b$$

$$x_k = P_{k-1}(A)b$$

$$P_{k-1}(A) = \alpha_0 b + \alpha_1 Ab + \cdots + \alpha_{k-1} A^{k-1} b$$

is a polynomial in  $A$  with coefficients  $\alpha_0, \dots, \alpha_{k-1}$ .

Ideally, we would choose the coefficients such that

$P_{k-1}(A)b \approx A^{-1}b$  in some way.

# Krylov Subspace Methods

The  $k$ th **Krylov subspace**  $\mathcal{K}_k(A, b)$  is defined as

$$\mathcal{K}_k(A, b) = \text{span}\{b, Ab, \dots, A^{k-1}b\}.$$

By

$$K_k = [b \quad Ab \quad A^2b \quad \dots \quad A^{k-1}b]$$

we denote the corresponding **Krylov matrix**.

How do we choose  $x_k = P_{k-1}(A)b$  from the  $k$ -th Krylov subspace? Define  $r_k = b - Ax_k$  to be the corresponding residual. We define the following Krylov subspace methods by the condition on the residual.

- Conjugate Gradients: The residual must be orthogonal to the  $k$ -th Krylov subspace.
- GMRES/MINRES: The norm of the residual is minimized over the  $k$ -th Krylov subspace.
- BICG: The residual is orthogonal to the Krylov subspace associated with  $A^\top$ .

# Arnoldi method

We want to create an orthogonal basis for the k-th Krylov subspace.

$$q_1 = b / \|b\|$$

for  $j = 1, \dots, k - 1$

Multiply to get next vector

$$t = A q_j$$

for  $i = 1, \dots, j$

$$h_{i,j} = q_i^T t$$

$$t = t - h_{i,j} q_i$$

$$h_{j+1,j} = \|t\|$$

$$q_{j+1} = t / h_{j+1,j}$$

Orthogonalize against  
previous vectors

end

Normalize

$$A [q_1 \ \dots \ q_j] = [q_1 \ \dots \ q_j \ q_{j+1}] \begin{bmatrix} h_{1,1} & h_{1,2} & \dots & h_{1,j} \\ h_{2,1} & h_{2,2} & \dots & h_{2,j} \\ 0 & h_{3,2} & \dots & \vdots \\ 0 & 0 & \dots & h_{j+1,j} \end{bmatrix}, \quad j = 1, \dots, k - 1$$

# Arnoldi method...

Rewrite the Arnoldi algorithm as

$$AQ_j = Q_j H_j + h_{j+1,j} q_{j+1} e_j^T$$

$$Q_j = [q_1 \quad \dots \quad q_j], \quad H_j = \begin{bmatrix} h_{1,1} & h_{1,2} & \dots & h_{1,j} \\ h_{2,1} & h_{2,2} & \dots & h_{2,j} \\ 0 & \ddots & \dots & h_{j,j} \end{bmatrix}$$

Orthogonal  
columns

Hessenberg  
matrix

By construction the columns of  $Q_j$  form an orthogonal basis of the  $j$ -th Krylov subspace.

# FOM – Full Orthogonalization Method

Consider the linear system  $Ax = b$  with initial approximation  $x_0$ .

We rewrite the linear system as  $Ay = b - Ax_0 = r_0$

for  $x = x_0 + y$ . We now define the Krylov subspace with starting vector  $r_0$  and write  $x_k = x_0 + Q_k \tilde{y}_k$

We now want to find  $\tilde{y}_k$  such that  $r_k \perp \mathcal{K}_k(A, r_0)$ , or equivalently

$$Q_k^T(b - Ax_k) = Q_k^T(b - Ax_k - AQ_k \tilde{y}_k) = Q_k^T(r_0 - AQ_k \tilde{y}_k) = 0.$$



$$Q_k^T AQ_k \tilde{y}_k = Q_k^T r_0$$

equivalently

$$H_k \tilde{y}_k = h_{11} e_1$$

# FOM – Full Orthogonalization Method

We have the following simple result for the residual

$$\begin{aligned} r_k &= b - Ax_k \\ &= b - Ax_0 - AQ_k \tilde{y}_k \\ &= h_{11}q_1 - Q_k H_k \tilde{y}_k - h_{k+1,k} e_k^T \tilde{y}_k q_{k+1} \\ &= h_{11}q_1 - h_{11}Q_k e_1 - h_{k+1,k} e_k^T \tilde{y}_k q_{k+1} \\ &= -h_{k+1,k} e_k^T \tilde{y}_k q_{k+1}. \end{aligned}$$

It follows that

$$\|r_k\| = |h_{k+1,k}|.$$

In practice, the full orthogonalization method is rarely used. Instead, the standard solver for many problems is GMRES - Generalized Minimum Residuals.

# GMRES

The idea of GMRES is simple:

- Perform an Arnoldi iteration
- In each step of the Arnoldi iteration minimize the residual

$$\|Ax_k - b\|_2 = \|A(x_0 + Q_k \tilde{y}_k - b)\|_2 = \|AQ_k \tilde{y}_k - r_0\|_2.$$

From the structure of the Arnoldi iteration we have equivalently

$$\begin{aligned}\|AQ_k \tilde{y}_k - r_0\|_2 &= \|Q_k H_k \tilde{y}_k - r_0 + h_{k+1,k} q_{k+1} e_k^T \tilde{y}_k\|_2 \\ &= \left\| \begin{bmatrix} H_k \\ h_{k+1,k} e_k^T \end{bmatrix} \tilde{y}_k - \|r_0\|_2 e_1 \right\|_2\end{aligned}$$



In each step minimize this expression.

# GMRES as Approximation Problem

$$\begin{aligned}\|r_k\|_2 &= \|r_0 - A\tilde{Q}\tilde{y}_k\|_2 \\ &= \|r_0 - AP_{k-1}(A)r_0\|_2 \\ &= \|(I - AP_{k-1}(A))r_0\|_2\end{aligned}$$



Polynomial of degree  $k$  such that  $P(0) = 1$ .

GMRES Approximation Problem: Find a polynomial  $P_k$  with  $P(0) = 1$  such that  $\|P(A)r_0\|_2 = \text{minimum}$

# GMRES Convergence

A simple observation:

$$\|r_{k+1}\|_2 \leq \|r_k\|_2$$

The reason is that GMRES minimises the residual over Krylov subspaces of growing dimension  $k$ .

Expressing the residual in terms of polynomials we obtain

$$\frac{\|r_k\|}{\|r_0\|} \leq \inf_{p_k \in \mathcal{P}_k} \|p_k(A)\|$$


Space of all polynomials of maximum degree  $k$  with  $p_n(0) = 1$ .

# GMRES Convergence...

Let  $A$  be diagonalizable with

$$A = V\Lambda V^{-1}.$$

At step  $k$  of the GMRES iteration, the residual  $r_k$  satisfies

$$\frac{\|r_k\|}{\|r_0\|} \leq \inf_{p_k \in \mathcal{P}_k} \|p_k(A)\| \leq \|V\| \cdot \|V^{-1}\| \inf_{p_k \in \mathcal{P}_k} \|p_k\|_{\Lambda(A)}$$

Rule of thumb: If the eigenvalues are bounded away from zero and not too far apart from each other then GMRES will converge quickly.

# GMRES Convergence...

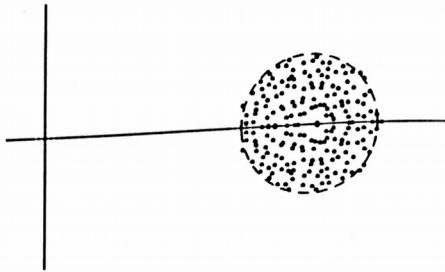


Figure 35.2. Eigenvalues of the  $200 \times 200$  matrix  $A$  of (35.17). The dashed curve is the circle of radius  $1/2$  with center  $z = 2$  in  $\mathbb{C}$ . The eigenvalues  $a$  approximately uniformly distributed within this disk.

GMRES for a matrix with tightly clustered spectrum around center  $z=2$ . We can observe fast exponential convergence.

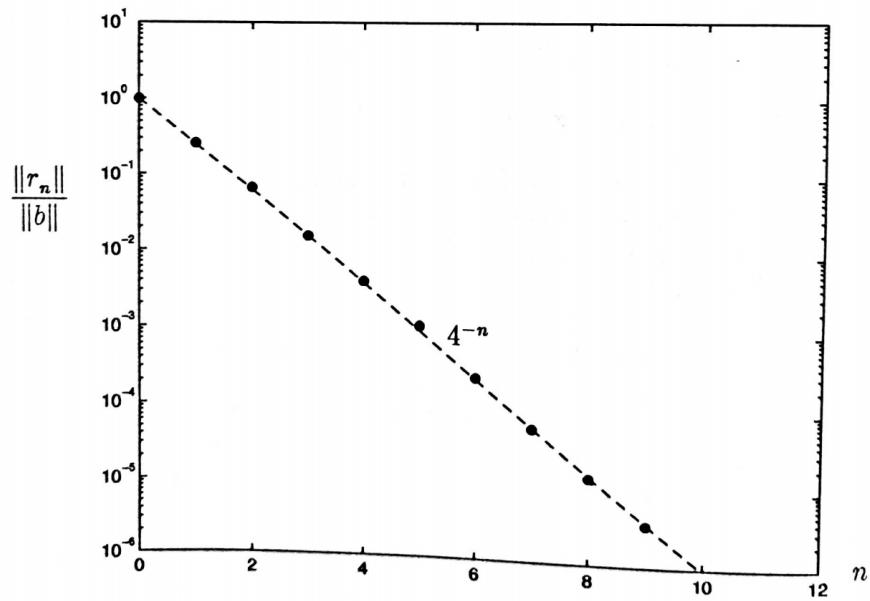


Figure 35.3. GMRES convergence curve for the same matrix  $A$ . This rapid, steady convergence is illustrative of Krylov subspace iterations under ideal circumstances, when  $A$  is a well-behaved (or well-preconditioned) matrix.

# GMRES Convergence...

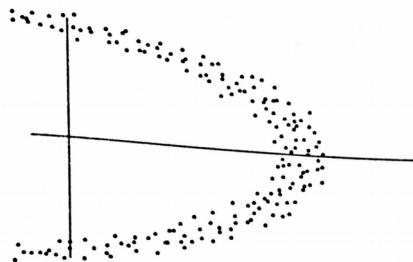


Figure 35.4. Eigenvalues of a  $200 \times 200$  matrix, like that of (35.17) except with a modified diagonal. Now the eigenvalues surround the origin on one side.

GMRES for a matrix,  
where the eigenvalues  
surround the origin.  
GMRES still converges,  
but very slowly.

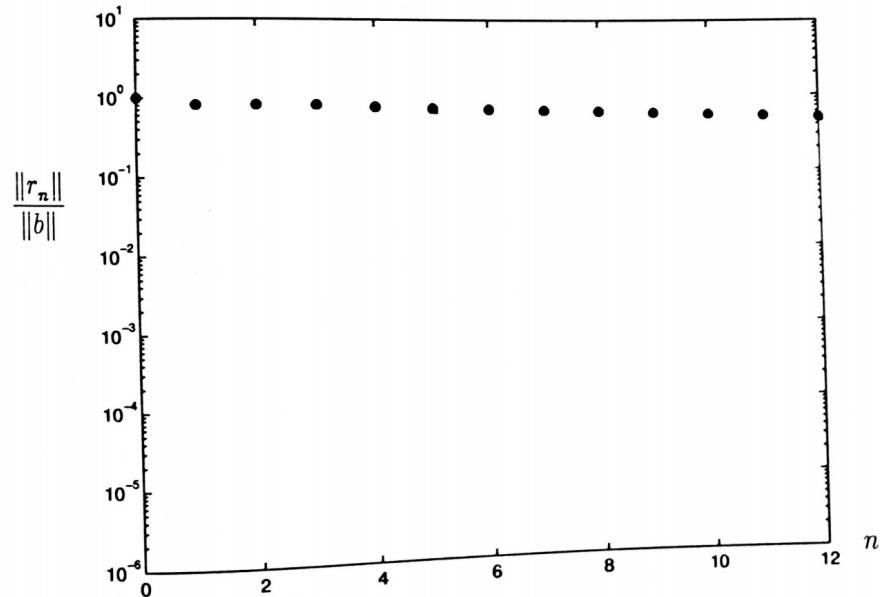


Figure 35.5. GMRES convergence curve for the matrix of Figure 35.4. The convergence has slowed down greatly. When an iterative method stagnates like this, it is time to look for a better preconditioner.

# Symmetric Problems

Let now the matrix A be symmetric.

$$A = A^T$$

From the symmetry we obtain

$$\begin{aligned} H_k^T &= [Q_k^T A Q_k]^T \\ &= Q_k^T A^T Q_k \\ &= Q_k^T A Q_k \\ &= H_k \end{aligned}$$

Hence, H is symmetric. Since it is also upper Hessenberg, it must be tridiagonal.

# The Lanczos Algorithm

We have seen that  $H$  is symmetric and tridiagonal.

Therefore, we have

$$H_k =: T_k = \begin{bmatrix} \alpha_1 & \beta_2 & & \\ \beta_2 & \alpha_2 & \beta_3 & \\ & \ddots & \ddots & \ddots \\ & \beta_{k-1} & \alpha_{k-1} & \beta_k \\ & & \beta_k & \alpha_k \end{bmatrix}.$$

This structure simplifies the Arnoldi method significantly. Instead of all previous vectors we only need to orthogonalize against the last two vectors.

# The Lanczos Algorithm...

$q_1 = b / \|b\|$ ,  $\text{beta}_1 = 0$ ,

$q_0 = 0$

for  $j = 1, \dots, k - 1$

$w = A q_j - \text{beta}_j q_{j-1}$

$\text{alpha}_j = w^T q_j$

$w_j = w - \text{alpha}_j q_j$

$\text{beta}_{j+1} = \|w_j\|$

$q_{j+1} = w_j / \text{beta}_{j+1}$

end

Much cheaper than Arnoldi. We do not need to keep all previous vectors in memory, but only two.

Only orthogonalize against the two previous vectors.

For symmetric matrices we can apply GMRES with the short recursion and arrive at MINRES, a cheap iterative solver for symmetric linear systems of equations.

# Conjugate Gradients

We now assume that  $A$  is symmetric positive definite.

Define the residuals such that

$$r_k \perp \mathcal{K}_k(A, r_0).$$

We have  $r_k = b - Ax_k$ . Hence,  $r_k \in \mathcal{K}_{k+1}(A, r_0)$   
However, since also  $r_k \perp \mathcal{K}_k(A, r_0)$  it follows  
that  $r_k$  shows in the same direction as  $q_{k+1}$ .

We have

$$r_i^T r_k = 0, \quad i < k$$

Define

$$\Delta r = r_k - r_{k-1}, \quad \Delta x = x_k - x_{k-1}$$

It follows that

$$(x_i - x_{i-1})^T (r_k - r_{k-1}) = 0, \quad i < k$$

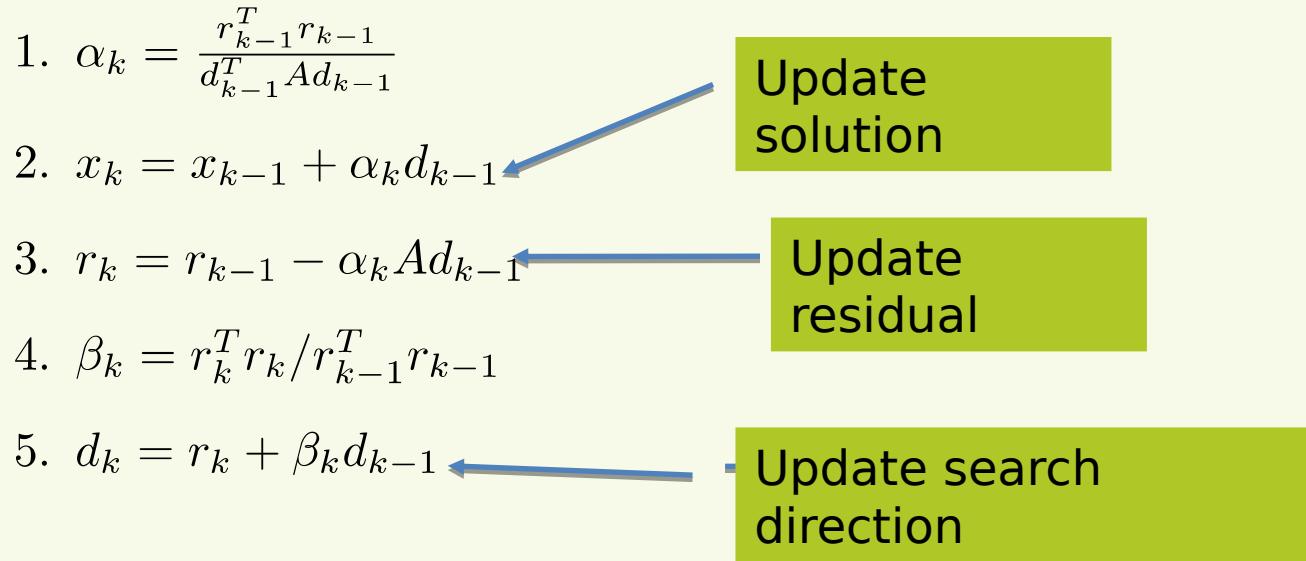
$$\begin{aligned} r_k - r_{k-1} &= (b - Ax_k) - (b - Ax_{k-1}) \\ &= -A(x_k - x_{k-1}) \end{aligned}$$



$$(x_i - x_{i-1})^T A(x_k - x_{k-1}) = 0, \quad i < k$$

# Conjugate Gradients...

Start with  $d_0 = r_0 = b - Ax_0$  and  $x_0$ .



- In each step one matvec is necessary
- CG is only stable for symmetric positive definite matrices
- Very interesting behavior in finite precision!

# CG as optimization method

CG can be interpreted as a line search numerical optimization method to minimize the function

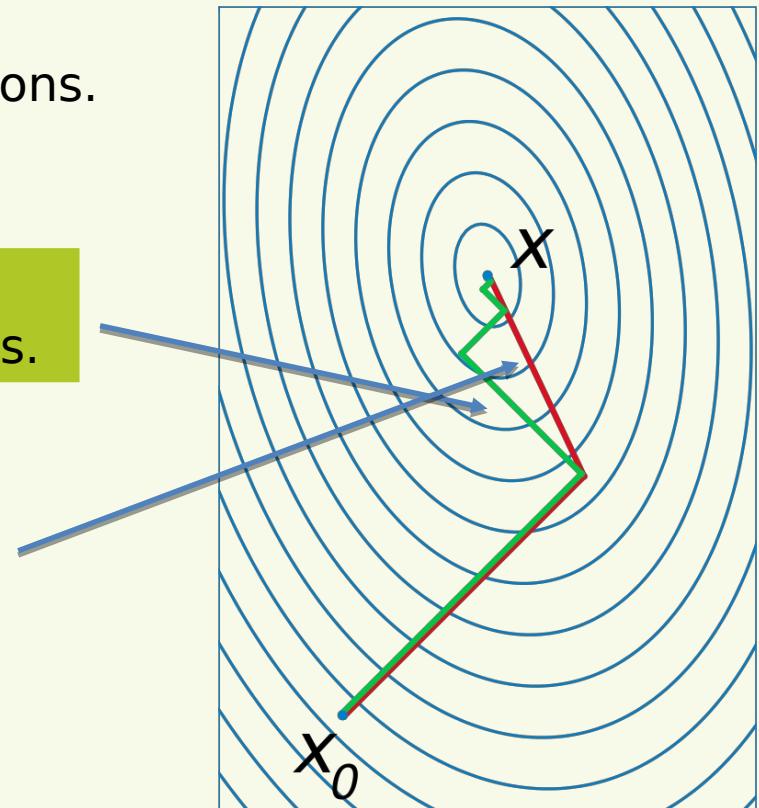
$$f(x) = \frac{1}{2}x^T Ax - x^T b$$

using a sequence of conjugate directions.

Line search with  
orthogonal directions.

Line search with  
conjugate directions.

This interpretation of CG can be generalized to nonlinear optimization problems.



# What method to use?

- Symmetric (Hermitian) Positive Definite: CG is the best choice
- Symmetric (Hermitian): Use Minres
- Nonsymmetric: If there are not too many iterations use GMRES. If memory and large number of iterations are a concern BiCGStab is a frequently used solver.
- Lots of other solvers exist (QMR, CGLS, etc.)
- None of these work well if the problem is ill-conditioned. In that case, preconditioning is necessary (to come...)

# Iterative Solvers in Scipy

Iterative methods for linear equation systems:

<code>bicg(A, b[, x0, tol, maxiter, xtype, M, ...])</code>	Use BIConjugate Gradient iteration to solve $A x = b$
<code>bicgstab(A, b[, x0, tol, maxiter, xtype, M, ...])</code>	Use BIConjugate Gradient STABilized iteration to solve $A x = b$
<code>cg(A, b[, x0, tol, maxiter, xtype, M, callback])</code>	Use Conjugate Gradient iteration to solve $A x = b$
<code>cgs(A, b[, x0, tol, maxiter, xtype, M, callback])</code>	Use Conjugate Gradient Squared iteration to solve $A x = b$
<code>gmres(A, b[, x0, tol, restart, maxiter, ...])</code>	Use Generalized Minimal RESidual iteration to solve $A x = b$ .
<code>lgmres(A, b[, x0, tol, maxiter, M, ...])</code>	Solve a matrix equation using the LGMRES algorithm.
<code>minres(A, b[, x0, shift, tol, maxiter, ...])</code>	Use MINimum RESidual iteration to solve $Ax=b$
<code>qmr(A, b[, x0, tol, maxiter, xtype, M1, M2, ...])</code>	Use Quasi-Minimal Residual iteration to solve $A x = b$

Iterative methods for least-squares problems:

<code>lsqr(A, b[, damp, atol, btol, conlim, ...])</code>	Find the least-squares solution to a large, sparse, linear system of equations.
<code>lsqr(A, b[, damp, atol, btol, conlim, ...])</code>	Iterative solver for least-squares problems.

# Preconditioning

An iterative solver for

$$Ax = b$$

converges best if the eigenvalues of  $A$  are tightly clustered away from zero. Most practical problems do not satisfy this property, e.g. FEM discretization of Laplace operator.

Preconditioning. Let  $P \approx A^{-1}$ .  $P$  is called a preconditioner. We now solve instead

$$PAx = Pb.$$

# Preconditioning...

- The preconditioner  $P$  needs not be explicitly given as matrix, but should be easy to evaluate.
- The best preconditioner is the inverse of  $A$ . But evaluating it is as expensive as solving the original system.
- Analytic Preconditioners: Derive  $P$  from a simpler to solve PDE problem that approximates the original problem in some sense.
- Algebraic Preconditioner: Algebraically derive an easy to invert approximation of  $A$ .

# Splitting Preconditioners

$$x_{k+1} = M^{-1}b + M^{-1}Nx_k$$

Substitute  $N = M - A$  to obtain

$$\begin{aligned} x_{k+1} &= M^{-1}b + M^{-1}(M - A)x_k \\ &= M^{-1}b + (I - M^{-1}A)x_k \end{aligned}$$

In the limit we have

$$x = M^{-1}b + (I - M^{-1}A)x$$

or equivalently

$$M^{-1}Ax = M^{-1}b.$$

We can solve this system directly with a Krylov subspace method.

# ILU – Incomplete LU Decomposition

ALGORITHM 10.2 Gaussian Elimination – IKJ Variant

```
1.      For  $i = 2, \dots, n$  Do:  
2.          For  $k = 1, \dots, i - 1$  Do:  
3.               $a_{ik} := a_{ik} / a_{kk}$   
4.              For  $j = k + 1, \dots, n$  Do:  
5.                   $a_{ij} := a_{ij} - a_{ik} * a_{kj}$   
6.              EndDo  
7.          EndDo  
8.      EndDo
```

- Based on the IKJ Variant of Gaussian Elimination
- Only update an entry if it is not part of a predetermined zero pattern
- ILU(0): Use the zero pattern of original matrix – no fill-in possible
- Use resulting approximate LU Decomposition as Preconditioner

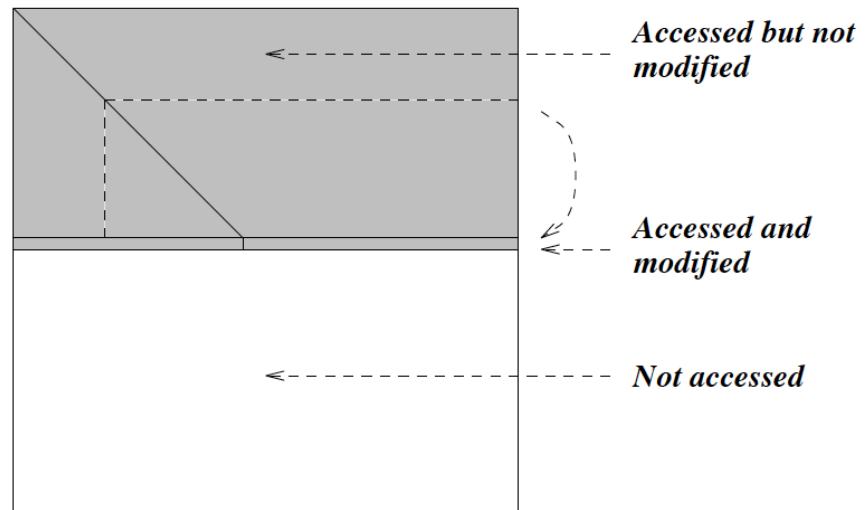


Figure 10.1: IKJ variant of the LU factorization.

# ILU(0)

Non-Zero Pattern of A

## ALGORITHM 10.4 ILU(0)

1.     For  $i = 2, \dots, n$  Do:  
       For  $k = 1, \dots, i - 1$  and for  $(i, k) \in NZ(A)$  Do:  
           Compute  $a_{ik} = a_{ik}/a_{kk}$   
           For  $j = k + 1, \dots, n$  and for  $(i, j) \in NZ(A)$ , Do:  
               Compute  $a_{ij} := a_{ij} - a_{ik}a_{kj}$ .  
       EndDo  
   EndDo

# ILU(P)

**Definition 10.5** The initial level of fill of an element  $a_{ij}$  of a sparse matrix  $A$  is defined by

$$lev_{ij} = \begin{cases} 0 & \text{if } a_{ij} \neq 0, \text{ or } i = j \\ \infty & \text{otherwise.} \end{cases}$$

Each time this element is modified in line 5 of Algorithm 10.2, its level of fill must be updated by

$$lev_{ij} = \min\{lev_{ij}, lev_{ik} + lev_{kj} + 1\}. \quad (10.17)$$

## ALGORITHM 10.5 ILU( $p$ )

1. For all nonzero elements  $a_{ij}$  define  $lev(a_{ij}) = 0$
2. For  $i = 2, \dots, n$  Do:
  3. For each  $k = 1, \dots, i - 1$  and for  $lev(a_{ik}) \leq p$  Do:
    4. Compute  $a_{ik} := a_{ik}/a_{kk}$
    5. Compute  $a_{i*} := a_{i*} - a_{ik}a_{k*}$ .
  6. Update the levels of fill of the nonzero  $a_{i,j}$ 's using (10.17)
  7. EndDo
  8. Replace any element in row  $i$  with  $lev(a_{ij}) > p$  by zero
  9. EndDo

# Multigrid Methods

- PDE adapted solver technique for non-oscillatory problems
- Grid independent convergence for iterative solvers
- Multigrid methods are tremendously successful and have impact far beyond just iterative solver techniques

# A model problem

$$-u''(x) = f(x), \quad x \in (0, 1)$$

$$u(0) = u(1) = 0$$

Choose n equally spaced points

$$x_i = i \times h, \quad i = 0, \dots, n + 1, \quad h = 1/(n + 1)$$

and approximate  $u''(x_i) \approx \frac{x_{i+1} - 2x_i + x_{i-1}}{h^2}$



$$Ax = b$$

$$A = \begin{bmatrix} 2 & -1 & & & \\ -1 & 2 & -1 & & \\ & \ddots & \ddots & \ddots & \\ & & -1 & 2 & -1 \\ & & & -1 & 2 \end{bmatrix}, \quad b = h^2 \begin{bmatrix} f(x_1) \\ f(x_2) \\ \vdots \\ f(x_{n-1}) \\ f(x_n) \end{bmatrix}$$

# Eigenvalues and Eigenvectors

We will analyze the behavior of iterative methods by considering the eigenvalues and eigenvectors of  $A$ .

$$Av_k = \lambda_k v_k$$

$$\lambda_k = 4 \sin^2 \frac{\theta_k}{2}, \quad \theta_k = \frac{k\pi}{n+1}$$

$$v_k = \sin(k\pi x_i)$$

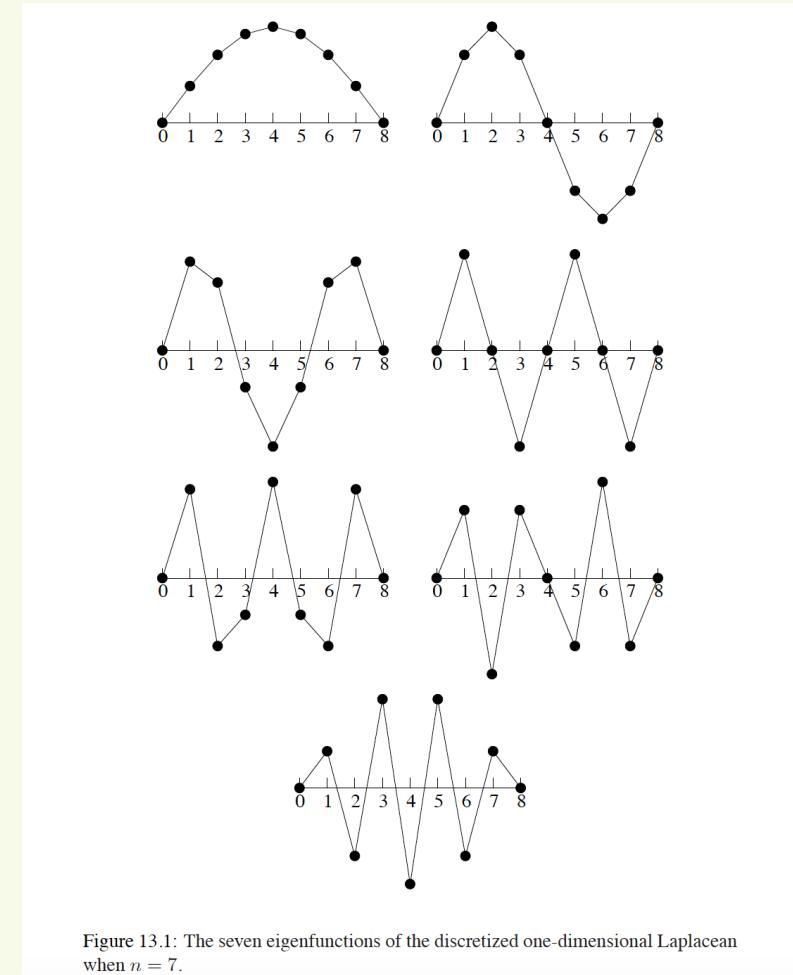


Figure 13.1: The seven eigenfunctions of the discretized one-dimensional Laplacean when  $n = 7$ .

# Richardson's iteration

$$\begin{aligned}x_{j+1} &= x_j + \omega(b - Ax_j) \\&= (I - \omega A)x_j + \omega b\end{aligned}$$

Bounds for eigenvalues  $\mu_i$  of iteration matrix  $M_\omega = I - \omega A$ .

$$1 - \omega \lambda_n \leq \mu_i \leq 1 - \omega \lambda_1$$



$$\omega < \frac{2}{\lambda_n}$$

Let  $\gamma > \lambda_n$  and choose  $\omega = \frac{1}{\gamma}$ .



Richardson's iteration converges.

# Error of Richardson's iteration

Error in iteration  $j$ :

$$e_j \equiv x^* - x_j$$


Limit of Richardson's iteration.

Expand the error  $e_0$  in the basis of eigenvectors of  $A$ .

$$e_0 = \sum_{k=1}^n \xi_k v_k$$

We obtain

$$e_j = M_\omega^j e_0 = \sum_{k=1}^n \left(1 - \frac{\lambda_k}{\gamma}\right)^j \xi_k v_k$$

Each error component is reduced by  $\left(1 - \frac{\lambda_k}{\gamma}\right)^j$ .

# Error of Richardson's iteration...

Choose  $\gamma = 4$ . In the model problem the reduction coefficient for the error with respect to the smallest eigenvalue is

$$1 - \sin^2 \frac{\pi}{2(n+1)} \approx 1 - \left(\frac{\pi h}{2}\right)^2 = 1 - \mathcal{O}(h^2).$$

As  $h$  decreases the rate of convergence becomes arbitrarily slow.

However, the oscillatory part (error components associated with eigenvalues  $k > n / 2$ ) are decreased rapidly since in that case.

$$\mu_k = 1 - \sin^2 \frac{k\pi}{2(n+1)} = \cos^2 \frac{k\pi}{2(n+1)} \leq \frac{1}{2}.$$

# Reduction coefficients for model problem

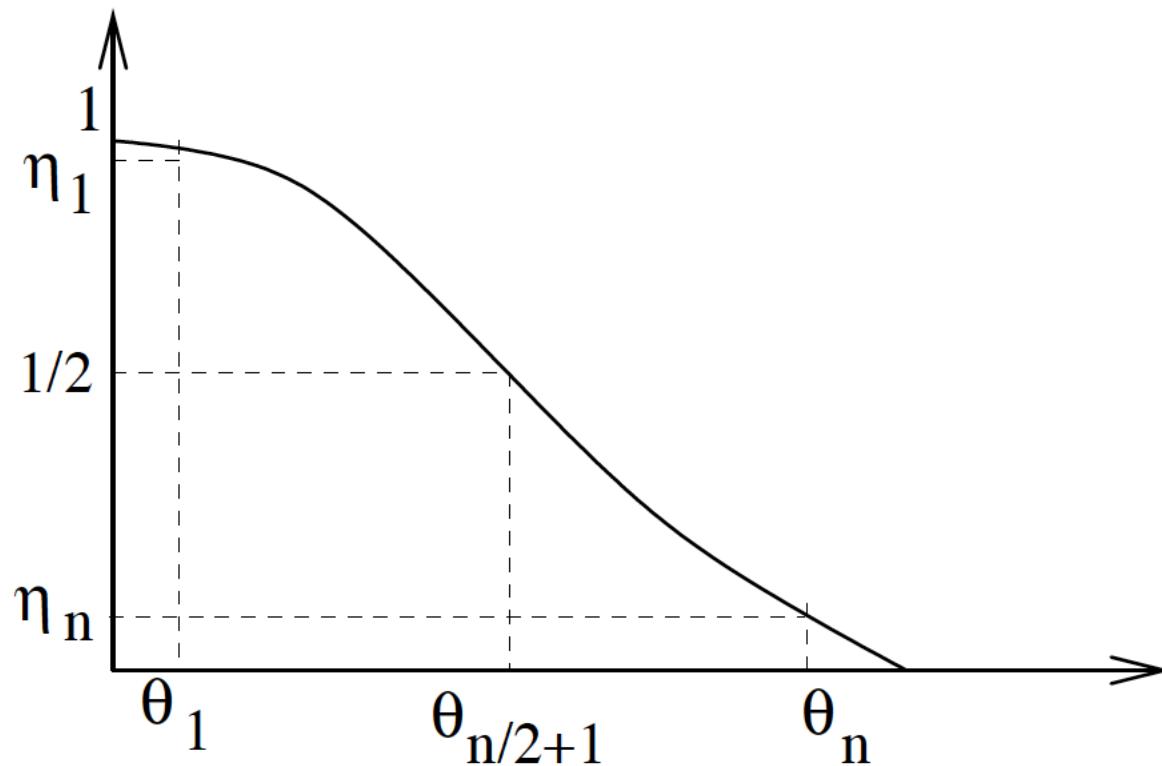


Figure 13.2: Reduction coefficients for Richardson's method applied to the 1-D model problem

# Coarse Grid Problem

Assume that  $n$  is odd. Denote by  $\Omega_{2h}$  the grid obtained by going from  $h$  to  $2h$ . We have

$$x_i^{2h} = i(2h), \quad x_i^{2h} = x_{2i}^h$$

and therefore

$$v_k(x_{2i}^h) = \sin(k\pi x_{2i}^h) = \sin(k\pi x_i^{2h}) = v_k^{2h}(x_i^{2h})$$

for  $k \leq n/2$ .

Some of the non-oscillatory modes on the fine grid become oscillatory modes on the coarse grid.

Multigrid idea: Eliminate high-frequency error components on the fine grid and then move to coarse grid to reduce low-frequency errors.

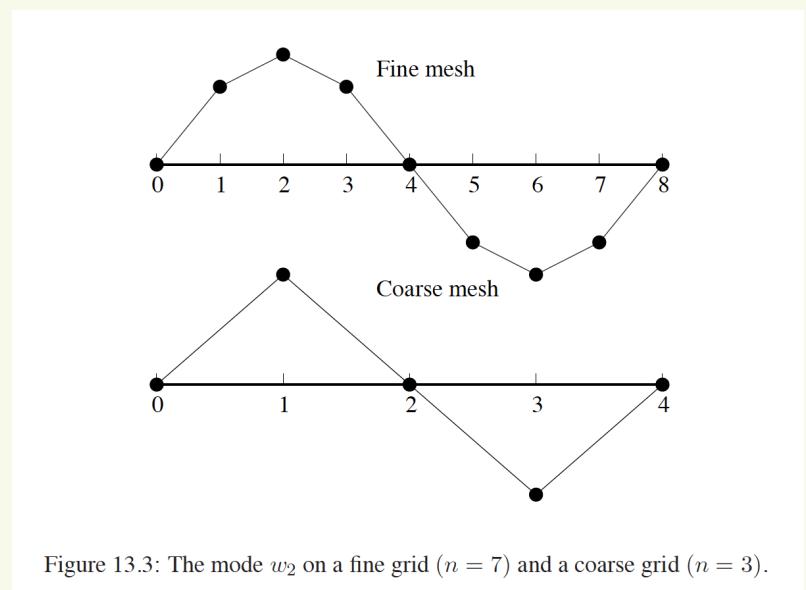


Figure 13.3: The mode  $w_2$  on a fine grid ( $n = 7$ ) and a coarse grid ( $n = 3$ ).

# Moving between fine grid and coarse grid

Fine grid:  $\Omega^h$  , Coarse grid: $\Omega^H$  ,  
 $H=2h$

Prolongation Operator

$$I_H^h : \Omega_H \rightarrow \Omega_h$$

We assume that  $n$  is odd and define the prolongation operation  
 $v^h = I_{2h}^h v^{2h}$  by

$$\begin{cases} v_{2j}^h &= v_j^{2h} \\ v_{2j+1}^h &= (v_j^{2h} + v_{j+1}^{2h})/2 \end{cases} \quad \text{for } j = 0, \dots, \frac{n+1}{2} .$$



Matrix form

$$v^h = \frac{1}{2} \begin{bmatrix} 1 & & & & \\ 2 & 1 & 1 & & \\ & 2 & 1 & 1 & \\ & & 1 & 1 & \\ & & & \vdots & \\ & & & & 1 \end{bmatrix} v^{2h} .$$

# Moving between fine grid and coarse grid...

Restriction of a function  $v^h$  onto  $v^{2h}$  :

$$v_j^{2h} = \frac{1}{4} (v_{2j-1}^h + 2v_{2j}^h + v_{2j+1}^h)$$

$$I_h^{2h} = \frac{1}{4} \begin{bmatrix} 1 & 2 & 1 & & & \\ & 1 & 2 & 1 & & \\ & & 1 & 2 & 1 & \\ & & & \dots & \dots & \dots \\ & & & & 1 & 2 & 1 \end{bmatrix}.$$

It follows that

$$I_{2h}^h = 2(I_h^{2h})^T$$

# The Multigrid Cycle

We want to solve

$$A_h u^h = f^h$$

on the fine grid.

Define the smoothing operation

$$u_\nu^h = \text{smooth}^\nu(A_h, u_0^h, f_h)$$

It applies  $\nu$  steps of a smoother on the linear system

$$A_h u^h = f^h$$

The smoothing operation is given by

$$u_{j+1}^h = S_h u_j^h + g^h \quad \longleftrightarrow \quad u_{j+1}^h = u_j^h + B_h(f^h - A_h u_j^h)$$

$$S_h \equiv I - B_h A_h, \quad B_h \equiv (I - S_h)^{-1} A_h^{-1}, \quad g^h \equiv B_h f^h$$

For Richardson's iteration:

# The Multigrid Cycle...

After the smoothing step for the error and residual have

$$e_\nu^h = (S_h)^\nu e_0^h = (I - B_h A_h)^\nu e_0^h, \quad r_h^\nu = (I - A_h B_h)^\nu r_0^h$$

## Two Grid Cycle

1. *Pre-smooth:*  $u^h := \text{smooth}^{\nu_1}(A_h, u_0^h, f^h)$
2. *Get residual:*  $r^h = f^h - A_h u^h$
3. *Coarsen:*  $r^H = I_h^H r^h$
4. *Solve:*  $A_H \delta^H = r^H$
5. *Correct:*  $u^h := u^h + I_H^h \delta^H$
6. *Post-smooth:*  $u^h := \text{smooth}^{\nu_2}(A_h, u^h, f^h)$

# The Multigrid Cycle

We want to write a single Two-Grid cycle as an iteration of the form

$$u_{new}^h = M_h u_0^h + g_{M_h}.$$

We obtain

$$u_{new}^h = S_h^{\nu_2} [S_h^{\nu_1} u_0^h + I_H^h A_H^{-1} I_h^H (-A_h S_h^{\nu_1} u_0^h)]$$

and therefore

$$M_H = S_h^{\nu_2} [I - I_H^h A_H^{-1} I_h^H A_h] S_h^{\nu_1}.$$

We define

$$T_h^H = I - I_H^h A_H^{-1} I_h^H A_h.$$

This matrix is called the coarse grid correction.

# V-Cycle

ALGORITHM 13.3  $u^h = \text{V-cycle}(A_h, u_0^h, f^h)$

1. *Pre-smooth:*  $u^h := \text{smooth}^{\nu_1}(A_h, u_0^h, f^h)$
2. *Get residual:*  $r^h = f^h - A_h u^h$
3. *Coarsen:*  $r^H = I_h^H r^h$
4. *If* ( $H == h_0$ )  
    *Solve:*  $A_H \delta^H = r^H$
6. *Else*
7.     *Recursion:*  $\delta^H = \text{V-cycle}(A_H, 0, r^H)$
8.     *EndIf*
9.     *Correct:*  $u^h := u^h + I_H^h \delta^H$
10.    *Post-smooth:*  $u^h := \text{smooth}^{\nu_2}(A_h, u^h, f^h)$
11.    *Return*  $u^h$

We recursively apply the Two-Cycle to solve the coarse grid system.

# General Multigrid cycle

ALGORITHM 13.4  $u^h = \mathbf{MG}(A_h, u_0^h, f^h, \nu_1, \nu_2, \gamma)$

1. Pre-smooth:  $u^h := \text{smooth}^{\nu_1}(A_h, u_0^h, f^h)$
2. Get residual:  $r^h = f^h - A_h u^h$
3. Coarsen:  $r^H = I_h^H r^h$
4. If ( $H == h_0$ )
5.     Solve:  $A_H \delta^H = r^H$
6. Else
7.     Recursion:  $\delta^H = \mathbf{MG}^\gamma(A_H, 0, r^H, \nu_1, \nu_2, \gamma)$
8. EndIf
9. Correct:  $u^h := u^h + I_H^h \delta^H$
10. Post-smooth:  $u^h := \text{smooth}^{\nu_2}(A_h, u^h, f^h)$
11. Return  $u^h$

Number of repetitions.

V-Cycle:  $\gamma = 1$

W-Cycle:  $\gamma = 2$

# General Multigrid cycle...

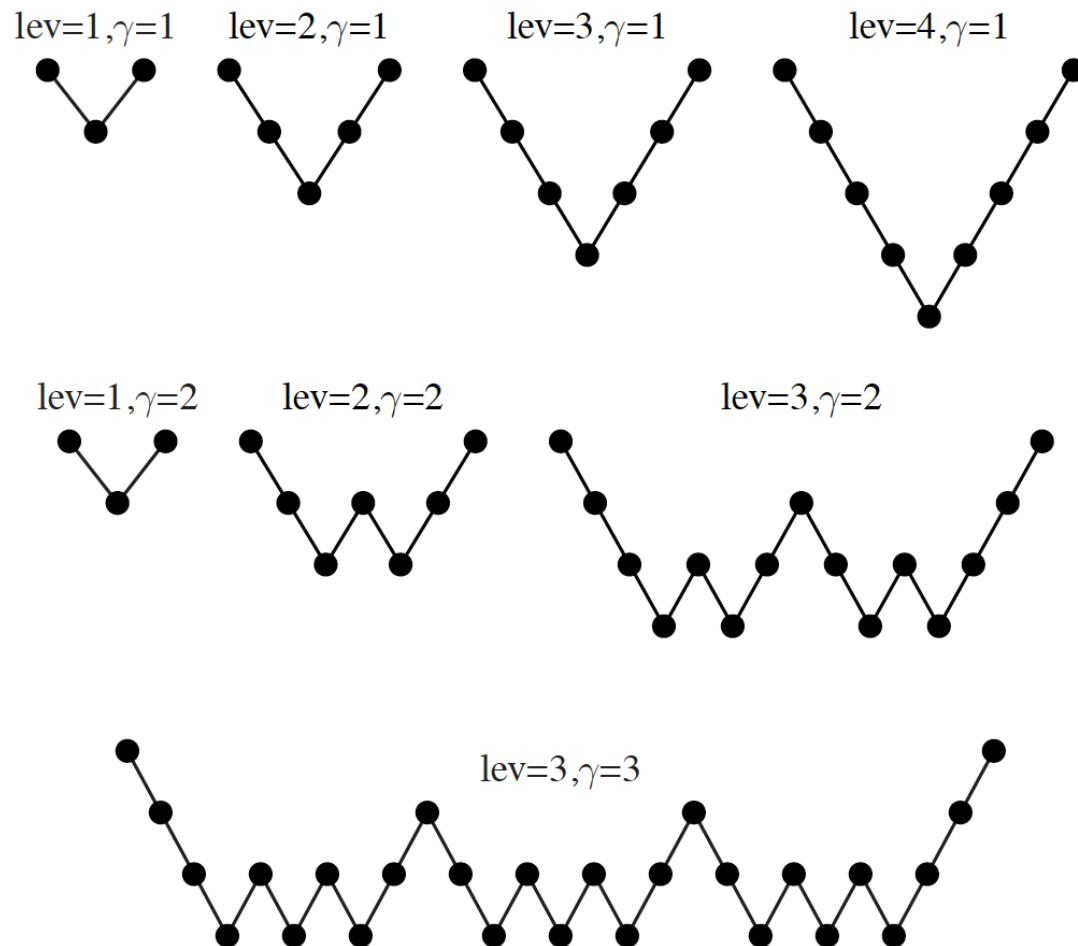


Figure 13.6: Representations of Various V-cycles and W-cycles

# Algebraic Multigrid

- Multigrid requires a sequence of meshes and restriction and prolongation operators between them.
- In practice this is often not available.
- We would like to have a Multigrid like technique that works directly on the matrix without requiring a hierarchy of meshes.

# Multigrid ingredients

Generically, multigrid methods require two ingredients:

1. A way to define a coarse subspace  $X_H$  from a fine subspace  $X_h$ .
2. An interpolation operator  $I_h^H$ .

Coarse grid problem

$$[I_h^H]^T A_h I_h^H u_H = [I_h^H]^T f_h$$

The diagram shows the coarse grid problem equation  $[I_h^H]^T A_h I_h^H u_H = [I_h^H]^T f_h$ . Two blue arrows point to the terms  $[I_h^H]^T$  and  $I_h^H$  respectively. The arrow pointing to  $[I_h^H]^T$  is labeled "Restriction" in a green box below it. The arrow pointing to  $I_h^H$  is labeled "Prolongation" in a green box below it.

# AMG Ideas

For smooth error  $s$  can show heuristically that in Multigrid

$$\sum_{j \neq i} \frac{|a_{ij}|}{|a_{ii}|} \left( \frac{s_i - s_j}{s_i} \right)^2 \ll 1$$

If  $\frac{|a_{ij}|}{|a_{ii}|}$  is large then the corresponding error  $s_i - s_j$

must be small, that is the error  $s$  changes  
slowly along strongly connected nodes in the  
adjacency graph.

Weak connection:  $|a_{ij}|/|a_{ii}| < \sigma$

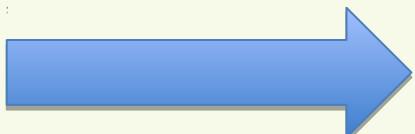
Given threshold  
parameter

Strong connection:

Around a given node  $i$  split up the neighbors into coarse  
grid nodes, strong connections and weak connections and  
treat those separately.

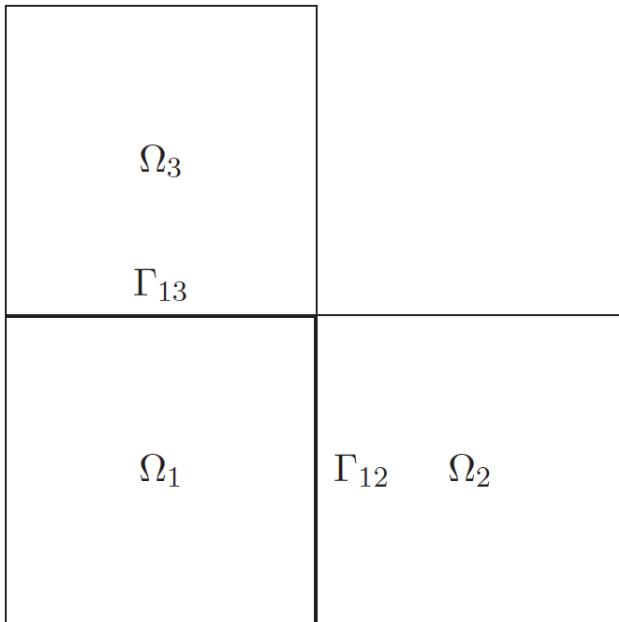
# Multigrid and Krylov solvers

- We can use Multigrid directly as solver
- Alternative to Krylov subspace methods
- Works extremely well for Poisson type problems
- In practice very often use AMG as preconditioner for a Krylov subspace method



Best of both  
worlds!!

# Domain Decomposition



## Model Problem

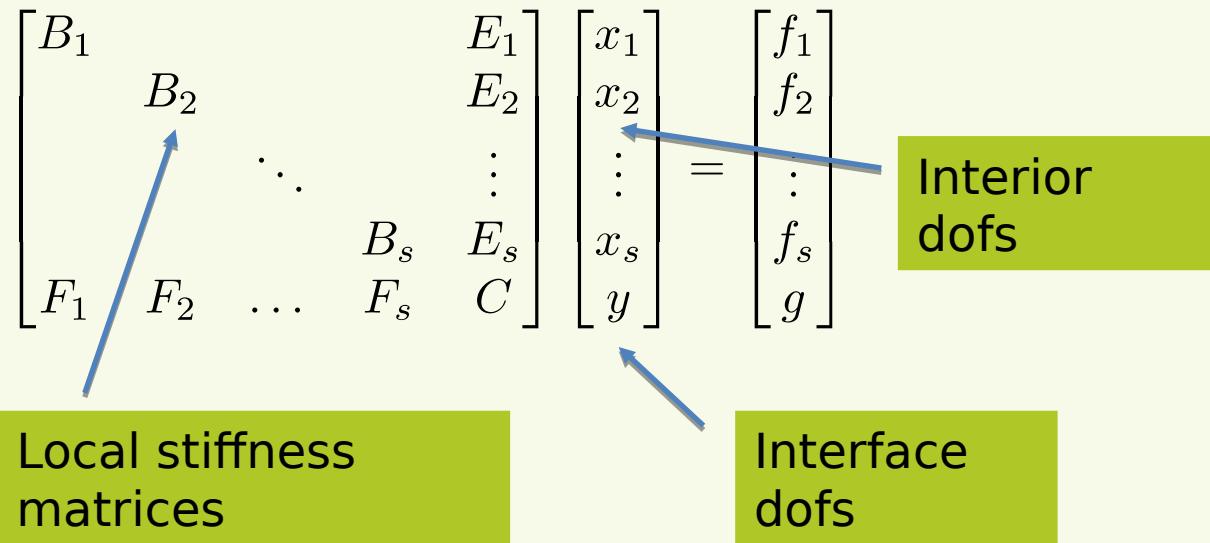
$$-\Delta u = f, \text{ in } \Omega$$

$$u = u_\Gamma, \text{ on } \Gamma$$

- Split complicated geometry into easier subdomains
- Allow parallel solves on the subdomains

$$\Omega = \bigcup_{i=1}^s \Omega_i$$

# Partitioning



$$A \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} f \\ g \end{bmatrix}$$

$$A = \begin{bmatrix} B & E \\ F & C \end{bmatrix}$$

- Sort dofs by interior and interface dofs
- Obtain an arrowhead structure
- We will reduce this system to a problem purely defined on the interfaces

# Partitioning...

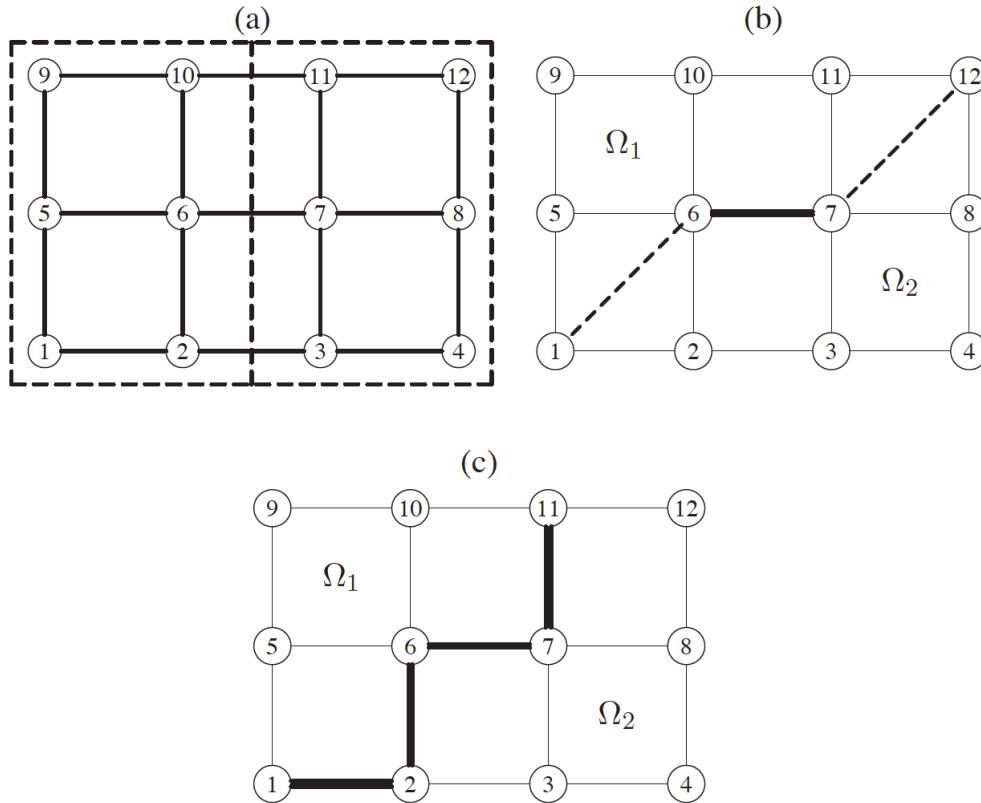


Figure 14.2: (a) Vertex-based, (b) edge-based, and (c) element-based partitioning of a  $4 \times 3$  mesh into two subregions.

# Domain Decomposition Types

- Type of Partitioning: Along edges, elements or vertices
- Overlap: Should subdomains overlap or not
- Treatment of interfaces: Schur complements, interface value updates,...
- Subdomain solution: Approximate solver or direct solver, local preconditioning...

# Schur complements

$$\begin{bmatrix} B & E \\ F & C \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} f \\ g \end{bmatrix}$$



$$x = B^{-1}(f - Ey)$$

Reduced system

$$(C - FB^{-1}E)y = g - FB^{-1}f$$

Define the Schur complement

$$S = C - FB^{-1}E$$

With  $E' = B^{-1}E$ ,  $f' = B^{-1}f$   
have

$$(C - FE')y = g - Ff'$$
$$x = f' - E'y$$

Parallel solve  
for interface  
dofs

Subdomain  
PDE solves

# Properties of the Schur Complement

Block LU Decomposition gives

$$\begin{bmatrix} B & E \\ F & C \end{bmatrix} = \begin{bmatrix} I & 0 \\ FB^{-1} & I \end{bmatrix} \begin{bmatrix} B & E \\ 0 & S \end{bmatrix}$$



Schur  
Complement

$$\begin{aligned} \begin{bmatrix} B & E \\ F & C \end{bmatrix}^{-1} &= \begin{bmatrix} B^{-1} & -B^{-1}ES^{-1} \\ 0 & S^{-1} \end{bmatrix} \begin{bmatrix} I & 0 \\ -FB^{-1} & I \end{bmatrix} \\ &= \begin{bmatrix} B^{-1} + B^{-1}ES^{-1}FB^{-1} & -B^{-1}ES^{-1} \\ -S^{-1}FB^{-1} & S^{-1} \end{bmatrix} \end{aligned}$$

The inverse of the Schur complement is the bottom right entry of the inverse of A.

# Properties of the Schur Complement...

Let  $A$  be nonsingular and partitioned as before, such that the submatrix  $B$  is nonsingular and let  $R_y$  be the restriction operator onto the interface variables, that is

$$R_y \begin{bmatrix} x \\ y \end{bmatrix} = y.$$

The following properties are true:

- The Schur complement matrix  $S$  is nonsingular
- If  $A$  is symmetric pos. definite, then so is  $S$
- For any  $y$ ,

$$S^{-1}y = R_y A^{-1} \begin{bmatrix} 0 \\ y \end{bmatrix}.$$

# Schur complement for FEM partitioning

Variational Problem

Find  $u \in V_h$  such that  $a(u, v) = (f, v)$  for all  $v \in V_h$ .

$$a(u, v) = \int_{\Omega} \nabla u \cdot \nabla v.$$

Assume disjoint subdomains.

Then

$$a_i(u, v) = \sum_{\Omega_i} \nabla u \cdot \nabla v$$

# Schur complement for FEM partitioning...

Assume standard P1 nodal basis functions.  
Sort dofs according to interface nodes and interior nodes to obtain.

$$\begin{bmatrix} B_1 & & E_1 \\ & B_2 & E_2 \\ & \ddots & \vdots \\ F_1 & F_2 & \dots & F_s & C \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_s \\ y \end{bmatrix} = \begin{bmatrix} f_1 \\ f_2 \\ \vdots \\ f_s \\ g \end{bmatrix}$$

On  $i$ th subdomain obtain local stiffness matrix

$$A_i = \begin{bmatrix} B_i & E_i \\ F_i & C_i \end{bmatrix}.$$

This matrix contains only contributions from elements in this subdomain.

# Schur complement for FEM partitioning...

For Schur complement  $S$  now have

$$\begin{aligned} S &= C - FB^{-1}E \\ &= C - \sum_{i=1}^s F_i B_i^{-1} E_i \\ &= \sum_{i=1}^s C_i - \sum_{i=1}^s F_i B_i^{-1} E_i \\ &= \sum_{i=1}^s [C_i - F_i B_i^{-1} E_i]. \end{aligned}$$

Can assemble global Schur complement  
 $S$  from  
local Schur complements

$$S_i = C_i - F_i B_i^{-1} E_i.$$

# Schwarz methods

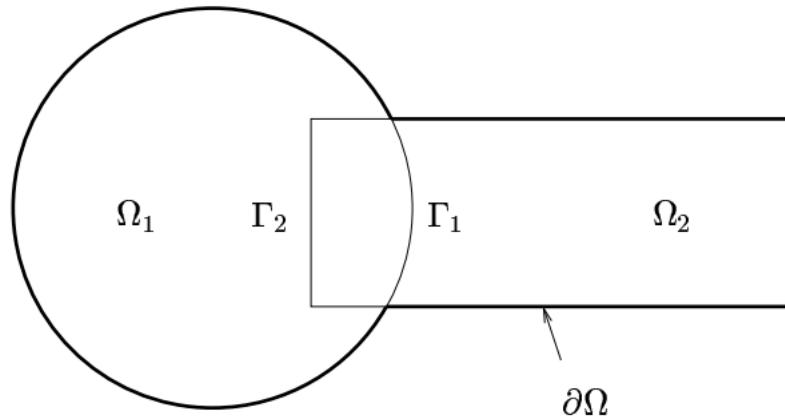


FIGURE 2.1. *The first domain decomposition method was introduced by Schwarz for a complicated domain, composed of two simple ones, namely a disk and a rectangle.*

Here, consider overlapping Schwarz methods, where the subdomains overlap.

# Schwarz Alternating Procedure

$$\begin{aligned}-\Delta u_1^{n+1} &= f_1, \text{ in } \Omega, \\ u_1^{n+1} &= u_2^n, \text{ on } \Gamma_1\end{aligned}$$

$$\begin{aligned}-\Delta u_2^{n+1} &= f_2, \text{ in } \Omega, \\ u_2^{n+1} &= u_1^{n+1}, \text{ on } \Gamma_1\end{aligned}$$

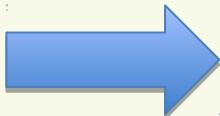
Matrix form:

$$\begin{bmatrix} A_1 & A_{12} \\ A_{21} & A_2 \end{bmatrix} \begin{bmatrix} u_1^{n+1} \\ u_2^n \end{bmatrix} = \begin{bmatrix} f_1 \\ f_2 \end{bmatrix}$$

$$\begin{bmatrix} A_1 & A_{12} \\ A_{21} & A_2 \end{bmatrix} \begin{bmatrix} u_1^{n+1} \\ u_2^{n+1} \end{bmatrix} = \begin{bmatrix} f_1 \\ f_2 \end{bmatrix}$$

$$A_1 u_1^{n+1} = f_1 - A_{12} u_2^n$$

$$A_2 u_2^{n+1} = f_2 - A_{21} u_1^{n+1}$$



# Schwarz Alternating Procedure...

Residual Form

$$A_1\delta_1^n = f_1 - A_{12}u_2^n - A_1u_1^n =: r_1^n \quad A_2\delta_2^n = f_2 - A_{21}u_1^{n+1} - A_2u_2^n =: r_2^n$$

$$u_1^{n+1} = u_1^n + \delta_1^n \quad u_2^{n+1} = u_2^n + \delta_2^n$$

Parallel Schwarz Method

$$A_1\delta_1^n = f_1 - A_{12}u_2^n - A_1u_1^n =: r_1^n \quad A_2\delta_2^n = f_2 - A_{21}u_1^n - A_2u_2^n =: r_2^n$$

$$u_1^{n+1} = u_1^n + \delta_1^n \quad u_2^{n+1} = u_2^n + \delta_2^n$$

# Algebraic Schwarz Methods

Define restriction operators

$$R_1 = \begin{bmatrix} 1 & & \\ & \ddots & \\ & & 1 \end{bmatrix}, \quad R_2 = \begin{bmatrix} & & 1 \\ & \ddots & \\ & & 1 \end{bmatrix},$$

and projections

$$A_j = R_j A R_j^T, \quad j = 1, 2$$

Multiplicative Schwarz  
Method

$$u^{n+\frac{1}{2}} = u^n + R_1^T A_1^{-1} R_1 (f - A u^n)$$

$$u^{n+1} = u^{n+\frac{1}{2}} + R_2^T A_2^{-1} R_2 (f - A u^{n+\frac{1}{2}})$$

# Multiplicative Schwarz

Without overlap obtain

$$\begin{bmatrix} A_1 & 0 \\ A_{21} & A_2 \end{bmatrix} \begin{bmatrix} u_1^{n+1} \\ u_2^{n+1} \end{bmatrix} = \begin{bmatrix} 0 & -A_{12} \\ 0 & 0 \end{bmatrix} \begin{bmatrix} u_1^n \\ u_2^n \end{bmatrix} + \begin{bmatrix} f_1 \\ f_2 \end{bmatrix}$$



Gauss-Seidel Iteration

Using suitable boundary conditions the multiplicative Schwarz method is identical to the Schwarz alternating procedure, even for overlapping boundaries.

For  $s$  subdomains multiplicative Schwarz takes the form

$$u^{n+\frac{j}{s}} = u^{n+\frac{j-1}{s}} + R_j^T A_j^{-1} R_j \left( f - A u^{n+\frac{j-1}{s}} \right), \quad j = 1, 2, \dots, s$$

# Multiplicative Schwarz...

Subtracting the iterates from the exact solution and defining the error vector  $d$  results in

$$\begin{aligned} u - u^{n+\frac{j}{s}} &= u - u^{n+\frac{j-1}{s}} - R_j^T A_j^{-1} R_j A \left( u - u^{n+\frac{j-1}{s}} \right), \\ d^{n+\frac{j}{s}} &= d^{n+\frac{j-1}{s}} - R_j^T A_j^{-1} R_j A d^{n+\frac{j-1}{s}}, \quad j = 1, 2, \dots, s \\ &= (I - P_j) d^{n+\frac{j-1}{s}} \end{aligned}$$

Hence,

$$d^{n+1} = Q_s d^n$$

with

$$Q_s = (I - P_s) \dots (I - P_1).$$

From  $u - u^{n+1} = Q_s(u - u^n)$  obtain

$$u^{n+1} = Q_s u^n + (I - Q_s)u.$$

# Multiplicative Schwarz...

Remember the fixed point iteration

$$u^{n+1} = M^{-1}f + (I - M^{-1}A)u^n$$

In the limit have

$$u = M^{-1}f + (I - M^{-1}A)u$$

or equivalently

$$M^{-1}Au = M^{-1}f.$$

Multiplicative Schwarz :  $u^{n+1} = Q_s u^n + (I - Q_s)u.$

We therefore obtain

$$M^{-1}A = I - Q_s$$

$$M^{-1}f = (I - Q_s)u = (I - Q_s)A^{-1}f$$

# Multiplicative Schwarz...

Define

$$Z_i = I - Q_i, \quad M_i = Z_i A^{-1}, \quad T_i = P_i A^{-1} = R_i^T A_i^{-1} R_i$$

We have  $M^{-1} = M_s$ ,  $M^{-1}A = Z_s$   
with the recurrence relation

$$Z_1 = P_1$$

$$M_1 = T_1$$

$$Z_i = Z_{i-1} + P_i(I - Z_{i-1}) \quad M_i = M_{i-1} + T_i(I - A M_{i-1})$$



$$Z_k = \sum_{j=1}^i P_j Q_{j-1}, \quad Q_0 \equiv I$$

# Multiplicative Schwarz Preconditioning

Solve the preconditioning multiplicative Schwarz system

$$(I - Q_s)u = (I - Q_s)A^{-1}f$$

by a Krylov subspace solver such as GMRES.

ALGORITHM 14.4 Multiplicative Schwarz Preconditioner

1.     Input:  $v$ ; Output:  $z = M^{-1}v$ .
2.      $z := T_1v$
3.     For  $i = 2, \dots, s$  Do:  
         $z := z + T_i(v - Az)$
4.     EndDo

Only solves on the subdomains are necessary!

ALGORITHM 14.5 Multiplicative Schwarz Preconditioned Operator

1.     Input:  $v$ , Output:  $z = M^{-1}Av$ .
2.      $z := P_1v$
3.     For  $i = 2, \dots, s$  Do  
         $z := z + P_i(v - z)$
4.     EndDo

# Additive Schwarz

The idea of additive Schwarz is similar to the parallel Schwarz alternating method. The new iteration is obtained as

$$u^{n+1} = u^n + \sum_{j=1}^s R_j A^{-1} R_j^T (f - A u^n).$$

Note: This method does not converge in the overlap!

We can still use additive Schwarz as an effective preconditioner though.

$$M^{-1} A u = M^{-1} f$$

$$M^{-1} A = \sum_{j=1}^s P_j$$

$$M^{-1} = \sum_{j=1}^s T_j$$

# Additive Schwarz

1. *For*  $i = 1, \dots, s$  *Do*  
    *Compute*  $\delta_i = R_i^T A_i^{-1} R_i(b - Ax)$
2. *EndDo*
3.  $x_{new} = x + \sum_{i=1}^s \delta_i$

## ALGORITHM 14.6 Additive Schwarz Preconditioner

1. *Input:*  $v$ ; *Output:*  $z = M^{-1}v$ .
2. *For*  $i = 1, \dots, s$  *Do:*  
    *Compute*  $z_i := T_i v$
3. *EndDo*
4. *Compute*  $z := z_1 + z_2 \dots + z_s$ .

# Schur Complement Preconditioners

Let the linear system be given as

$$\begin{bmatrix} B & E \\ F & C \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = \begin{bmatrix} f \\ g \end{bmatrix}$$

Schur complement system

$$(C - FB^{-1}E)y = g - FB^{-1}f$$

How can we efficiently solve this Schur complement system with the left-hand side

$$S = C - FB^{-1}E$$

Applying S to a vector:

1. Compute  $v' = Ev$ ,
2. Solve  $Bz = v'$
3. Compute  $w = Cv - Fz$ .

Need a good preconditioner for the Schur complement.

# Induced Preconditioner

Restriction Operator

$$R_y \begin{bmatrix} x \\ y \end{bmatrix} = y.$$

Let  $P$  be a good preconditioner for the original matrix  $A$ .

We define the induced Schur complement preconditioner as

$$S^{-1} \approx R_y P R_y^T.$$

Let  $P^{-1} = L_A U_A$  be an ILU preconditioner for  $A$ . Then the preconditioner  $P_S$  for  $S$  induced by  $P$  is given by

$$P_S^{-1} = L_S U_S, \quad L_S = R_y L_A R_y^T, \quad U_S = R_y U_A R_y^T.$$

# An introduction to PETSc

- Scalable software library for the solution of algebraic systems arising from PDEs
- Developed in C. But very good Python interface available (`petsc4py`)
- FEniCS by default uses PETSc as linear algebra backend

This and the following slides incorporate information and graphs from the PETSc tutorial available at  
<https://www.mcs.anl.gov/petsc/documentation/tutorials/PetscTu06.pdf>

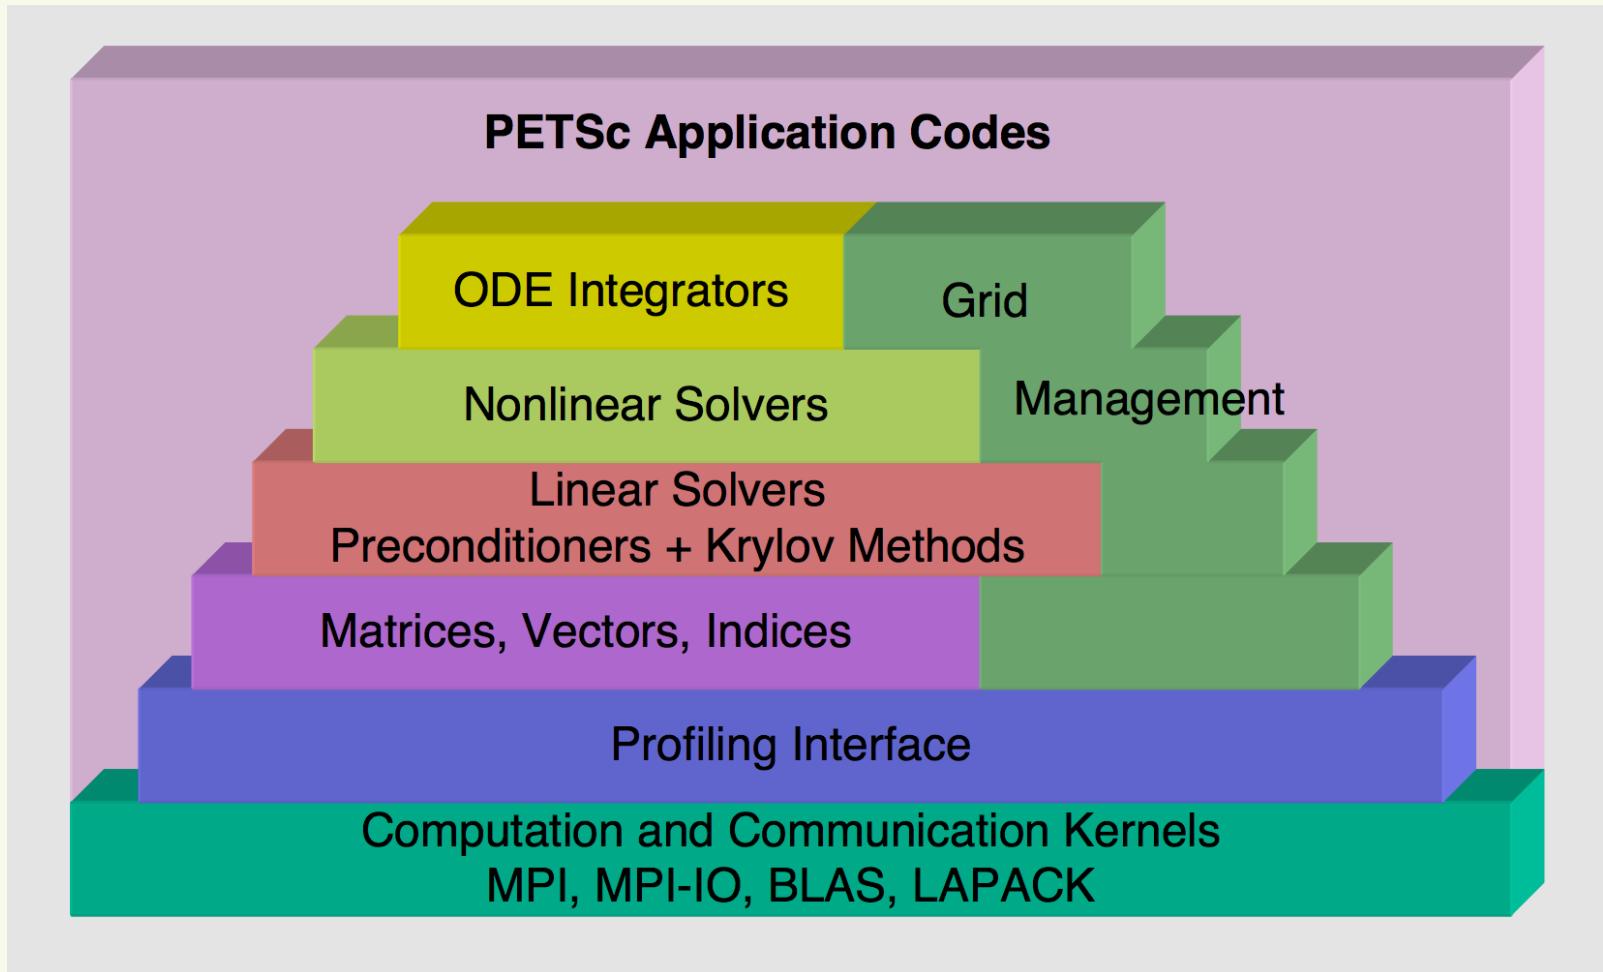
# Main features

- Interfaces to a large range of numerical packages
- Parallel vector/array operations
- A range of parallel matrix formats
- A huge selection of built-in iterative solvers
- Nonlinear solvers
- Some ODE integration capabilities
- Interfaces to Algebraic Multigrid
- PETSc is used for massive Peta-Scale applications

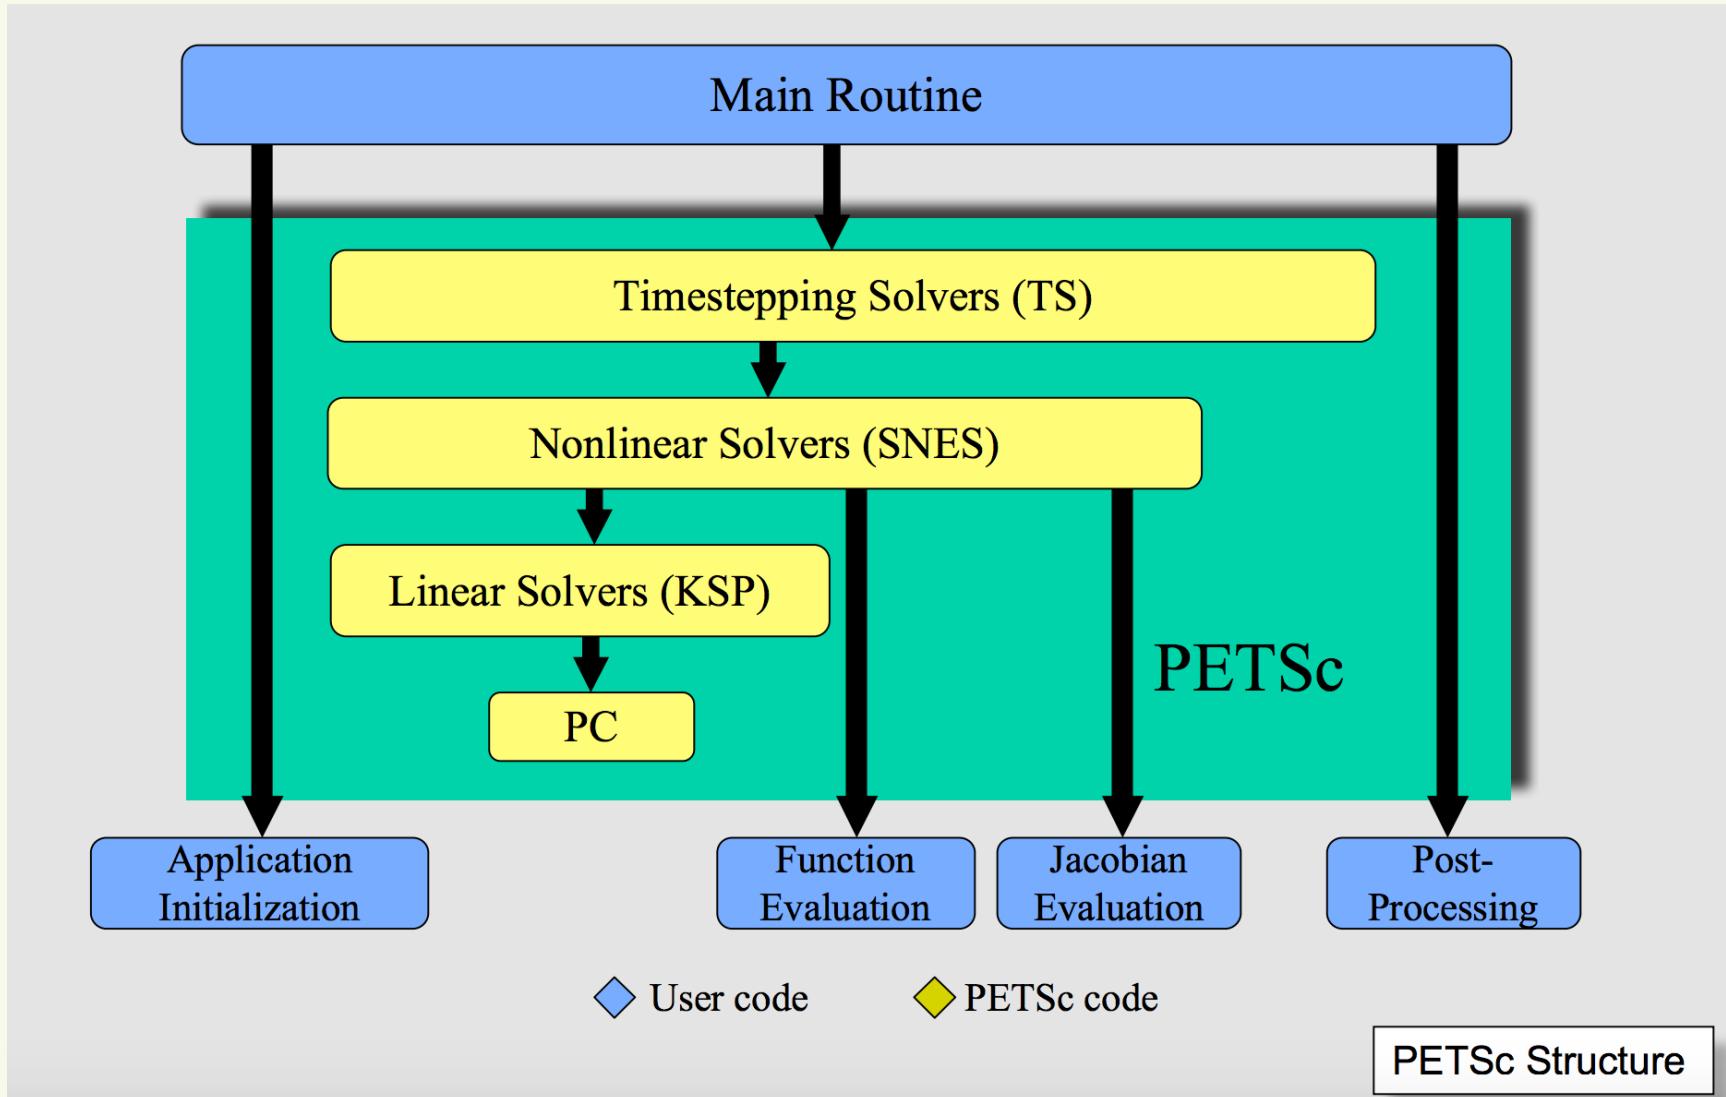
# Feature Matrix

- <https://www.mcs.anl.gov/petsc/documentation/linearsolvertable.html>

# PETSc Structure



# Overview of a PDE Simulation



# Vectors

- Sequential and Parallel Vectors.
- Parallel vectors are distributed across all processes.

```
# Creating a vector in Python
```

```
from petsc4py import PETSc
v_mpi = PETSc.createMPI(10)
v_seq = PETSc.createSeq(10)
```

```
# Get the local indices
```

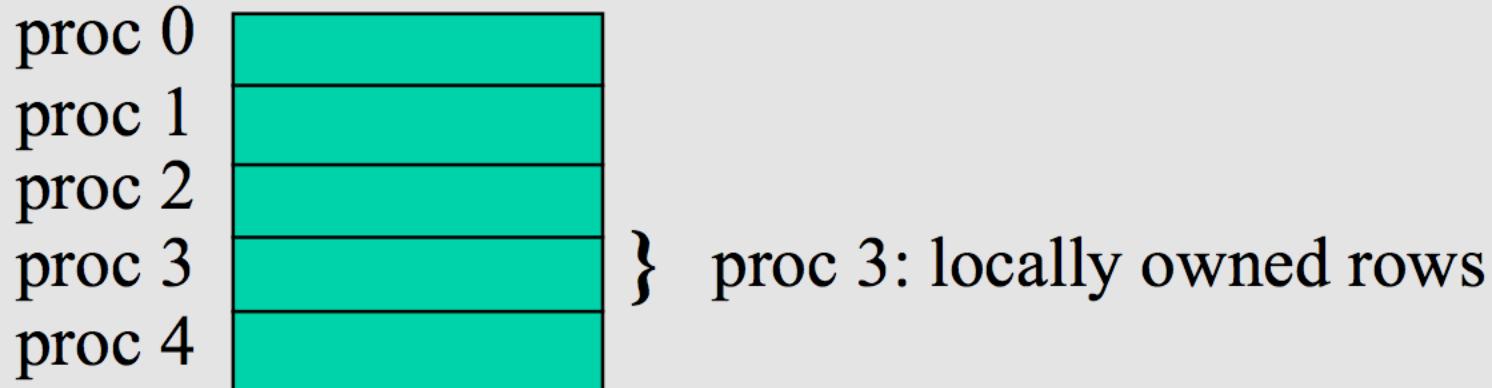
```
s, e = v.getOwnershipRange()
```

```
# Assign all local elements of  
the vector the value 1
```

```
with v as a:  
    a[:] = 1
```

# Matrices

Rows are distributed among processes



`A = PETSc.Mat().createAIJ([m, n], nnz=10)`

Can create matrix also directly  
from CSR data.

Upper bound on  
nonzeros per row

To finalize assembly and  
communicate data use  
`A.AssemblyBegin()`  
`A.AssemblyEnd()`

# The KSP Solver Module

```
#Initialize a solver for GMRES with ILU  
Preconditioner
```

```
ksp = PETSc.KSP().create()  
ksp.setType('gmres')  
ksp.setOperators(A)  
ksp.getPC().setType('ilu')  
ksp.setFromOptions()  
ksp.solve(b, x)
```

Disable iterations.  
Only use  
preconditioner

#Initialize a solver for a direct LU solve

```
ksp = PETSc.KSP().create()  
ksp.setType('preonly')  
ksp.setOperators(A)  
ksp.getPC().setType('lu')  
ksp.setFromOptions()  
ksp.solve(b, x)
```

# Trilinos

- Sandia Based C++ Library for computational PDEs.
- More functionality than PETSc, but less stable development cycle (more disruptive changes over time)
- <https://trilinos.org>

# Some Trilinos Packages

Epetra  
(Legacy  
distributed  
linear algebra  
module)

Tpetra  
(Templated  
distributed  
linear algebra  
module)

Thyra  
(Linear Algebra  
adapters to  
external codes)

Amesos(2)  
(Direct solver  
interface)

Belos  
(Iterative solver  
interface)

Intrepid  
(PDE  
Discretization  
Library)

Nox  
(Nonlinear  
Solvers)

ML  
(Multigrid  
Library)

PyTrilinos  
(Python  
interface to  
Legacy  
Libraries)

# PyTrilinos – Direct Solver Example

```
#! /usr/bin/env python
from PyTrilinos import Amesos, Triutils, Epetra

Comm = Epetra.PyComm()
Map, A, x, b, Exact = Triutils.ReadHB("fidap035.rua", Comm)

Problem = Epetra.LinearProblem(A, x, b);
Factory = Amesos.Factory()
SolverType = "MUMPS"
Solver = Factory.Create(SolverType, Problem)
AmesosList = {
    "MaxProcs": 2,
    "PrintStatus": True
}
Solver.SetParameters(AmesosList)
Solver.SymbolicFactorization()
Solver.NumericFactorization()
Solver.Solve()
```

All solvers can be accessed  
in parallel through a Python  
script with no effort

```
$ mpirun -np 4 python my-script.py
```

# PyTrilinos – Iterative Solver Example

```
#!/usr/bin/env python
from PyTrilinos import AztecOO, Triutils, Epetra

Comm = Epetra.PyComm()
Map, A, x, b, Exact = Triutils.ReadHB("fidap035.rua", Comm)

Solver = AztecOO.AztecOO(A, x, b)
Solver.SetAztecOption(AztecOO.AZ_solver, AztecOO.AZ_cg)
Solver.SetAztecOption(AztecOO.AZ_precond,
                      AztecOO.AZ_dom_decomp)
Solver.SetAztecOption(AztecOO.AZ_subdomain_solve,
                      AztecOO.AZ_icc)
Solver.SetAztecOption(AztecOO.AZ_graph_fill, 1)

Solver.Iterate(1550, 1e-5)
```

```
$ mpirun -np 4 python my-script.py
```

# Finite Difference Methods

- We will give an overview on solving stationary and time-domain problems using finite difference methods
- We will not go deeply into the Mathematical Theory
- Focus on tools and implementation
- Books:
  - K.W. Morton and D.F. Mayers, Numerical Solution of Partial Differential Equations
  - H. P. Langtangen and S.Linge, Finite Difference Computing with PDEs – A modern software approach (available online)

# 1D Model Problem

Consider the 1d oscillator

$$u'' + \omega^2 u = 0$$

$$u(0) = U_0$$

$$u'(0) = 0$$

Exact solution

$$u(t) = U_0 \cos(\omega t)$$

$$u''(t_n) + \omega^2 u(t_n) = 0, \quad n = 1, \dots, N$$

$$u''(t_n) \approx \frac{u^{n+1} - 2u^n + u^{n-1}}{\Delta t^2}$$

$$\frac{u^{n+1} - 2u^n + u^{n-1}}{\Delta t^2} = -\omega^2 u^n$$

$$\frac{u^1 - u^{-1}}{2\Delta t} = 0$$

1. Discretize the domain
2. Satisfy the equation at discrete time points
3. Replace derivatives by finite differences
4. Formulate algorithm

# Time-Stepping Algorithm

Solution values computed by the following iteration:

$$u^{n+1} = 2u^n - u^{n-1} - \Delta t^2 \omega^2 u^n$$

From initial conditions have

$$u^{-1} = u^1$$

and therefore

$$u^1 = 2u^0 - u^1 - \Delta t^2 \omega^2 u^0 = u^0 - \frac{1}{2} \Delta t^2 \omega^2 u^0.$$

# Analysis

Assume that

$$u^n = U_0 A^n, \quad A = e^{i\tilde{\omega}\Delta t}$$

Unknown numerical  
wavenumber

$$\begin{aligned} \frac{u^{n+1} - 2u^n + u^{n-1}}{\Delta t^2} &= U_0 \frac{A^{n+1} - 2A^n + A^{n-1}}{\Delta t^2} \\ &= \frac{U_0}{\Delta t^2} \left( e^{i\tilde{\omega}(t_n + \Delta t)} - 2e^{i\tilde{\omega}t_n} + e^{i\tilde{\omega}(t_n - \Delta t)} \right) \\ &= U_0 e^{i\tilde{\omega}t_n} \frac{2}{\Delta t^2} (\cos(\tilde{\omega}\Delta t) - 1) \\ &= -U_0 e^{i\tilde{\omega}t_n} \frac{4}{\Delta t^2} \sin^2 \left( \frac{\tilde{\omega}\Delta t}{2} \right) \end{aligned}$$

We obtain

$$-U_0 e^{i\tilde{\omega}t_n} \frac{4}{\Delta t^2} \sin^2 \left( \frac{\tilde{\omega}\Delta t}{2} \right) + \omega^2 U_0 e^{i\tilde{\omega}t_n} = 0.$$

# Analysis...

We obtain

$$\tilde{\omega} = \pm \frac{2}{\Delta t} \sin^{-1} \left( \frac{\omega \Delta t}{2} \right).$$

Taylor series expansion gives

$$\tilde{\omega} = \omega \left( 1 + \frac{1}{24} \omega^2 \Delta t^2 \right) + \mathcal{O}(\Delta t^4).$$

Need to keep  $\omega \Delta t$  small. In practice choose between 20 and 30 points per wavelength, even more for longer simulations.

Global error

Error at nth time step

$$\xrightarrow{} e^n = \frac{1}{24} n \omega^3 \Delta t^3$$

Error until time T

$$\xrightarrow{} \|e^n\|_{l^2} = \frac{1}{24} \sqrt{T^3} 3 \omega^3 \Delta t^2$$

# First-Order Schemes

First order formulation of oscillator

$$\begin{aligned} u' &= v \\ v' &= -\omega^2 u \end{aligned} \quad \xrightarrow{\hspace{1cm}} \quad \vec{u}' = A\vec{u}, \quad \vec{u} = \begin{bmatrix} u \\ v \end{bmatrix}, \quad A = \begin{bmatrix} 0 & 1 \\ -\omega^2 & 0 \end{bmatrix}$$

Euler's method

$$\vec{u}^{n+1} = (I + \Delta t A) \vec{u}^n \quad u^{n+1} = 2u^n - u^{n-1} - \Delta t^2 \omega^2 u^{n-1}$$

Implicit Euler

$$(I - \Delta h A) \vec{u}^{n+1} = u^n \quad \frac{u^{n+1} - 2u^n + u^{n-1}}{\Delta t^2} = -\omega^2 u^{n+1}$$

Crank-Nicholson

$$(I - \frac{\Delta t}{2} A) \vec{u}^{n+1} = (I + \frac{\Delta t}{2} A) \vec{u}^n \quad \frac{u^{n+1} - 2u^n + u^{n-1}}{\Delta t^2} = -\omega^2 u^n + \mathcal{O}(\Delta t^2)$$

# Wave Equation

$$\begin{aligned} u_{tt} &= c^2 u_{xx} \\ u(x, 0) &= U_0 \\ u_t(x, 0) &= 0 \\ u(0, t) &= 0 \\ u(L, t) &= 0 \end{aligned}$$

- Use equispaced points in space with spacing  $\Delta x$ .
- Use equispaced points in time with spacing  $\Delta t$ .
- $u_i^n := u(x_i, t_n)$

Finite difference discretization

$$\frac{u_i^{n+1} - 2u_i^n + u_i^{n-1}}{\Delta t^2} = c^2 \frac{u_{i+1}^n - 2u_i^n + u_{i-1}^n}{\Delta x^2}$$

Recursion

$$u_i^{n+1} = -u_i^{n-1} + 2u_i^n + C^2 \frac{u_{i+1}^n - 2u_i^n + u_{i-1}^n}{\Delta x^2}$$

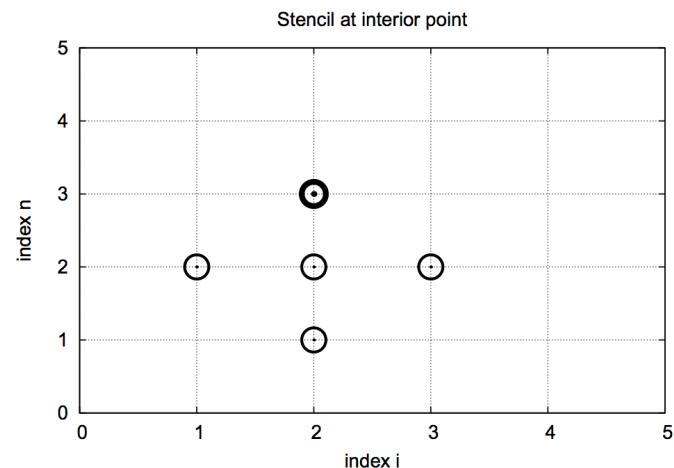
$$C = c \frac{\Delta t}{\Delta x}$$

Initial Conditions:

$$\begin{aligned} u_i^{-1} &= u_i^1 \\ u_i^0 &= U_0(x_i) \end{aligned}$$

Courant Number

# Stencil

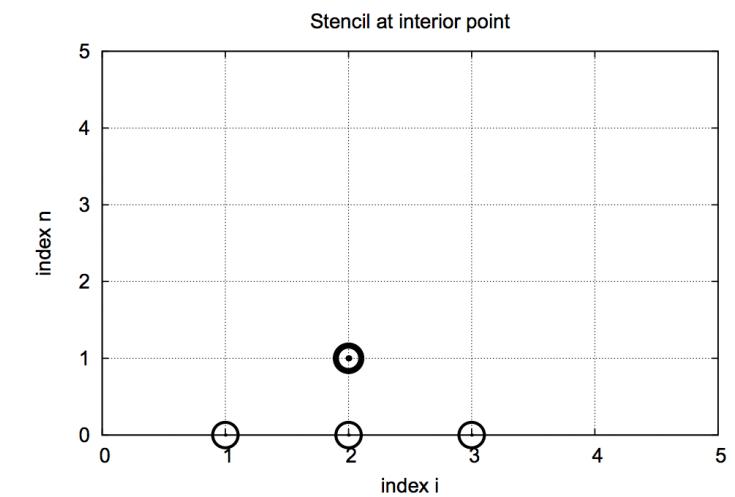


Space-Time dependency of discrete solution.

CFL Condition: The numerical domain of dependence must contain the analytic domain of dependence.

$$\rightarrow c \frac{\Delta t}{\Delta x} \leq 1$$

Stencil for the initial condition



# Wave Equation in Multiple Dimensions

$$u_{tt} = c^2 \Delta u$$

$$\Delta u = u_{xx} + u_{yy}, \quad (d = 2)$$

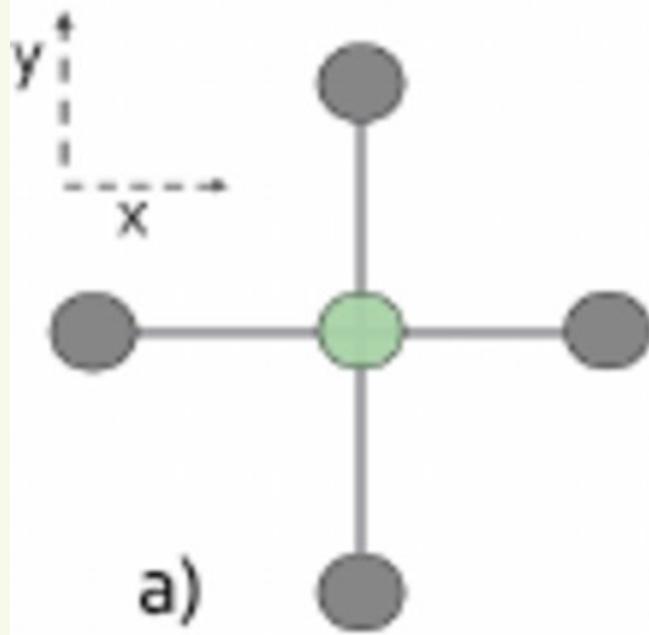
$$\Delta u = u_{xx} + u_{yy} + u_{zz}, \quad (d = 3)$$

+ initial and boundary conditions.

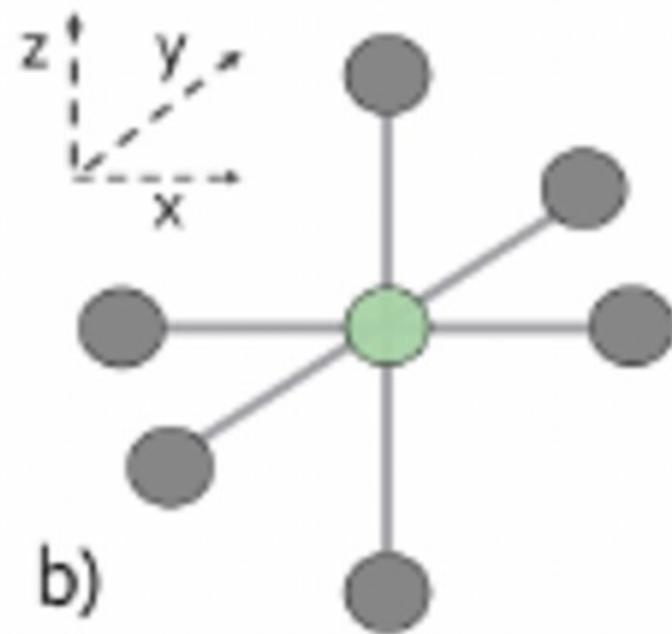
Discretization in three space dimensions.

$$u_{i,j,k}^{n+1} = 2u_{i,j,k}^n - u_{i,j,k}^{n-1} + c^2 \Delta t^2 [D_x D_x u + D_y D_y u + D_z D_z u]_{i,j,k}^n$$

# Stencils



5 point  
stencil



7 point  
stencil

# Regular grids in PETSc

- PETSc allows creation of regular grid in 1, 2 or 3 space dimensions.
- Grids are normal vectors, but support grid based indexing.
- Ghost cells from neighboring processes are automatically synchronized.

```
# Creating a global 2d grid  
da = PETSc.DA().create([10, 10])
```

How to distribute grids in parallel?

# Distributed Arrays

Processor 2

26	27	28	29	30
21	22	23	24	25
16	17	18	19	20
11	12	13	14	15
6	7	8	9	10
1	2	3	4	5

Processor 3

Processor 2

22	23	24	29	30
19	20	21	27	28
16	17	18	25	26
7	8	9	14	15
4	5	6	12	13
1	2	3	10	11

Processor 3

Processor 0

Processor 1

Natural Ordering

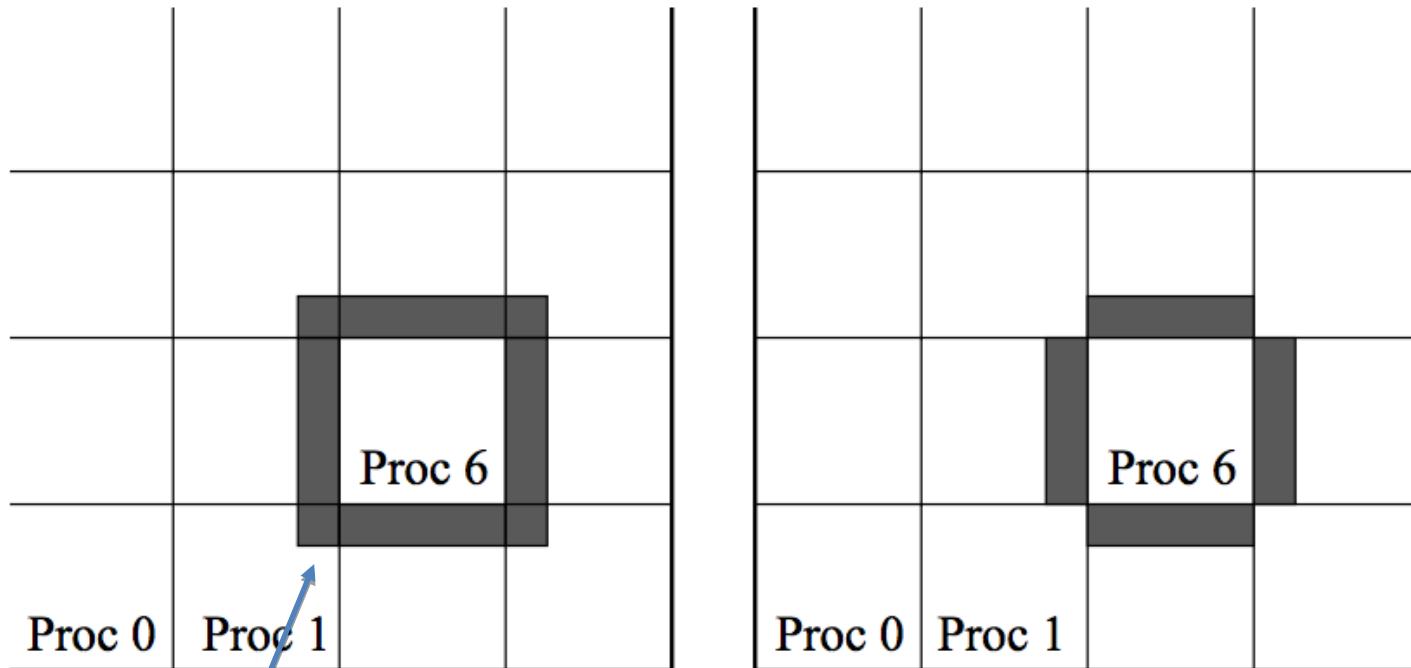
Processor 0

Processor 1

PETSc Ordering

Array distributed on a 2x 2  
process grid

# PETSc Stencil Types



Ghost elements will automatically be handled

PETSc manual, <http://www.mcs.anl.gov/petsc/petsc-current/docs/manual.pdf>

Use Star-type pencil for standard 5 or 7 point Finite Difference stencil